

GRUNDLÆGGENDE PROGRAMMERING

SQL I: Introduction to SQL and Databases!



Claus Brabrand

`(((brabrand@itu.dk)))`

Associate Professor, Ph.D.

`(((Software and Systems)))`

 **IT University of Copenhagen**

A G E N D A

■ Introduction to Databases:

TEORI

- 1) Create a database: 'CREATE TABLE'
- 2) Insert data into database: 'INSERT INTO'
- 3) Query a database: 'SELECT FROM'

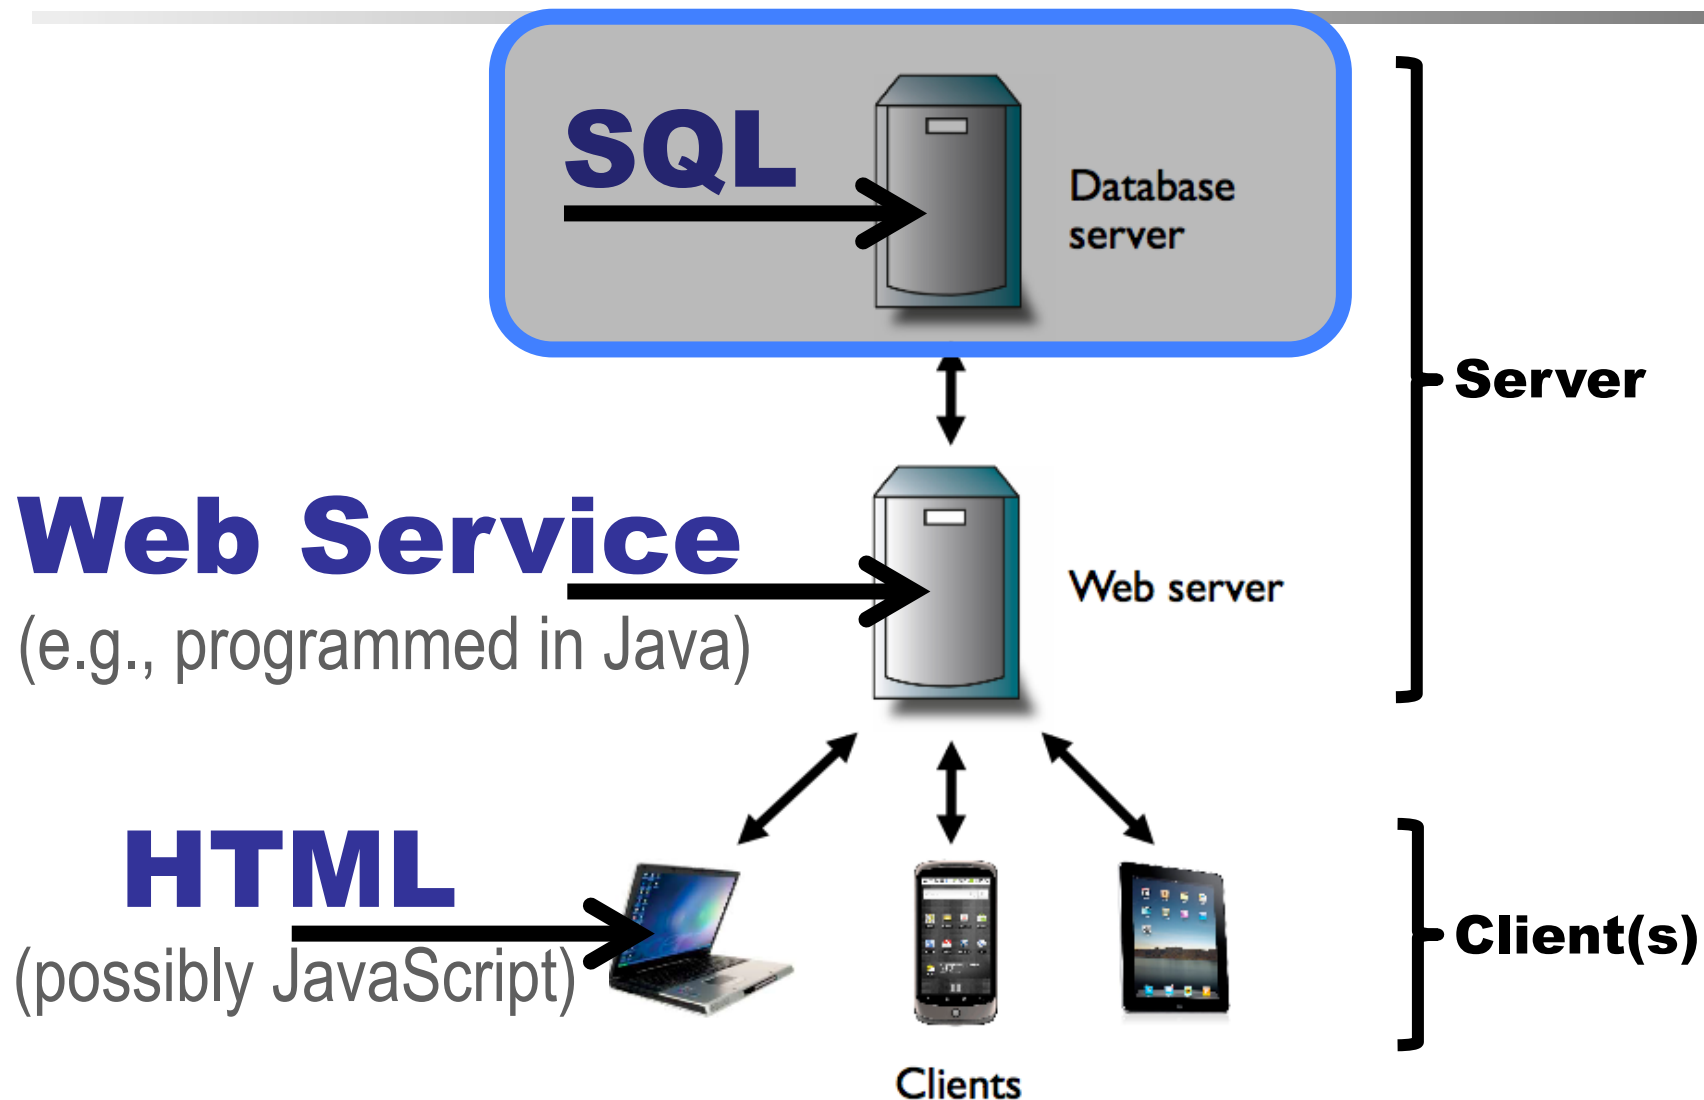
■ How to create + connect to a MySQL database

■ Introduction to Databases:


DEMO

- 1) Create a database: 'CREATE TABLE'
- 2) Insert data into database: 'INSERT INTO'
- 3) Query a database: 'SELECT FROM'

Web Service Architecture



Java Program with Database



SQL (Database Language)

- **SQL** is a (relational) **Database Language**:
 - Originally "SEQUEL" (Structured English Query Language)
 - Invented in the 1970s and popular since 1980s
 - Most common way to access (relational) databases
 - (Other kinds of databases: "*object-oriented*", "*spatial*")

Example:

```
SELECT first_name FROM users WHERE last_name = 'Hansen';
```

- We will use **MySQL** (free SQL implementation):
 - *robust, efficient, free* (www.mysql.com)
 - ...and used a lot in the ***software industry!***

Entities (aka, Tables)

A relational database is built on...:

■ *Entities (aka, Tables):*

(Sort of kind of like an Excel spreadsheet !)

stud_name	stud_id	address	course_id	course_name	teacher
Anna	1234	Somewhere 3	GRPRO	Grundlæggende Programmering	Claus
Bent	0001	Homestreet 4	GRPRO	Grundlæggende Programmering	Claus
Carl	0002	Nowhere 9b	KKRT	IT-Kulturkontekstsreflektionsteori	TBA
Anna	1234	Somewhere 3	BPAK	Projektarbejde og Kommunikation	Mathias

row:
known as
a *record*

column:

has a **name** (here: 'stud_id')
and a **type** (here: 'int')

(svarer til "int stud_id;" i Java)

a table cell is
called a *field*
(and it contains
the actual data)

■ ...and *relations* between entities/tables:

Relations (between Entities)

■ One-to-one relations:

CUSTOMERS:

customer	creditcard_id
Anna	1234567890
Bent	2345678901
Carl	3456789012

one-to-one

CREDIT_CARDS:

creditcard_id	type	expires
1234567890	VISA	07/17
2345678901	Master Card	09/16
3456789012	VISA	11/17

■ One-to-many relations:

ZIPCODES:

zip code	area name
0999	Copenhagen C
2300	Copenhagen S
1257	Copenhagen K

one-to-many

STREETS:

street_name	zip code	speed limit
Amagerbrogade	2300	50
Amalienborg Slotsplads	1257	40
Emil Holms Kanal	0999	20
Rued Langgards Vej	2300	50

*"A street has one zipcode", whereas
"A zipcode has many streets"*

Problem with Many-to-Many

■ Many-to-many relations:

STUDENTS:

student	student_id
Anna	1234
Bent	0001
Carl	0002



COURSES:

course_id	course_name
GRPRO	Grundlæggende Programmering
KKRT	IT-Kulturkontekstsreflektionsteori
BPAK	Projektarbejde og Kommunikation

■ Note: *many-to-many* relations **cannot** be captured in/between tables!

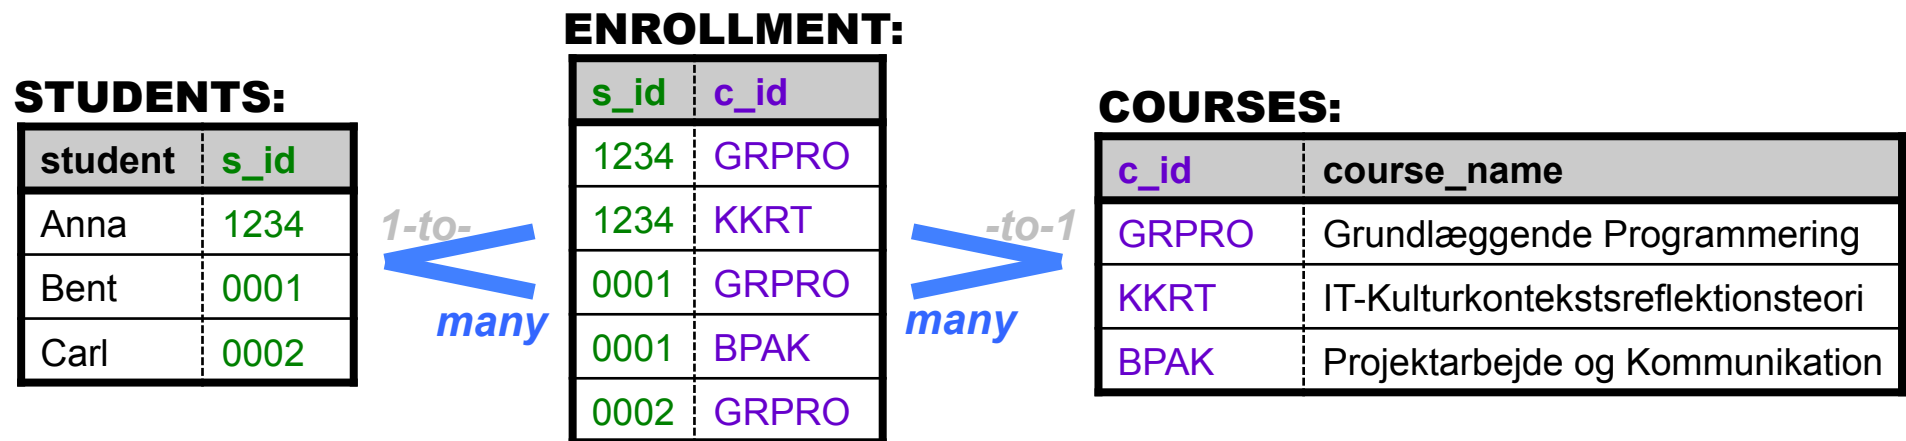
- **Solution**: make a new table which will **be** the many-to-many relation; (i.e., how **students** ↔ **courses**):

ENROLLMENT:

student_id	course_id
1234	GRPRO
1234	KKRT
0001	GRPRO
0001	BPAK
0002	GRPRO

Solution: New Table for M-to-M

- **SOLUTION:** create a *new table* which then *becomes the many-to-many relation:*



- The "enrollment" table now *is the many-to-many relation* (how: **students** ↔ **courses**)!

"A **student** has many **courses**"

"A **course** has many **students**"

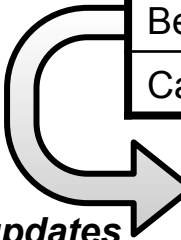
Redudance and Normalization

- Data redundancy is dangerous!
 - Updates might lead to *data inconsistency*!:

CUSTOMERS:

customer	id	address
Anna	111111	Somewhere 3

Be
Ca



updates
may cause
inconsistency!

customer	id	address
Anna	111111	Anotherplace 5
Bent	222222	Homestreet 4
Carl	333333	Nowhere 9b

ACCOUNTS:

id	amount	address
111111	1.250,75	Somewhere 3

22		
33		
id	amount	address
111111	1.250,75	Somewhere 3
222222	587,33	Homestreet 4
333333	-3.125,07	Nowhere 9b

- We ***normalize*** a DB to avoid data redundancy
 - A good DB has little or no data redundancy!

Normalization Example (I/III)

- Let's make a database containing info about students and the courses they are attending:
- Data:
 - student name
 - student id (unique for each student)
 - home address
 - course id (unique for each course)
 - course name
 - course teacher

Normalization Example (II/III)

- We might be tempted to make *one big table*:

stud_name	stud_id	address	course_id	course_name	teacher
Anna	1234	Somewhere 3	GRPRO	Grundlæggende Programmering	Claus
Bent	0001	Homestreet 4	GRPRO	Grundlæggende Programmering	Claus
Carl	0002	Nowhere 9b	BPAK	Projektarbejde og Kommunikation	Mathias
Anna	1234	Somewhere 3	KKRT	IT-Kulturkontekstsreflektionsteori	TBA

- However, such a table would have *lots of redundancy!*

Normalization Example (II/III)

- We might be tempted to make *one big table*:

stud_name	stud_id	address	course_id	course_name	teacher
Anna	1234	Somewhere 3	GRPRO	Grundlæggende Programmering	Claus
Bent	0001	Homestreet 4	GRPRO	Grundlæggende Programmering	Claus
Carl	0002	Nowhere 9b	BPAK	Projektarbejde og Kommunikation	Mathias
Anna	1234	Somewhere 3	KKRT	IT-Kulturkontekstsreflektionsteori	TBA

- However, such a table would have *lots of redundancy!*
- We can eliminate redundancy by breaking it into independent tables and relations:

STUDENTS:

stud_name	stud_id	address
Anna	1234	Somewhere 3
Bent	0001	Homestreet 4
Carl	0002	Nowhere 9b

COURSES:

course_id	course_name	teacher
GRPRO	Grundlæggende Programmering	Claus
KKRT	IT-Kulturkontekstsreflektionsteori	TBA
BPAK	Projektarbejde og Kommunikation	Mathias

Normalization Example (III/III)

STUDENTS:

stud_name	stud_id	address
Anna	1234	Somewhere 3
Bent	0001	Homestreet 4
Carl	0002	Nowhere 9b

COURSES:

course_id	course_name	teacher
GRPRO	Grundlæggende Programmering	Claus
KKRT	IT-Kulturkontekstsreflektionsteori	TBA
BPAK	Projektarbejde og Kommunikation	Mathias

- Now, all we need is *m-to-m relation* (table) that says "*how **students** ↔ **courses** are related*":

ENROLLMENT:

student_id	course_id
1234	GRPRO
0001	GRPRO
0002	BPAK
1234	KKRT

BENEFITS:

- reduced redundancy! :-)
- courses** can exist without **students** :-)
- students** can exist without **courses** :-)

Note: we could have further reduced redundancy in our example (e.g., teacher could get own table)

Exercise: Eliminate Redundancy !

- **Exercise:**
Eliminate redundancy in the following table:

pizza	name	description	ordered_by	creditcard	date
1	Margherita	Tomato and Cheese	Marco	724124825	23/10/2014
15	Carnivore	Meatballs, Beef, Chicken, and Bacon	Stefano	123456789	23/10/2014
22	Vegetariana	Rucola, Cucumber, and Brussels sprouts	Filippo	329942217	22/10/2014
1	Margherita	Tomato and Cheese	Marco	724124825	22/10/2014
15	Carnivore	Meatballs, Beef, Chicken, and Bacon	Filippo	329942217	21/10/2014
1	Margherita	Tomato and Cheese	Marco	724124825	24/12/2013

- ...by "**normalization**" (i.e., break it into):
 - independent tables; *and*
 - relations between them

A G E N D A

■ Introduction to Databases:

TEORI

- 1) Create a database: 'CREATE TABLE'
- 2) Insert data into database: 'INSERT INTO'
- 3) Query a database: 'SELECT FROM'

■ How to create + connect to a MySQL database

■ Introduction to Databases:

DEMO

- 1) Create a database: 'CREATE TABLE'
- 2) Insert data into database: 'INSERT INTO'
- 3) Query a database: 'SELECT FROM'

SQL: Main Commands

■ Data definition:

- CREATE TABLE // creates a new table
- DROP TABLE // deletes a table

■ Data manipulation:

- SELECT .. FROM .. WHERE // retrieves information
- INSERT INTO .. VALUES (..) // insert record(s)
- DELETE FROM // delete record(s)
- UPDATE // changes record(s)

Example:

```
SELECT first_name FROM users WHERE last_name = 'Hansen';
```

- SQL commands are **case in-sensitive** (but it's common to use upper-case)
- Names of tables and records are **case sensitive** (let's use lower-case)

Table Creation (by Example)

■ Example: 'mailing_list'

```
CREATE TABLE mailing_list (  
  name VARCHAR(50) NOT NULL,  
  email VARCHAR(100) NOT NULL  
)
```

mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	obama@hotmail.com
John	jdoe@notmail.com

name of field
(here: 'email')

type of field 'email' (here up to
max 100 chars of variable length)

field in record cannot be 'null'
(i.e., must have a value)

'null' is a *special* SQL value:
≠ 0 (zero) or "" (empty string)

■ Another example: 'phone_numbers'

```
CREATE TABLE phone_numbers (  
  email VARCHAR(100) NOT NULL,  
  phone VARCHAR(20) NOT NULL  
)
```

phone_numbers:

email	phone
brabrand@itu.dk	7218 5076
obama@hotmail.com	212-555-0000
jdoe@notmail.com	123-456-7890

Primary Keys



- Primary key is a field which *uniquely* identifies a record (table row)

- 'email' is a good candidate here:

```
CREATE TABLE mailing_list (  
    name VARCHAR(100) NOT NULL,  
    email VARCHAR(100) PRIMARY KEY  
)
```

mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	obama@hotmail.com
John	jdoe@notmail.com
John	john1928@yahoo.com

(Note: a primary key field automatically implies that it cannot be null)

- **Primary keys** are declared at table creation
- **Note:** without **primary keys**, records cannot be uniquely identified (so we need them!)

Special: Auto Increment

- MySQL special function: AUTO_INCREMENT
 - ...that we can use to *auto-generate* unique values:

```
CREATE TABLE students (  
  stud_id INT PRIMARY KEY AUTO_INCREMENT,  
  name VARCHAR(50) NOT NULL,  
  age INT NOT NULL,  
)
```

students:

stud_id	name	age
1	Anna	20
2	Brian	99
3	Claire	87

- This simplifies insertion of new records:
 - every time a new record is inserted it automatically gets a fresh number (1, 2, 3, ...) so we don't need to bother

Question: Primary Keys?

STUDENTS:

stud_name	stud_id	address
Anna	1234	Somewhere 3
Brian	0001	Homestreet 4
Claire	0002	Nowhere 9b

COURSES:

course_id	course_name	teacher
DSDS	Introduction to Scripting...	Claus
GSD	Global Software Development	Claus
DC	Digital Culture	Somebody

ENROLLMENT:

student_id	course_id
1234	DSDS
1234	GSD
0001	DSDS
0001	DC
0002	DSDS

Common SQL Types

■ Common SQL Types:

Type	Description	Example value
INT	Integer (number)	42
DOUBLE	Floating-point (decimal) number	3.1415
VARCHAR(80)	Text up to 80 characters	"John Doe"
VARCHAR	Text up to 255 characters	"John Doe"
TEXT	Long text (max 65535 characters)	"Once upon a time, ..."
DATE	Date	2014-10-23
TIME	Time	12:59:59
DATETIME	Date & Time	2014-10-23 12:59:59

■ ...and many more (see MySQL specification)

Inserting Records (by Example)

■ Insert records:

```
INSERT INTO mailing_list (name, email)
VALUES ('Claus', 'brabrand@itu.dk') ;

INSERT INTO mailing_list (name, email)
VALUES ('Barack', 'obama@hotmail.com') ;

INSERT INTO mailing_list (name, email)
VALUES ('John', 'johndoe@notmail.com') ;
```

mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	obama@hotmail.com
John	jdoe@notmail.com

```
INSERT INTO phone_numbers (email, phone)
VALUES ('brabrand@itu.dk', '72185076') ;

INSERT INTO phone_numbers (email, phone)
VALUES ('obama@hotmail.com', '212-555-0000') ;

INSERT INTO phone_numbers (email, phone)
VALUES ('jdoe@notmail.com', '123-456-7890') ;
```

phone_numbers:

email	phone
brabrand@itu.dk	7218 5076
obama@hotmail.com	212-555-0000
jdoe@notmail.com	123-456-7890

Info Retrieval (I)

students:

name	id	address
Anna	1234	Somewhere 3
Brian	0001	Homestreet 4
Claire	0002	Nowhere 9b

■ Information retrieval:

```
SELECT * FROM students ;
```



name	id	address
Anna	1234	Somewhere 3
Brian	0001	Homestreet 4
Claire	0002	Nowhere 9b

```
SELECT name, address FROM students ;
```



name	address
Anna	Somewhere 3
Brian	Homestreet 4
Claire	Nowhere 9b

Info Retrieval (II)

students:

name	id	address
Anna	1234	Somewhere 3
Brian	0001	Homestreet 4
Claire	0002	Nowhere 9b

■ Information retrieval:

```
SELECT name FROM students WHERE id > 1;
```

name

Anna

Claire

```
SELECT id, address FROM students WHERE name = 'Brian' ;
```

id

address

0001

Homestreet 4

```
SELECT * FROM students WHERE id = '1234' ;
```

name

id

address

Anna

1234

Somewhere 3

Info Retrieval (III)

students:

name	id	address
Anna	1234	Somewhere 3
Brian	0001	Homestreet 4
Claire	0002	Nowhere 9b

■ Information retrieval:

```
SELECT * FROM students ORDER BY id ;
```



name	id	address
Brian	0001	Homestreet 4
Claire	0002	Nowhere 9b
Anna	1234	Somewhere 3

```
SELECT * FROM students ORDER BY name DESC ;
```

desc (aka descending) => backwards



name	id	address
Claire	0002	Nowhere 9b
Brian	0001	Homestreet 4
Anna	1234	Somewhere 3

Deleting and Updating records

■ Deleting records:

```
DELETE FROM mailing_list WHERE name='John' ;
```

mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	obama@hotmail.com
John	jdoe@notmail.com



mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	obama@hotmail.com

■ Updating records:

```
UPDATE mailing_list SET email='barack@gmail.com' WHERE name='Barack' ;
```

mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	obama@hotmail.com



mailing_list:

name	email
Claus	brabrand@itu.dk
Barack	barack@gmail.com

Dropping Tables



■ Dropping tables:

```
DROP TABLE mailing_list ;
```


mailing_list:

name	email
Claus	brahbrand@itu.dk
Barack	obama@notmail.com
John	jdoe@notmail.com

Remember to
validate input!



SQL Injection Attacks :-)



A G E N D A

■ Introduction to Databases:

TEORI

- 1) Create a database:
- 2) Insert data into database:
- 3) Query a database:

'CREATE TABLE'

'INSERT INTO'

'SELECT FROM'

■ How to create + connect to a MySQL database

■ Introduction to Databases:

DEMO

- 1) Create a database:
- 2) Insert data into database:
- 3) Query a database:

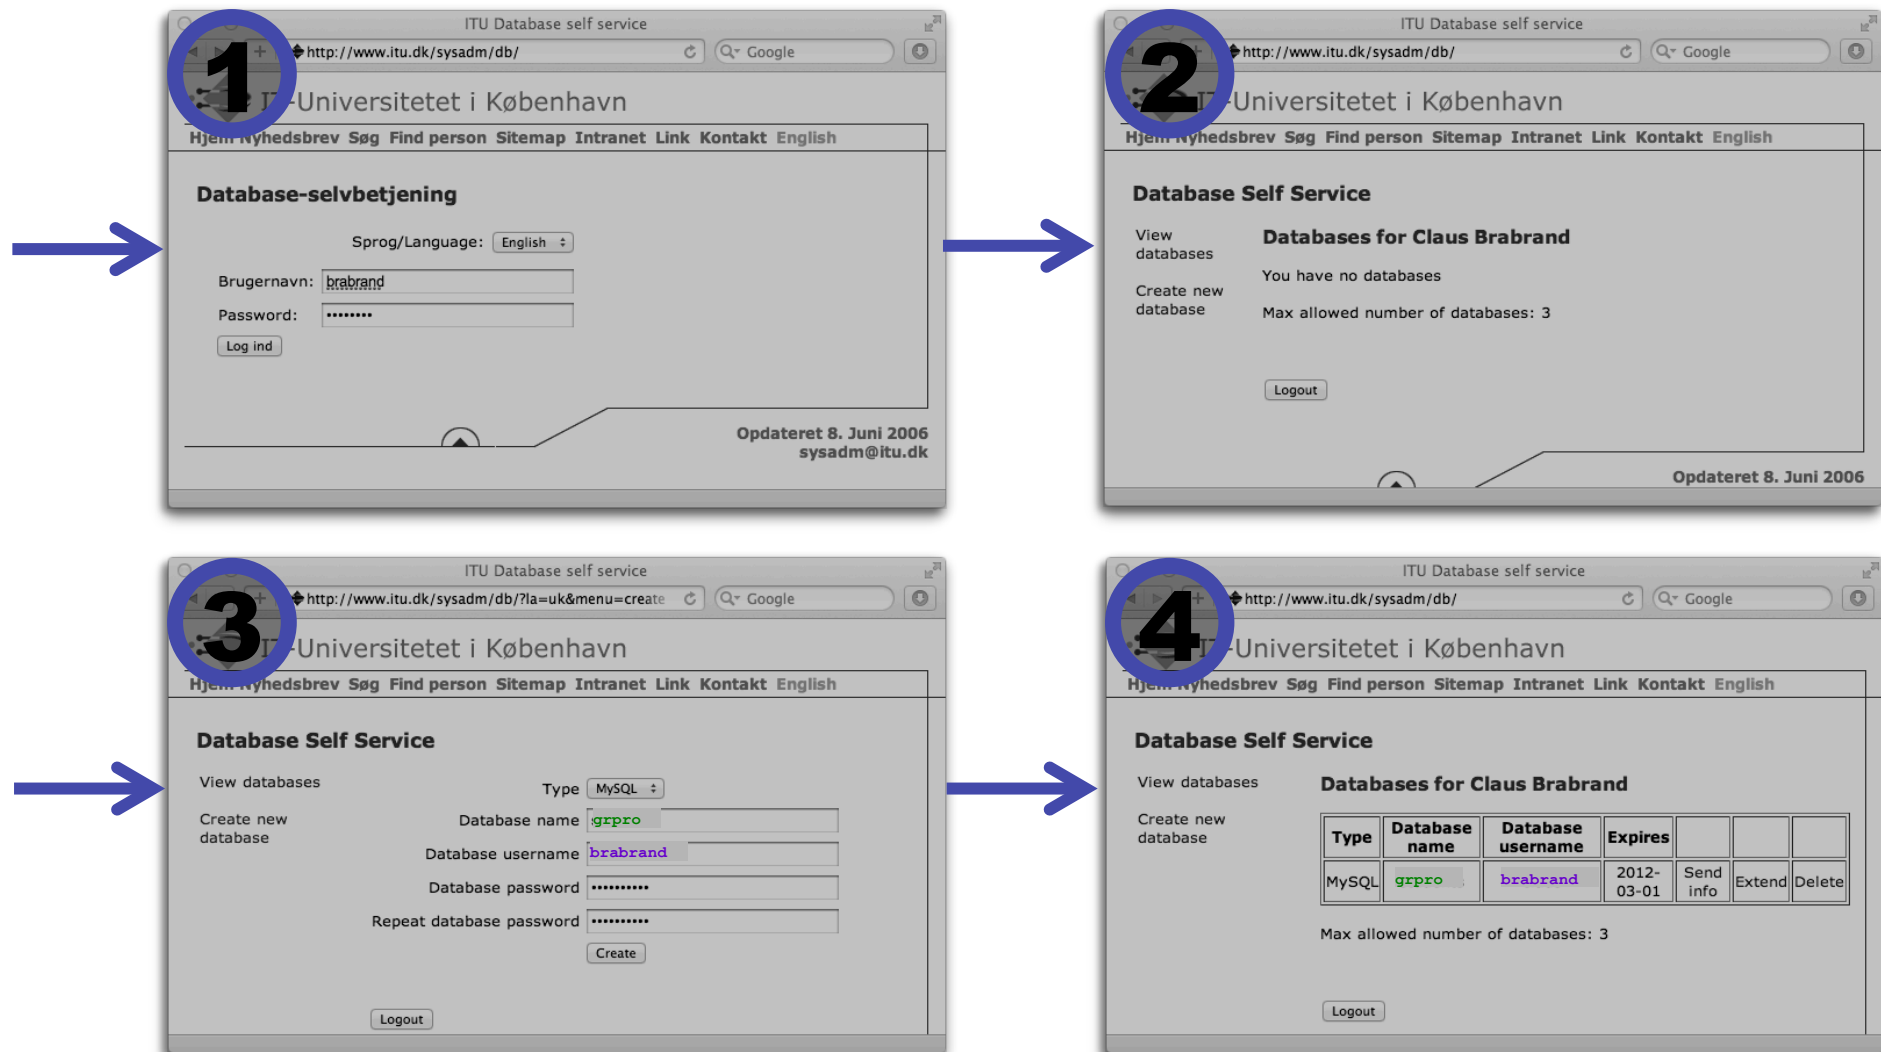
'CREATE TABLE'

'INSERT INTO'

'SELECT FROM'

(<http://www.itu.dk/sysadm/db/>)

SQL Setup



SQL Log in

- Start a "terminal" and write the following:

```
Last login: Fri Nov  4 11:18:54 on ttys000
adm2-66:~ brabrand$ ssh brabrand@ssh.itu.dk
brabrand@ssh.itu.dk's password:
Last login: Fri Nov  4 11:19:10 2011 from 130.226.142.243

[brabrand@cypher ~]$ mysql -h mysql.itu.dk -u brabrand -p grpro
Enter password:
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 9696431
Server version: 5.0.67 Source distribution

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql>
```

*you write **the blue stuff** and the
computer writes **the gray stuff***

A G E N D A

■ Introduction to Databases:

TEORI

- 1) Create a database: 'CREATE TABLE'
- 2) Insert data into database: 'INSERT INTO'
- 3) Query a database: 'SELECT FROM'

■ How to create + connect to a MySQL database

■ Introduction to Databases:

DEMO

- 1) Create a database: 'CREATE TABLE'
- 2) Insert data into database: 'INSERT INTO'
- 3) Query a database: 'SELECT FROM'

SQL Demo: Create & Insert

```
mysql> CREATE TABLE mailing_list ( name VARCHAR(100) NOT NULL,  
->                                     email VARCHAR(100) NOT NULL ) ;  
Query OK, 0 rows affected (0.00 sec)  
  
mysql> INSERT INTO mailing_list (name, email)  
->     VALUES ('Claus', 'brabrand@itu.dk') ;  
Query OK, 1 row affected (0.00 sec)  
  
mysql> INSERT INTO mailing_list (name, email)  
->     VALUES ('Barack', 'obama@hotmail.com') ;  
Query OK, 1 row affected (0.00 sec)  
  
mysql> INSERT INTO mailing_list (name, email)  
->     VALUES ('John', 'jdoe@notmail.com') ;  
Query OK, 1 row affected (0.00 sec)  
  
mysql>
```

SQL Demo: Select

```
mysql> SELECT * FROM mailing_list ;
```

+	-----	+	-----	+
	name		email	
+	-----	+	-----	+
	Claus		brabrand@itu.dk	
	Barack		obama@hotmail.com	
	John		jdoe@notmail.com	
+	-----	+	-----	+

```
3 rows in set (0.00 sec)
```

```
mysql> SELECT name FROM mailing_list ;
```

+	-----	+
	name	
+	-----	+
	Claus	
	Barack	
	John	
+	-----	+

```
3 rows in set (0.00 sec)
```

SQL Demo: Select (cont'd)

```
mysql> SELECT email FROM mailing_list WHERE name='Barack';
```

```
+-----+
| email          |
+-----+
| obama@hotmail.com |
+-----+
1 row in set (0.00 sec)
```

```
mysql> SELECT * FROM mailing_list ORDER BY name DESC ;
```

```
+-----+-----+
| name   | email                |
+-----+-----+
| John   | jdoe@notmail.com     |
| Claus  | brabrand@itu.dk      |
| Barack | obama@hotmail.com    |
+-----+-----+
3 rows in set (0.00 sec)
```

```
mysql>
```

SQL Demo: Delete

```
mysql> DELETE FROM mailing_list WHERE name='John';
```

```
Query OK, 1 row affected (0.00 sec)
```

```
mysql> SELECT * FROM mailing_list ;
```

```
+-----+-----+
| name   | email                |
+-----+-----+
| Claus  | brabrand@itu.dk      |
| Barack | obama@hotmail.com    |
+-----+-----+
```

```
2 rows in set (0.00 sec)
```

```
mysql>
```

SQL Demo: Update

```
mysql> UPDATE mailing_list SET email='barack@gmail.com'
-> WHERE name='Barack' ;
Query OK, 1 row affected (0.00 sec)
```

```
mysql> SELECT * FROM mailing_list ;
```

```
+-----+-----+
| name   | email                |
+-----+-----+
| Claus  | brabrand@itu.dk      |
| Barack | barack@gmail.com     |
+-----+-----+
2 rows in set (0.00 sec)
```

```
mysql>
```

SQL Demo: Show Columns

```
mysql> SHOW COLUMNS FROM mailing_list ;
```

Field	Type	Null	Key	Default	Extra
name	varchar(100)	NO		NULL	
email	varchar(100)	NO		NULL	

```
2 rows in set (0.00 sec)
```

```
Query OK, 1 row affected (0.00 sec)
```

```
mysql>
```

A G E N D A



■ Bonus: More SQL commands

SHOW TABLES & DESCRIBE

- Show tables in database:

```
SHOW TABLES ;
```

Tables_in_my_database

mailing_list
phone_numbers
students

- Show information about a particular table:

```
DESCRIBE students ;
```

Field	Type	Null	Key	Default	Extra
name	varchar(80)	NO		NULL	
age	int(11)	NO		NULL	

- Quit:

```
QUIT ;
```

LIMIT and ORDER-BY

- We can limit the number of records displayed:

- Retrieve a hundred recipes:

```
SELECT * FROM recipes LIMIT 100 ;
```

- Retrieve the youngest student:

```
SELECT * FROM students ORDER BY age LIMIT 1 ;
```

- Retrieve three (alphabetically) last courses:

```
SELECT * FROM courses ORDER BY name DESC LIMIT 3 ;
```

COUNT

- Given the table:

```
SELECT * FROM students ;
```

stud_id	name	age
1234	Anna	20
5678	Brian	25
9999	Claire	23

- COUNT (will count number of records):

```
SELECT COUNT(*) FROM students ;
```

COUNT(*)

3

- COUNT AS (gives the count a name):

```
SELECT COUNT(*) AS num_courses FROM students ;
```

num_courses

3

MIN, MAX, SUM, AVG

stud_id	name	age
1234	Anna	20
5678	Brian	25
9999	Claire	23

■ MIN (minimum):

```
SELECT MIN(age) FROM students ;
```

MIN(age)

20

■ MAX (maximum):

```
SELECT MAX(age) AS oldest FROM students ;
```

oldest

25

■ SUM:

```
SELECT SUM(age) FROM students ;
```

SUM(age)

68

■ AVG (average):

```
SELECT AVG(age) AS average FROM students ;
```

average

22.6667

Thx



- *Questions?* -

Next time: "JOIN" & "GROUP BY"

Exercises 9

■ Exercise 9.1:

- Create own MySQL database
(hosted on the ITU MySQL server)
- Log into it via the MySQL text client

■ Exercise 9.2:

- Create SQL tables for ITU Courses and Teachers
- ...and populate your database with real'ish data

■ Exercise 9.3:

- Try out the various queries from these slides