



# GROUP FELLOWSHIP OF THE RING

## Assignment 2

Dennis Thinh Tan Nguyen William Diedrichsen Marstrand, Thor  
Valentin Olesen, Pernille Lous, Nicklas Johansen

## Problem 1: Relational Algebra

### Task 1

- (b)

### Task 2

- (b)

### Task 3

- The expression joins the two relations Endorsement and Member which gives us a relation with the following attributes: *userID*, *Name*, *skillID*, *prof* and *n* where *n* is less than four.
- There are five columns with the following attributes: *userID*, *name*, *skillID*, *prof* and *n*.
- There are 3 tuples

d)

Name	userID	skillID	prof	n
Jesper	U001	S4	15	1
Jan	U006	S4	21	1
Jens	U005	S4	31	3

### Task 4

- The expression selects all *userID* and *skillID* from the table Endorsement and divides the relation with the *skillID* belonging to the *userID* U005 in the table Endorsement.
- One column with *UserID*.
- 5 tuples

d)

userID
U001
U002
U003
U005
U006

### Task 5

- a) Creates two temporary tables of Endorsement named E and E1.  
Joins E and E1 where *skillID* in E is the same as *skillID* in E1 and *userID* in E is greater than *userID* in E1.  
Then selects all *userID* in E and E1.
- b) A table with two columns *SkillID* and *userID*
- c) There are 10 tuples

userID.E	userID.E1
U002	U001
U003	U001
U003	U002
U005	U001
U005	U002
U005	U003
U006	U001
U006	U002
U006	U003
U006	U005

d)

## Problem 2: Relational Tuple Calculus

### Task 1

- a) The expression gets a tuple t from the table endorsement where the *skillID* is equal to S1.
- b) 4 columns named *userID*, *skillID*, *prof* and *n*.
- c) 1 tuple

userID	skillID	prof	n
U003	S1	18	5

d)

### Task 2

- a) This question contains errors and has there been omitted (*By request of TA.*)

### Task 3

- a) This expression gives a set of tuples where *userID* from endorsement is equal to *userID* from Members and *skillID* in endorsement tuple is equal to "S1" and *prof* in endorsement is greater than 20. *name* in tuple t is equal to *name* in Member. Since there is in fact only one tuple with *skillID* equal to "S1", but the tuple has a *prof* attribute less than 20, t will be an empty set.
- b) Since t is an empty set the answer is 0
- c) No tables has been created.

## Problem 3: SQL Queries

### Exercise 1:

```
SELECT COUNT(*) FROM danishMovies;
```

### Exercise 2:

```
SELECT year, COUNT(imdbVotes) FROM movie GROUP BY year;
```

### Exercise 3:

```
-- Helper-query:
```

```
SELECT DISTINCT role FROM involved;
```

```
--Only roles are actor and director
```

```
-- Query:
```

```
SELECT COUNT(*) FROM person WHERE name LIKE 'C%';
```

```
-- Comment: DISTINCT could have been used upon duplicate occurrences.
```

### Exercise 4:

```
SET @pulpID = (SELECT id FROM movie WHERE title LIKE 'Pulp Fiction');
SELECT name,birthdate FROM person
INNER JOIN involved ON personId = id WHERE movieId = @pulpID AND birthdate IS NOT NULL ORDER BY
birthdate;
```

```
--Some of the persons dosen't have a birthdate in the database. These are thrown away in the
query.
```

### Exercise 5:

```
SET @UmaID = (select id from person where name like 'Uma Thurman'); #Uma ID
SET @JohnID = (select id from person where name like 'John Travolta'); #John ID
```

```
--Exclusive starring
```

```
SELECT title FROM
  (SELECT movieId FROM involved WHERE personId=@UmaID
   AND movieId IN
     (SELECT movieId FROM involved WHERE personId=@JohnID))
  AS ids LEFT JOIN movie ON ids.movieId=id;
```

**Exercise 6:**

```
SELECT name FROM person, involved AS personInv1
WHERE person.id = personInv1.personId AND personInv1.role = "Actor" AND personInv1.movieId =
(SELECT id FROM movie WHERE title = "Pulp Fiction")AND NOT EXISTS
(SELECT 1 FROM involved AS personInv2, involved AS personInv3, involved AS personInv4 WHERE
personInv2.role = "Actor" AND personInv3.role = "Actor" AND personInv4.role = "Actor" AND
personInv2.movieId = personInv1.movieId AND personInv2.personId != personInv1.personId AND
personInv3.personId = personInv1.personId AND personInv4.personId = personInv2.personId AND
personInv3.movieId != personInv1.movieId AND personInv4.movieId = personInv3.movieId);
```

*-- This should return an empty list but we seem to get an endless loop.*

**Exercise 7:**

```
SET @SamuelID = (select id from person where name like 'Samadu Jackson');
SET @JohnID = (select id from person where name like 'John Travolta');

SELECT title, year FROM
(select movieId from involved where personId=@SamuelID
AND movieId IN
(select movieId from involved where personId=@JohnID))
AS ids LEFT JOIN movie on ids.movieId=id WHERE year >= 1980 ORDER BY title, year;
```

**Exercise 8:**

```
SELECT title, ImdbRank FROM movie WHERE year >= 1990
AND year < 2000 ORDER BY imdbRank DESC LIMIT 5;
```

**Exercise 9:**

```
SELECT movie.language, ROUND(AVG(imdbRank), 2) AS 'Average Rating'
FROM movie WHERE movie.year = 1994 GROUP BY movie.language ORDER BY AVG(imdbRank) DESC;
```

**Exercise 10:**

```
SET @JohnId = (SELECT id FROM person WHERE name LIKE 'John Travolta');
SELECT TITLE, MAX(imdbRank) FROM movie WHERE id
IN(SELECT movieId FROM involved WHERE personId = @JohnId);
```

**Exercise 11:**

```
SET @CCbirth = (SELECT birthdate FROM person WHERE name LIKE 'Charles Chaplin');
SET @CCdeath = (SELECT deathdate FROM person WHERE name LIKE 'Charles Chaplin');

SELECT count(*) FROM person WHERE gender = 'f' AND birthdate>@CCdeath OR deathdate<@CCbirth;
```

**Exercise 12:**

```
SELECT name, MIN(height) FROM person WHERE id IN (SELECT personId FROM involved WHERE movieId in
(SELECT movieId FROM genre WHERE genre='action') AND role='actor');
```

**Exercise 13:**

```
SELECT genre, AVG(imdbRank) FROM (SELECT * FROM genre INNER JOIN movie ON genre.movieId = movie.id) as avgRankByGenre GROUP BY avgRankByGenre.genre;
```

**Exercise 14:**

```
SELECT genre.genre, sum(approvedMovies.count) votes FROM genre,  
    (SELECT movie.id, userRating.count FROM movie,  
    (SELECT movie.id, count(*) count FROM movie, ratings WHERE movie.id = ratings.movieId GROUP  
    BY movie.id)  
userRating WHERE movie.id = userRating.id AND userRating.count > 10) AS approvedMovies WHERE  
approvedMovies.id = genre.movieId GROUP BY genre.genre ORDER BY votes DESC;
```

**Exercise 15:**

```
SELECT movie.title, count(mr2.toId) AS refNumber FROM movieref mr1, movieref mr2, movie  
WHERE mr1.fromId = movie.id and mr1.toId = mr2.fromId GROUP BY mr1.fromId  
ORDER BY refNumber desc LIMIT 1;
```

**Exercise 16:**

```
SELECT Count(*) AS "ActorAsDirector" FROM person  
WHERE person.id IN (SELECT Distinct personID FROM involved WHERE role = 'director')  
AND person.id IN (SELECT distinct personID FROM involved WHERE role = 'actor');
```

**Exercise 17:**

```
SELECT a.genre, b.genre, count(*) FROM genre a INNER JOIN genre b ON a.movieId = b.movieId  
AND a.genre < b.genre GROUP BY a.genre, b.genre ORDER BY count(*) DESC LIMIT 1
```

## Problem 4: Indexing

To determine what tables would benefit from indexing, we have looked at the size of each table;

### Size chart

<b>Involved</b>	1.997.903
<b>Person</b>	167.010
<b>Movieref</b>	140.820
<b>Genre</b>	129.962
<b>Movie</b>	59.284
<b>Ratings</b>	486
<b>danishMovies</b>	251
<b>Popular</b>	74

### Example

This tells us that the involved table should probably be indexed if used extensively in our queries. By way of example, exercise 6 requires for us to find all people starred in the '*Pulp Fiction*' movie and filter out those who are starred together only in that movie. This requires extensive use of lookups in the involved table that holds the *id* for movies and people. The query takes around 1 minute without indexing. However, if we create an index on *movieId*, *personId* and a composite index of both, it would only take 0,01 seconds.

### Index creation

**Notation:** CREATE index <indexname> ON <tablename>(<attribute>);

Indices:

```
CREATE index movieIndex ON involved (movieId);
CREATE index personIndex ON involved (personId);
CREATE index personMovieIndex ON involved (personId,movieId);
```

### EXPLAIN

```
mysql> explain select distinct movieId from involved where personId in (select personId from involved where movieId = @PulpFictionId);
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | involved | ref | movieIndex, personIndex, personMovieCompIndex | movieIndex | 5 | const | 45 | Using where; Using temporary; Start temporary |
| 1 | SIMPLE | involved | ref | movieIndex, personIndex, personMovieCompIndex | personMovieCompIndex | 5 | imdb.involved.personId | 5 | Using index; End temporary |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
2 rows in set (0.00 sec)
```

The '*EXPLAIN*' keyword shows us that the movieIndex and composite key of movieIndex and personIndex is used to execute the query.

### Timing

Query for finding movies that people from Pulp Fiction have starred in:

```
SET @PulpFictionId = (select id from movie where title = 'pulp fiction');  
SELECT DISTINCT movieId from involved WHERE personID IN  
    (SELECT personId FROM involved WHERE movieId = @PulpFictionId);
```

	Time
<b>Without Indexing</b>	~73 seconds
<b>With indexing</b>	0,01 seconds
<b>Time to create indices</b>	10 seconds

To conclude, the index speeds up the query from 1 minute and 13 seconds to 0,01 seconds.