

# Hashing and Sorting

Carsten Schürmann

# Example

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

- Find the a student with CPR number 151197-2352

# Idea

Compute the location from the key!!!

Collision resolution method

151197-2352

Hash function

Utilization

0	
1	
2	
3	
4	151197-2352, Claus, Dragør
5	
6	
7	
8	

Hash table size

Hash bucket

# Hash Functions

Notation:  $h(x)$

Goal: Uniform distribution over buckets

Division hashing  $h(x) = (a * x + b) \bmod M$

Example:  $h(cpr) = cpr \bmod 1000$

# Choice of Parameters

$$h(x) = (a * x + b) \bmod M$$

- $a$  arbitrary
- $b$  arbitrary
- $M$  prime number



Concern:  $M = 2$ , last digit of cpr reveals sex

# Size of the Hash Table

Example:

50000 students, 10 students per page


Utilization: 90%

In average  $50000/10/0.9 = 5555$

Pick  $M$  as closest prime to 5555.

# Collision Resolution

- Linear probing
- Re-hashing
- Separate chaining




# Extendible Hashing

Keys are turned into bitstrings

$$h(k1) = 100100$$

$$h(k2) = 010110$$


$$h(k3) = 110110$$



Directory



# Inserting $k_3$ , doubling the buckets




Next, insert  $h(k_4) = 011110$

# Idea of Depth

$$h(k4) = 011110$$


Global depth  
Local depth



# Overfull 01 bucket

$$h(k_4) = 011110$$

Global depth  
Local depth

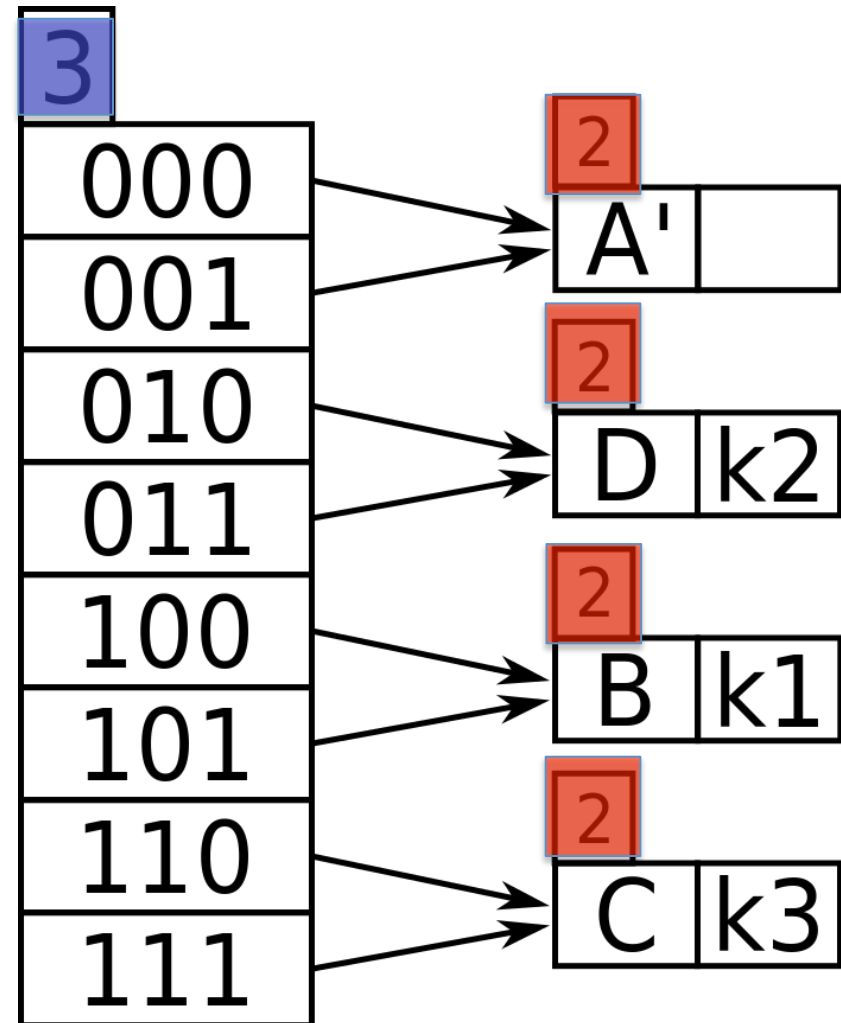


# Splitting

$$h(k_4) = 011110$$

Global depth

Local depth

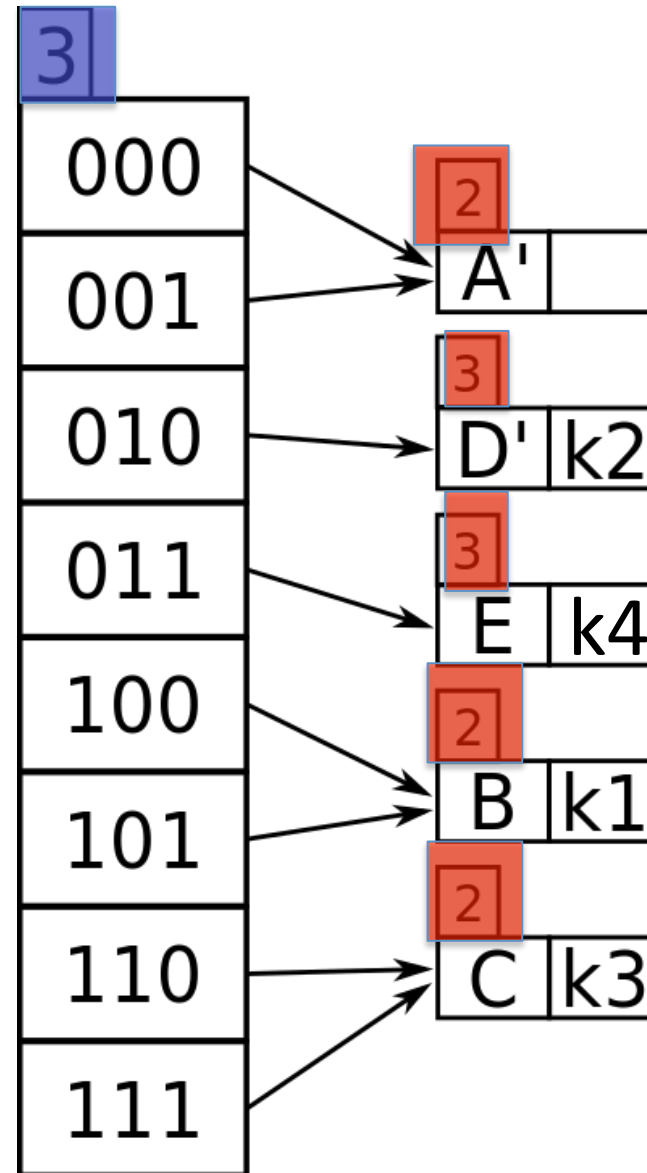


# Insertion

$$h(k4) = 011110$$

Global depth

Local depth



# Summary

## Extendible hashing


- Directory doubles on demand
- Shrink on demand
- Local and global depth as indicators
  - If both are the same, double

# Linear Hashing

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13

Utilization: 50%


Splitting policy: Hash table  $> 75\%$  full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13

Utilization: 50%

Splitting policy: Hash table > 75% full



Insert **10**, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8, <b>10</b>
1	13

Utilization: 75%

Splitting policy: Hash table > 75% full

Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8, 10
1	13, 15

Utilization: 100%

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8, 10
1	13, 15
2	

mod 4

mod 2

mod 4

Utilization: 100%

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13, 15
2	10

mod 4

mod 2

mod 4

Utilization: 100%

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13, 15
2	10

mod 4

mod 2

mod 4

Utilization: 4/6

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13, 15
2	10

mod 4

mod 2

mod 4

Utilization: 4/6

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13, 15
2	10

mod 4  
mod 2  
mod 4

Overflow
19

Utilization: 5/6

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13, 15
2	10
3	

mod 4

mod 4

mod 4

mod 4

Overflow

19

Utilization: 5/6

Splitting policy: Hash table > 75% full



# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13
2	10
3	15, 19

mod 4

mod 4

mod 4

mod 4

Utilization: 5/8


Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13
2	10
3	15, 19

mod 4

mod 4

mod 4

mod 4

Utilization: 5/8


Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13
2	10, 22
3	15, 19

mod 4

mod 4

mod 4

mod 4

Utilization:  $6/8 = 75\%$

Splitting policy: Hash table  $> 75\%$  full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13
2	10, 22
3	15, 19

mod 4  
mod 4  
mod 4  
mod 4

Overflow
18

Utilization: 7/8

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13
2	10, 22
3	15, 19

mod 8

mod 4

mod 4

mod 4

mod 8

Overflow

18

Utilization: 7/16

Splitting policy: Hash table > 75% full

# Insert 10, 15, 19, 22, 18

Number of Buckets: 2

Bucket Capacity: 2

Next bucket  
to be split



Bucket	Content
0	8
1	13
2	10, 22
3	15, 19

mod 8  
mod 4  
mod 4  
mod 4  
mod 8

Overflow
18

Utilization: 7/16

Splitting policy: Hash table > 75% full

# Summary

- Linear hashing
  - Algorithms for search, insertion and deletion
  - Growth and contracts one bucket at the time

# B+ trees versus Hashing

- Speed search
    - Exact match queries
    - Range queries
  - No reorganization
- Speed
    - Very fast on exact match queries
  - No good behavior on
    - Range queries
    - Ordered by queries



# Hash Index Support

- Some DBMS' support hash indexes

```
create index test3  
on Movie(id)  
using hash;
```

- But not mysql.

# Sorting

# When To Sort

- Duplicate elimination
- `select ... order by`
- `select ... group by`
- Sort-merge join from last lecture involves sorting

What to do when the data does  
not fit into memory?

# Two Way External Sort

- Files consist of  $n$  pages
- Buffer size  $b$


## Algorithm

- Pass  $0$  : Read a page, sort it, Write it
- Pass  $i$  :
  - Read two input pages (in one time step)
  - Write one output page (in two time steps)


# Step 0

3,4	6,2	9,4	8,7	5,6	3,1	2
3,4	2,6	4,9	7,8	5,6	1,3	2

# Step 1



# Step 2



3,4 6,2 9,4 8,7 5,6 3,1 2

3,4 2,6 4,9 7,8 5,6 1,3 2

2,3  
4,6

4,7  
8,9

1,3  
5,6

2

2,3  
4,4  
6,7  
8,9


1,2  
2,3  
3,4  
4,5  
6,6  
7,8  
9

1,2  
3,5  
6





# Step 3

- Each pass:
  - Read each page
  - write each page
- Number of passes  
 $\lceil \log_2 n \rceil + 1$
- I/O cost  
 $2n (\lceil \log_2 n \rceil + 1)$
- Note: This scales also to combining  $n$  pages



# Heap Order


A heap is a data structure where left and right subtree contain only smaller elements than the root.



# Heap Sort

- Alternative way to sort
- Used in databases
- Guarantees a more continuous use of **b** buffers
- Supports also double buffering

More about this in RG  
Chapter 13



# State of the Art in Sorting

Current Champions 2014 (tie):

- TritonSort (UCSD)
  - 100 TB in 1,378 seconds (4.3 TB/min)
  - 186 Amazon EC2 i2.8xlarge nodes
- Apache Spark
  - 100 TB in 1,406 seconds (4.27 TB/min)
  - 207 Amazon EC2 i2.8xlarge nodes

More Info: <http://sortbenchmark.org>