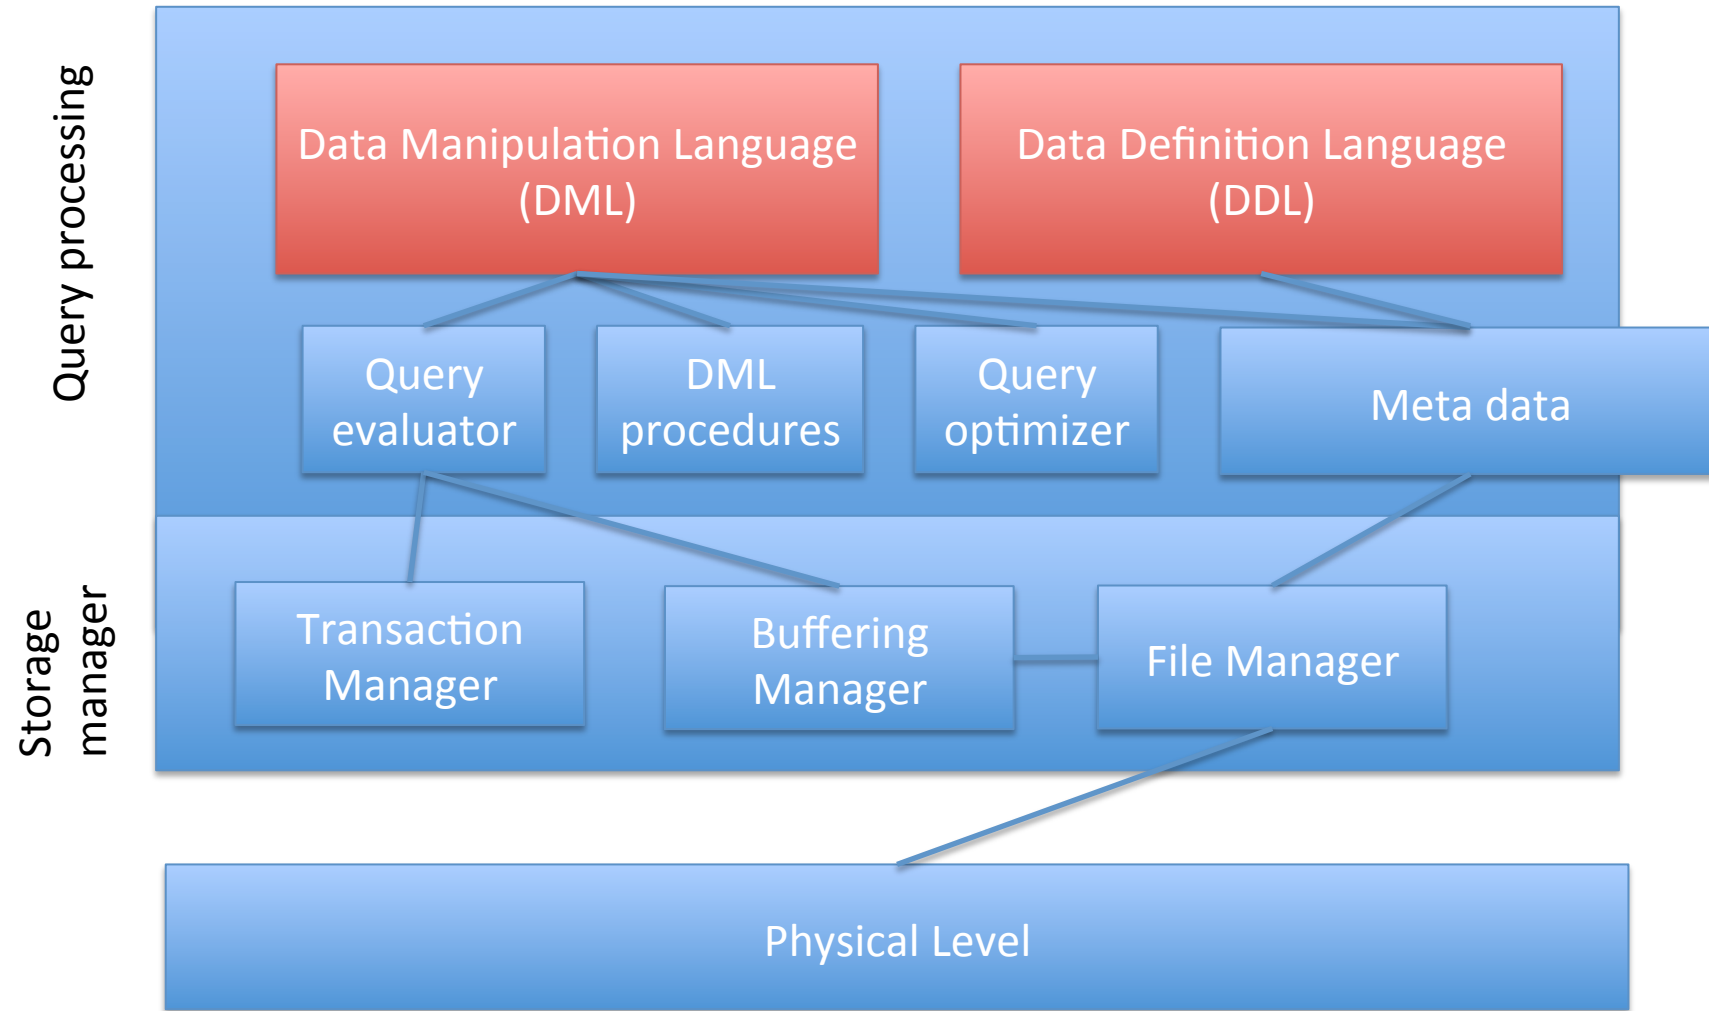


# SQL

Carsten Schürmann

# Database System Architecture



# SQL

The most widely used relational query language.

- Major standard is SQL-1999 (=SQL3)
- Introduced “Object-Relational” concepts
- SQL 2003, SQL 2008 have small extensions
- SQL 92 is a basic subset

# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregate functions
- Nested subqueries
- Insertion
- Joins

# SQL Syntax

**select** A1, A2, ... An

**from** R1, R2, ... Rm

**where** P

# Our Database for Today

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12

## Class

course	name	etcs
SIDD	Databases	7.5
CSYS	Critical Systems	7.5

# Query

Find the cprs of all students named Claus!

$$\pi_{\text{cpr}}(\sigma_{\text{name} = \text{'Claus'}}(\text{Student}))$$

```
select cpr
from Student
where name = "Claus"
```

# General Form

$$\pi_{A_1, \dots, A_n}(\sigma_P(R_1 \times R_2 \dots R_m))$$

select  $A_1, \dots, A_n$   
from  $R_1, R_2 \dots R_m$   
where  $P$



# Looking at the Condition

Find the cprs of all students named Claus from Copenhagen!

$\pi_{\text{cpr}}(\sigma_{\text{name} = \text{'Claus'} \wedge \text{address} = \text{'Copenhagen'}}(\text{Student}))$

```
select  cpr
from    Student
where   name = "Claus"
and     address = "Copenhagen"
```

# More logical operations

Boolean operators (and or not ...)

Comparison operators (< > = ...)

String operators (like % \_ ...)

Date operators (DATE, YEAR, ...)

Find all people at addresses starting with "D"

```
select cpr
```

```
from Student
```

```
where address like "D%"
```

# Our Database for Today

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12

## Class

course	name	etcs
SIDD	Databases	7.5
CSYS	Critical Systems	7.5

# Query

Find the names of all students taking SIDD!

```
select name
from     Student, Takes
where    Student.cpr = Takes.cpr
and      Takes.class = "SIDD"
```

# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregate
- Nested subqueries
- Insertion
- Joins

# Explicit Renaming as Shorthand

Find the names of all students taking SIDD!

```
select name
from    Student as S, Takes as T
where   S.cpr = T.cpr
and     T.class = "SIDD"
```

# Explicit Renaming for Self-Join

Who are Jesper's  
grand-parents?

Parent

cpr	parent	child
140298-1234	Jesper	Lotte
041297-5367	Nikoline	Jesper
151197-2352	Claus	Nikoline
050596-1142	Martin	Lotte

```
select S.name
from   Parent as S, Parent as T
where  S.child = T.parent
and    T.child = "Jesper"
```

# Explicit Renaming for Self-Join

Find courses with more ECTS than SIDD

```
select S.course  
from   Class as S, Class as T  
where  S.ects > T.ects  
and    T.course = "SIDD"
```



# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregate
- Nested subqueries
- Insertion
- Joins

# Our Database for Today

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12

## Class

course	name	etcs
SIDD	Databases	7.5
CSYS	Critical Systems	7.5

# Set operations

Find the CPRs of people taking both SIDD and Critical Systems

```
(select cpr  
  from   takes  
  where  course = "SIDD")
```

**intersect**

```
(select cpr  
  from   takes  
  where  course = "CSYS")
```

# Set operations

Find the CPRs of people taking SIDD or Critical Systems

```
(select cpr  
  from takes  
 where course = "SIDD")
```

union

```
(select cpr  
  from takes  
 where course = "CSYS")
```

# Set operations

Find the CPRs of people taking both SIDD but not Critical Systems

```
(select cpr  
  from takes  
 where course = "SIDD")
```

except

```
(select cpr  
  from takes  
 where course = "CSYS")
```

# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregate
- Nested subqueries
- Insertion
- Joins

# Ordering

Find all students sorted by name!

```
select *  
from Student  
order by name asc
```

asc is default

# Ordering

Find all students lexicographically sorted by name and cpr (in reverse)

```
select *  
from Student  
order by name, cpr desc
```



# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- **Aggregates**
- Nested subqueries
- Insertion
- Joins

# Aggregates

- How to compute the average grade for SIDD?
- How to compute how many students are taking SIDD?
- How to compute the worst grade CSYS?

We are leaving the confines of relational algebra!

# Aggregates

How to compute the average grade for SIDD?

Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12
140789-7612	SIDD	10
140789-7612	CSYS	4
041297-5367	CSYS	10
010292-3333	SIDD	4

```
select avg(grade)
```

```
from Takes
```

```
Where course = "SIDD"
```

avg(grade)
9

# Aggregates

- Average `avg (grade)`
- Counting `count (grade)`
- Minimum `min (grade)`
- Maximum `max (grade)`

None of these aggregate functions is expressible in relational algebra.

# Aggregates

Find the total number of students for each course

Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12
140789-7612	SIDD	10
140789-7612	CSYS	4
041297-5367	CSYS	10
010292-3333	SIDD	4

```
select course, count(*)  
from Takes  
group by course
```

course	count(*)
SIDD	4
CSYS	2

# Aggregates - Renaming

Find the total number of students for each course

Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12
140789-7612	SIDD	10
140789-7612	CSYS	4
041297-5367	CSYS	10
010292-3333	SIDD	4

```
select course, count(*) as size
```

```
from Takes
```

```
group by course
```

course	size
SIDD	4
CSYS	2

# Aggregates

What is the grade average of each student?

Takes

	cpr	course	grade	
	140298-1234	SIDD	10	
	041297-5367	SIDD	12	
	140789-7612	SIDD	10	
	140789-7612	CSYS	4	
	041297-5367	CSYS	10	
	010292-3333	SIDD	4	

```
select cpr, avg(grade)
from Takes
group by cpr
```

cpr	avg(grade)
140298-1234	10
041297-5367	11
140789-7612	7
010292-3333	4

# Aggregates

Which students have  
an average grade  
above 10?

Takes

	cpr	course	grade	
	140298-1234	SIDD	10	
	041297-5367	SIDD	12	
	140789-7612	SIDD	10	
	140789-7612	CSYS	4	
	041297-5367	CSYS	10	
	010292-3333	SIDD	4	

```
select cpr, avg(grade)
from Takes
group by cpr
having avg(grade) > 10
```

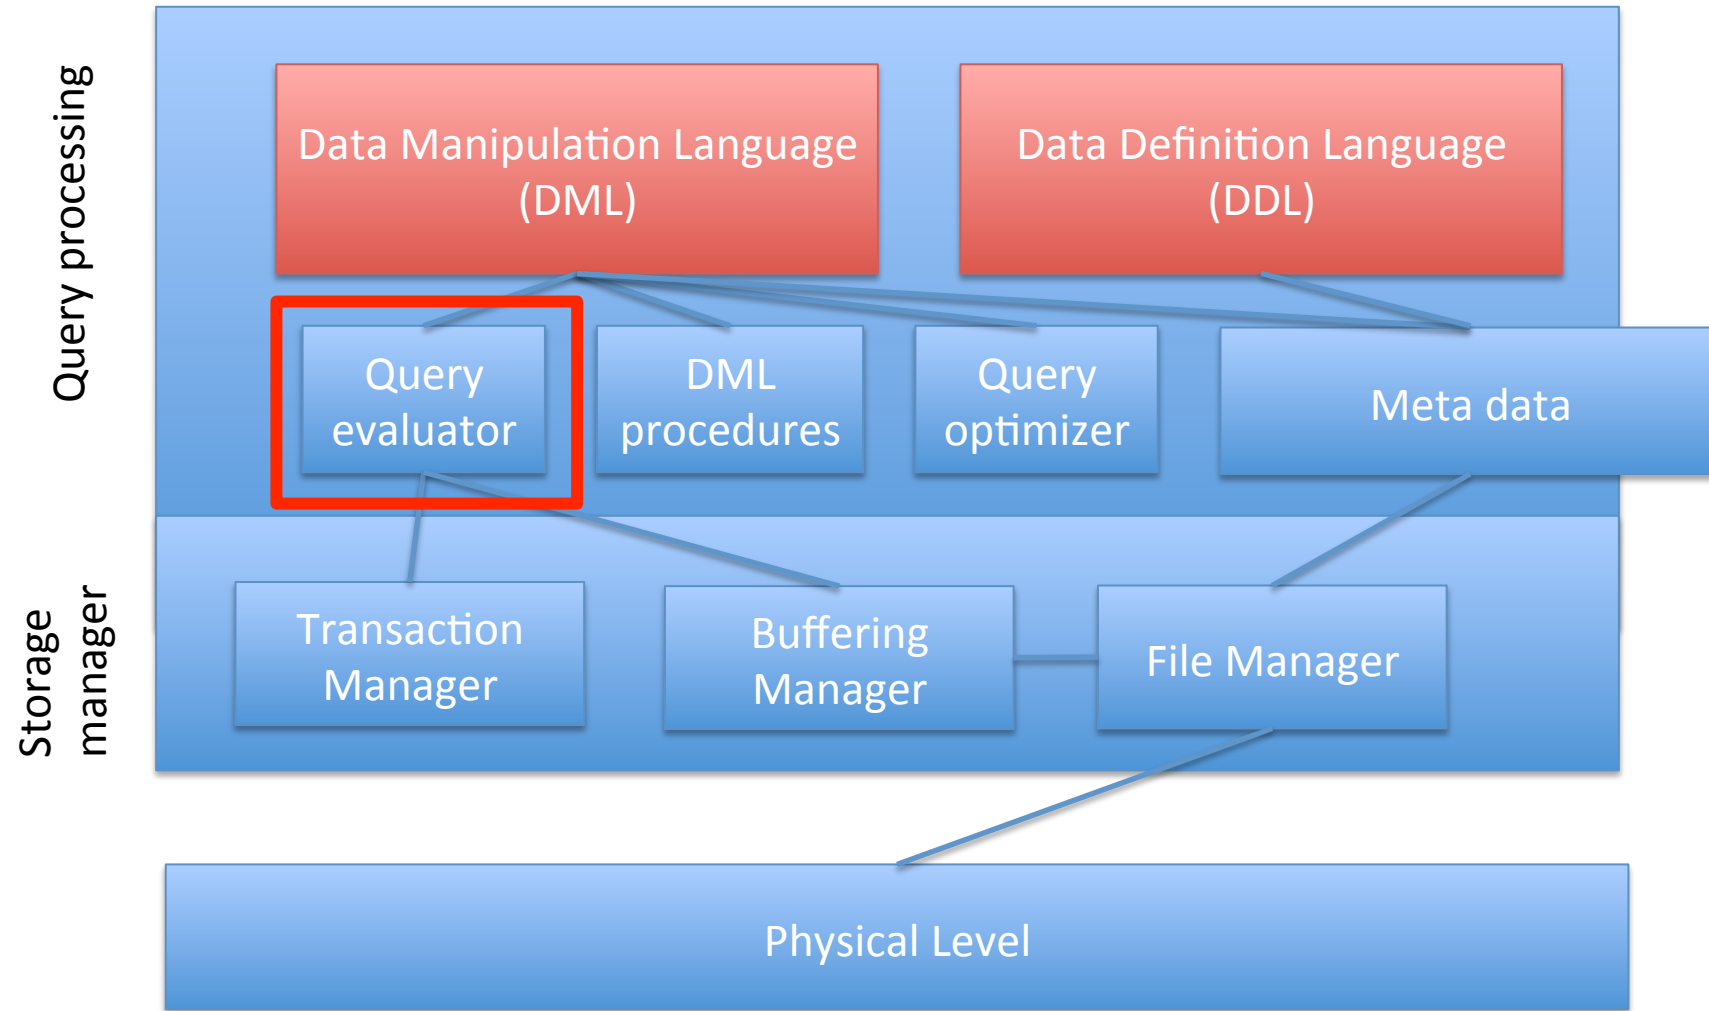
cpr	avg(grade)
140298-1234	10
041297-5367	11
140789-7612	7
010292-3333	4

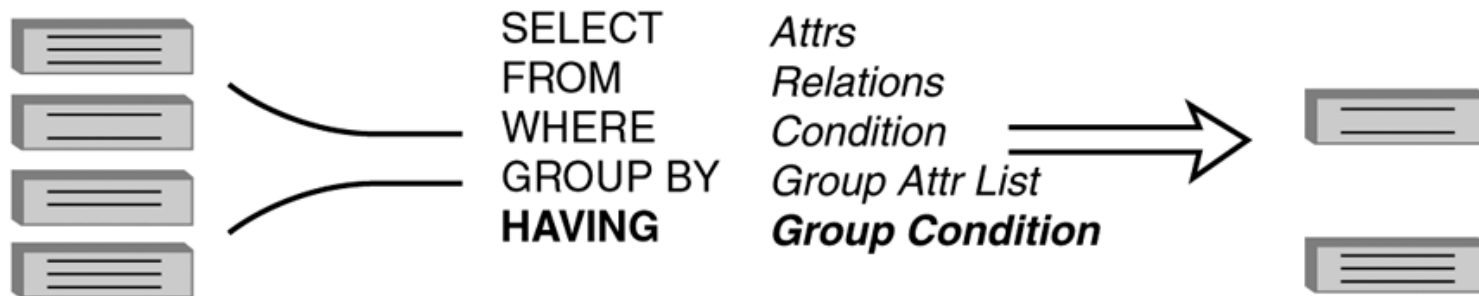
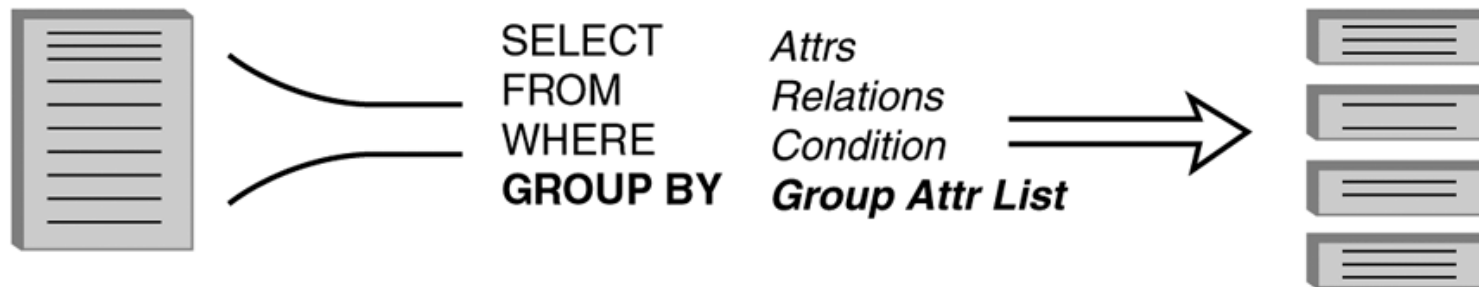
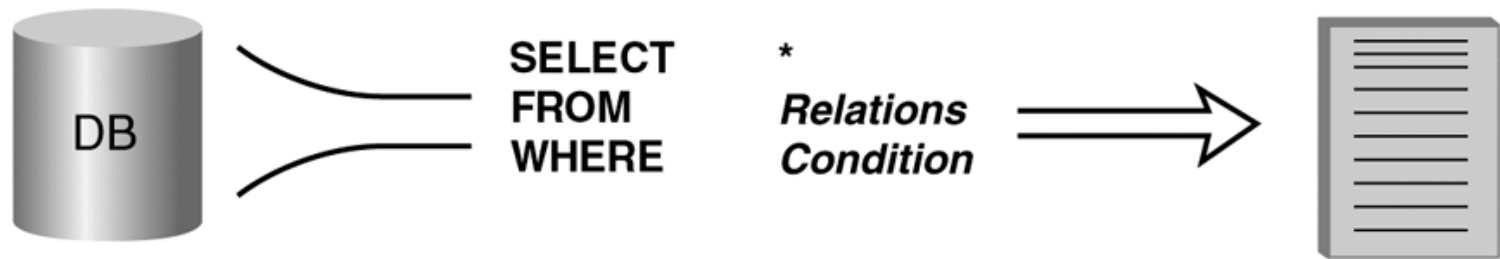


# Refined SQL Syntax

**select** A1, A2, ... An  
**from** R1, R2, ... Rm  
**where** P  
**order by** B1, B2, ... Bn  
**group by** C1, C2, ... Cn  
**having** Q

# Database System Architecture





# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregates
- Nested subqueries
- Insertion
- Joins

# Our Database for Today

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12

## Class

course	name	etcs
SIDD	Databases	7.5
CSYS	Critical Systems	7.5

# Nested Subqueries

Find all **names** of students **who take SIDD**.

Alternatively

Find the **names** for the **CPRs** **that are contained**  
**in the set of of CPRs** of students **who take SIDD**.

# Nested Subqueries

Find the **names** for the **CPRs** that are contained in the set of of **CPRs** of students who take **SIDD**.

```
select name
from Student
where cpr "is contained in"
    select cpr
    from Takes
    where course = "SIDD"
```

# Nested Subqueries

Find the **names** for the **CPRs** that are contained in the set of of **CPRs** of students who take **SIDD**.

```
select name
from Student
where cpr in
    (select cpr
     from Takes
     where course = "SIDD")
```



## That's Not New ...

Find the **names** for the **CPRs** that are contained in the set of of **CPRs** of students who take **SIDD**.

```
select name
from Student, Takes
where Student.cpr = Takes.cpr
and course = "SIDD"
```

.... But Convenient

# Nested Subqueries

Find the student with the highest CPR number.

Option #1

# Nested Subqueries

Find the student with the highest CPR number.

```
select name
from Student
where cpr "is greater or equal to any cpr in"
      ( select cpr
        from Takes
      )
```

# Nested Subqueries

Find the student with the highest CPR number.

```
select name
from Student
where cpr >= all
      ( select cpr
        from Takes
      )
```

# Nested Subqueries

Find the student with the highest CPR number without nested subqueries.

Option #2

# \$100 question

Find the student with the highest CPR number without nested queries!

*Idea: Find all but the highest CPR!*

```
select S.cpr, S.name, S.address  
from Student as S, Student as T  
where S.cpr < T.cpr
```

# \$100 question

Find the student with the highest CPR number without nested queries!

*Idea: Take all students **except** those from before*

```
(select * from Student)
```

**except**

```
(select S.cpr, S.name, S.address  
  from Student as S, Student as T  
 Where S.cpr < T.cpr)
```

# Nested Subqueries

Find the student with the highest CPR number using aggregates.

Option #3



# Nested Subqueries

Find the student with the highest CPR number.

```
select name
from Student
where cpr is equal to
      ( select max(cpr)
        from Takes
      )
```

# Nested Subqueries

Find the student with the highest CPR number.

```
select name  
from Student  
where cpr in
```

```
( select max(cpr)  
  from Takes  
)
```

Because of what we  
already discussed!

# Nested Subqueries

Find the student with the highest CPR number.

```
select name  
from Student  
where cpr =
```

Because this relation  
contains only one  
element



```
( select max(cpr)  
  from Takes  
)
```

# Exercise

Find the **CPR** number of the student with the **highest average grade**.

```
select cpr, avg(grade)
from Takes
group by cpr
having avg(grade) is the highest ever
```

# Exercise

Find the **CPR** number of the student with the **highest average grade**

```
select cpr, avg(grade)
from Takes
group by cpr
having avg(grade) >= all
    (select avg(grade)
     from Takes
     group by cpr)
```

**Relation of all  
grade averages**

# Nesting Operators

- $A < \text{all } (R)$      $A$  is less than every tuple in  $R$
- $A <> \text{all } (R)$      $A$  not in  $R$
- $A > \text{some } (R)$      $A$  is greater than a tuple in  $R$
- $A = \text{some } (R)$      $A$  is in  $R$
- $\text{Exist } (R)$         checks if  $R$  is inhabited

# Our Database for Today

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12

## Class

course	name	etcs
SIDD	Databases	7.5
CSYS	Critical Systems	7.5

# Exists

Find all courses nobody is enrolled in.

```
select course
from class
where not exists
(select *
 from Takes
 where Takes.course = Class.course)
```



# Derived relations

Find the **CPR** number of the student with the **highest average grade**

```
select cpr, avg(grade)
from Takes
group by cpr
having avg(grade) >= all
```

```
(select avg(grade)
from Takes
group by cpr)
```

Both  
relations  
are  
pretty  
similar

# Derived relations

Find the **CPR** number of the student with the **highest average grade** assuming we have pre computed an auxiliary table Aux(cpr,average)

```
select cpr, average
From Aux
where average =
    (select max(average)
     from Aux)
```

Aux

cpr	average
140298-1234	10
041297-5367	11
140789-7612	7
010292-3333	4

# Derived relations

How to precompute Aux?

```
select cpr, avg(grade) as average  
from Takes  
group by cpr
```

Aux

cpr	average
140298-1234	10
041297-5367	11
140789-7612	7
010292-3333	4

# Putting all together

Find the CPR number of the student with the highest average grade

```
select cpr, average
from (select cpr, avg(grade) as average
      from Takes
      group by cpr) as Aux
where average =
      (select max(average)
       from Aux)
```

# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregates
- Nested subqueries
- Insertion
- Joins

# Insertions

```
insert into Student  
values ("010192-1329", "Lone", "Lyngby")
```

```
insert into Student(ssn, name, address)  
values ("031193-1339", "Jens", "Odense")
```

```
insert into Student  
select ssn, name, address  
from Foreign-student
```

# Deletions

```
delete from Student  
where name="Lone"
```

Deletes all records where the name is "Lone"!

```
delete from Student
```

Deletes all records where the name is "Lone"!

# Updates

Jesper got a 12 in SIDD.

```
update Takes
set      grade=12
where    name="Jesper"
and      course="SIDD"
```



# Data Manipulation Language DML

- Select – From – Where
- Renaming
- Set operation
- Ordering
- Aggregates
- Nested subqueries
- Insertion
- Joins

# Natural Join

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12
280296-2222	CSYS	8

## Student ⋈ Takes

cpr	name	address	course	grade
140298-1234	Jesper	Copenhagen	SIDD	10
041297-5367	Nikoline	Aarhus	SIDD	12

# Natural Join

Verbose, but ok:

```
select cpr, name, address,  
       course, grade  
from Student Join Takes  
     on Student.cpr = Takes.cpr
```

# Natural Join

Verbose, but ok:

```
select cpr, name, address,  
       course, grade  
from Student Join Takes  
     on Student.cpr = Takes.cpr
```

Also called: Inner Join

# However, observe:

## Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

## Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12
280296-2222	CSYS	8

Student ⋈ Takes

cpr	name	address	course	grade
140298-1234	Jesper	Copenhagen	SIDD	10
041297-5367	Nikoline	Aarhus	SIDD	12

Lost  
Information

# Ways to Safe Information

- Outer join
- Left outer join

# Outer Joins

Student

cpr	name	address
140298-1234	Jesper	Copenhagen
041297-5367	Nikoline	Aarhus
151197-2352	Claus	Dragør
050596-1142	Martin	Copenhagen

Takes

cpr	course	grade
140298-1234	SIDD	10
041297-5367	SIDD	12
280296-2222	CSYS	8

Student ⋈ Takes

cpr	name	address	course	grade
140298-1234	Jesper	Copenhagen	SIDD	10
041297-5367	Nikoline	Aarhus	SIDD	12
151197-2352	Claus	Dragør	NULL	NULL
050596-1142	Martin	Copenhagen	NULL	NULL
280296-2222	NULL	NULL	CSYS	8

Right  
Outer  
Join

# Be Aware of Null Values

Null values are semantically ambiguous!

$cond_1$	$cond_2$	$cond_1$ AND $cond_2$	$cond_1$ OR $cond_2$
true	true	true	true
true	false	false	true
true	unknown	unknown	true
false	true	false	true
false	false	false	false
false	unknown	false	unknown
unknown	true	unknown	true
unknown	false	false	unknown
unknown	unknown	unknown	unknown

$cond$	NOT $cond$
true	false
false	true
unknown	unknown



# Conclusion

Complete guide to SQL as a  
Data Manipulation Language (DML)

Next time

SQL as a Data Definition Language (DDL)