

Transactions

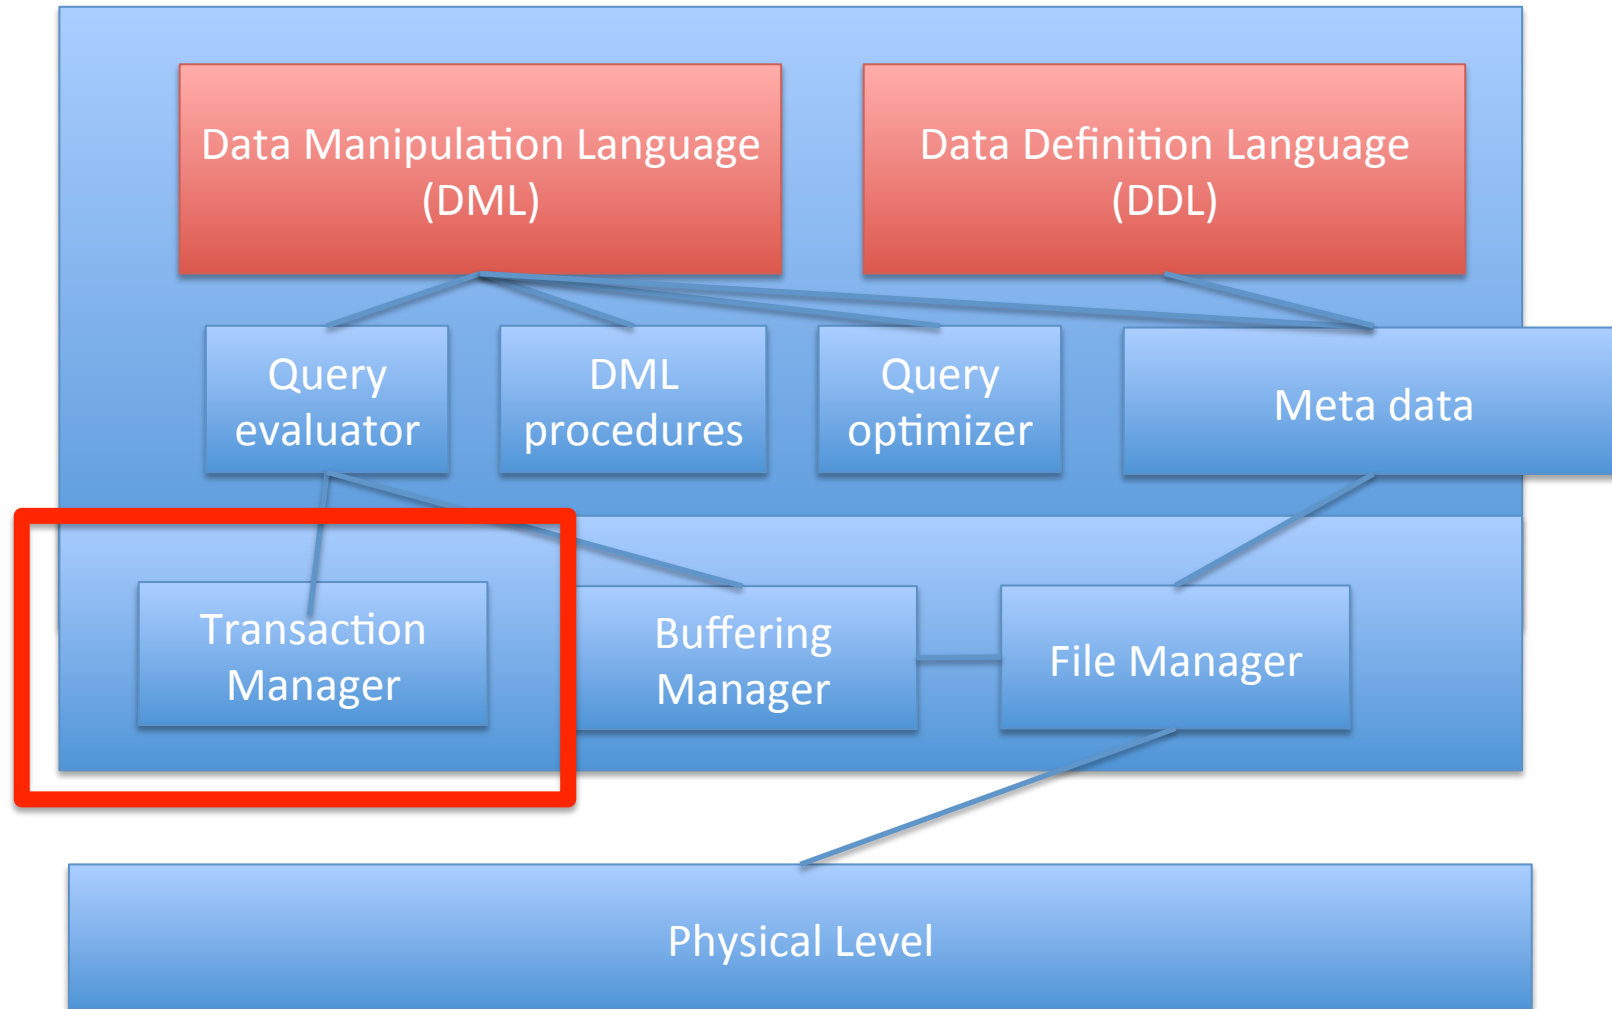
Carsten Schürmann

Motivation

- Two update the same record ("Jesper")
 - ("Jesper", 10), ("Jesper", 20)
 - Race condition!
- Transfer 100 k from one account to another
 - how does it happen?

DBMSs automatically handle both issues

Database System Architecture



Transactions


- Execution of one or more operations
- Implements a higher-level function
- May temporarily break consistency of database
- Transactions form unit
 - Either all steps are executed
 - Or no steps are executed

Transaction Example


Move 1000 kr from Jesper's account to mine!

1. Check whether Jesper's got 1000 kr
2. Deduct 1000 kr from Jesper's account
3. Deposit 1000 kr on Carsten's account

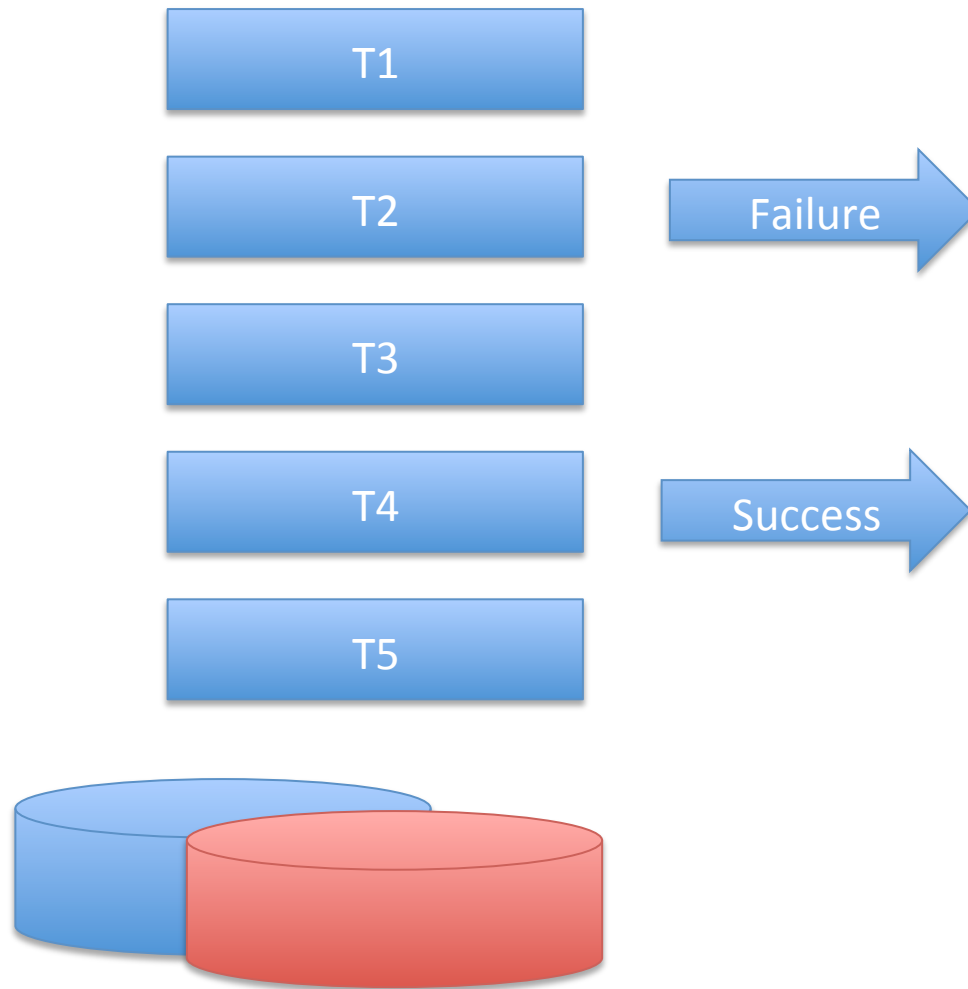
Perfect System



Perfect System



Perfect System



Thoughts

- Sequentiality is good
- Concurrency is better
- Parallel execution of independent transactions
- Interleaving of transactions
- Correctness
- Fairness
- Persistency

ACID

Atomicity: All actions in the defining the transaction happen, or none happen.

Consistency: If each transaction is consistent and the DB starts consistent, then it ends up consistent.

Isolation: Execution of one transaction is isolated from that of other transactions

Durability: If a transaction finishes successfully, its effects persist.

ACID

Atomicity: *all or nothing*

Consistency: *it looks correct to me*

Isolation: *as if alone*

Durability: *survive failures*

Basic Database Operations

Database: Fixed set of names (A, B, C...)

- Read: $R(A)$
- Write: $W(A)$

Transaction commands

1. begin (start the transaction)
2. commit (make state persistent)
3. abort (roll state back to the beginning)

Atomicity

Assume first two step executes

Then power failure

Rewind the transactions

1. Check whether Jesper's got 1000 kr
2. Deduct 1000 kr from Jesper's account
3. Deposit 1000 kr on Carsten's account

Recreate the database as it was before start

Mechanisms: Logging!

(Audit trail and efficiency)

ACID

Atomicity: All actions in the defining the transaction happen, or none happen.

Consistency: If each transaction is consistent and the DB starts consistent, then it ends up consistent.

Isolation: Execution of one transaction is isolated from that of other transactions

Durability: If a transaction finishes successfully, its effects persist.

Consistency

Data Consistency

Data in the DBMS is accurate in modeling the real world and follows *integrity constraints*!

Transaction Consistency

If the database is consistent before the transaction starts, it will be after also.

Level of Consistency

Strong Consistency

Guaranteed to see all writes immediately, but transactions will be slower

Weak Consistency

Writes will take time to become visible, but transactions will be faster

ACID

Atomicity: All actions in the defining the transaction happen, or none happen.

Consistency: If each transaction is consistent and the DB starts consistent, then it ends up consistent.

Isolation: Execution of one transaction is isolated from that of other transactions

Durability: If a transaction finishes successfully, its effects persist.

Isolation

A transaction runs as if it was running by itself

Pessimistic Isolation Strategy

Don't let problems arise in the first place

Optimistic Isolation Strategy

Since conflicts are rare, deal with them as they arise


Invariant

$$A + B = 2120$$

There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together.

The net effect must be equivalent to these two transactions running **serially** in some order.



T1	T2
BEGIN	BEGIN
$A = A + 100$	$A = A * 1.06$
$B = B - 100$	$B = B * 1.06$
COMMIT	COMMIT



Scheduling


T1
BEGIN
$A = A + 100$
$B = B - 100$
COMMIT

T2
BEGIN
$A = A * 1.06$
$B = B * 1.06$
COMMIT



T2
BEGIN
$A = A * 1.06$
$B = B * 1.06$
COMMIT

T1
BEGIN
$A = A + 100$
$B = B - 100$
COMMIT



Schedule

T1	T2
BEGIN	
A = A + 100	
B = B - 100	
COMMIT	
	BEGIN
	A = A * 1.06
	B = B * 1.06
	COMMIT

Schedule

T1	T2
	BEGIN
	$A = A * 1.06$
	$B = B * 1.06$
	COMMIT
BEGIN	
$A = A + 100$	
$B = B - 100$	
COMMIT	

Interleaving Schedule (Good)

T1	T2
BEGIN	
A = A + 100	
	BEGIN
	A = A * 1.06
B = B - 100	
COMMIT	
	B = B * 1.06
	COMMIT


=

T1	T2
BEGIN	
A = A + 100	
B = B - 100	
COMMIT	
	BEGIN
	A = A * 1.06
	B = B * 1.06
	COMMIT

Interleaving Schedule (Bad)

T1	T2
BEGIN	
A = A + 100	
	BEGIN
	A = A * 1.06
	B = B * 1.06
	COMMIT
B = B - 100	
COMMIT	

Someone lost 6kr.



Correctness

Serial schedule

A schedule that does not interleave the actions of different transactions

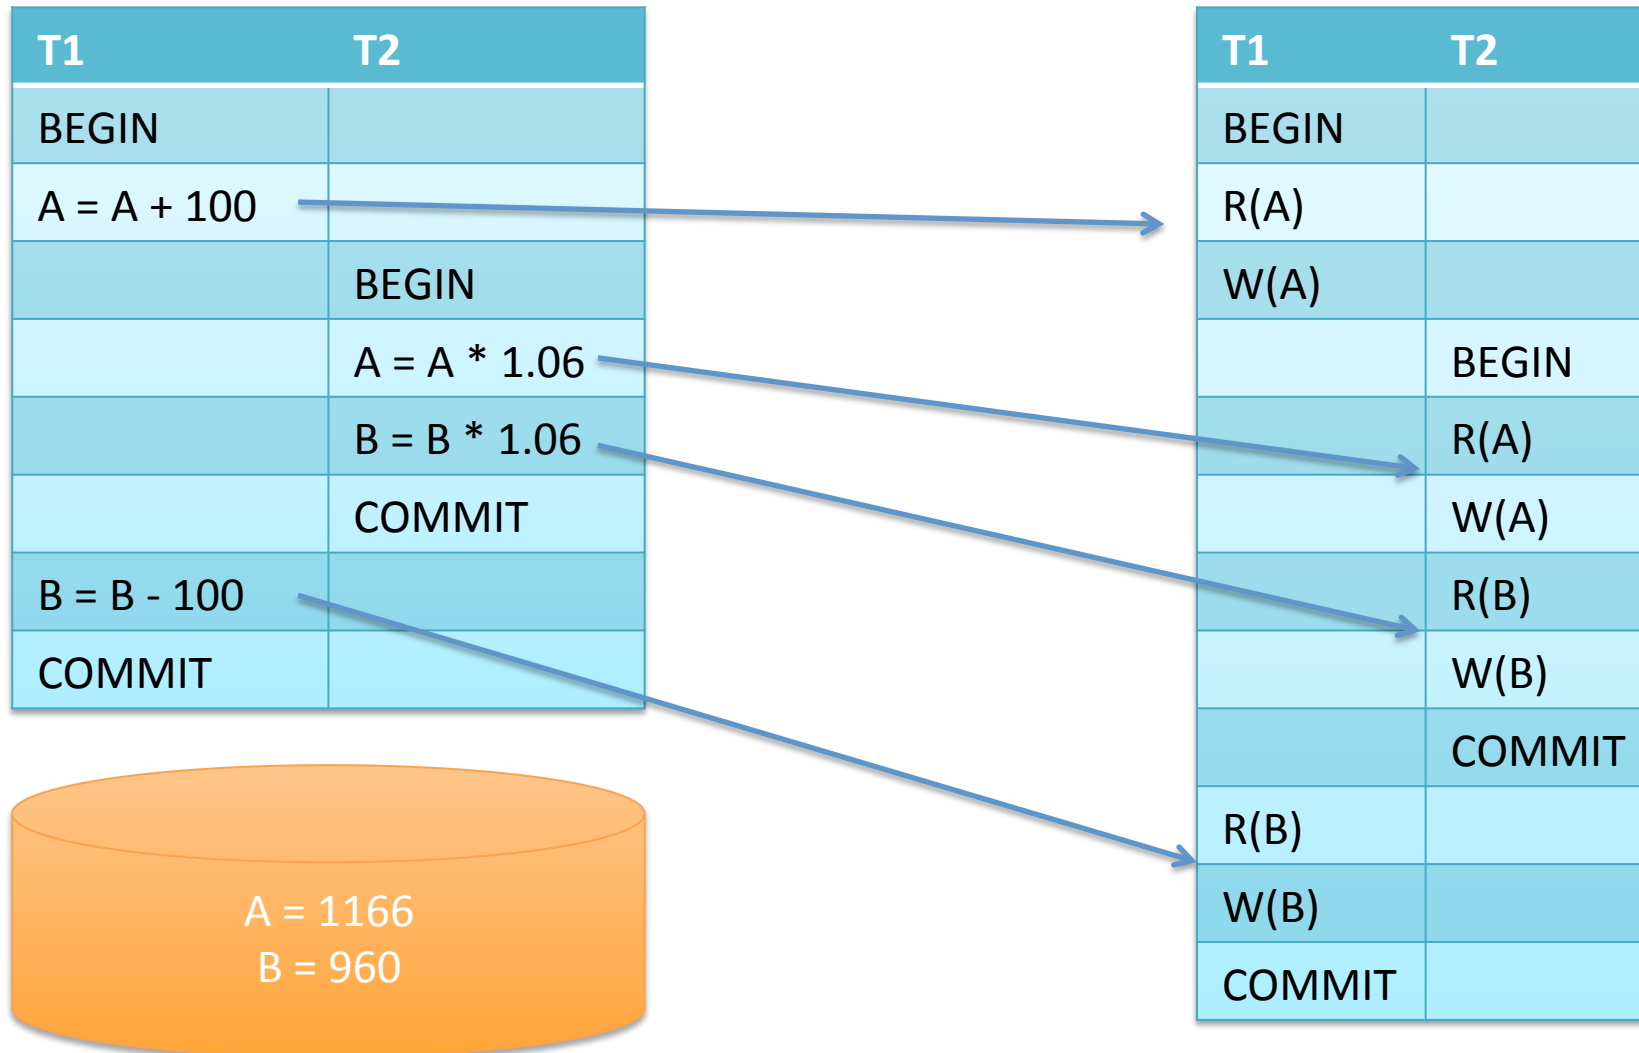
Equivalent schedule $S1 \equiv S2$

For any database state, the effect of executing the schedule **S1** is identical to the effect of executing the schedule **S2**

Serializability

A schedule is *serializable* if and only if it is equivalent to a serial schedule.

Interleaving Schedule (Bad)



Interleaved Execution Anomalies

- **Read-Write** conflicts (R-W)
- **Write-Read** conflicts (W-R)
- **Write-Write** conflicts (W-W)
- Note, no **Read-Read** conflicts

Write Read Conflicts

Reading Uncommitted Data
“Dirty Reads”

T1	T2
BEGIN	
R(A)	
W(A)	
	BEGIN
	R(A)
	W(A)
	COMMIT
R(B)	
W(B)	
ABORT	

Read Write Conflicts

“Unrepeatable reads”

	T1	T2	
	BEGIN		
50kr	R(A)		
		BEGIN	
		R(A)	50kr
		W(A)	100kr
		COMMIT	
100kr	R(A)		
	COMMIT		

Write Write Conflicts


Overwriting uncommitted data

	T1	T2	
	BEGIN		
50kr	W(A)		
		BEGIN	
		W(A)	590kr
		W(B)	Jesper
		COMMIT	
Mette	W(B)		
	COMMIT		

Locks

Locks

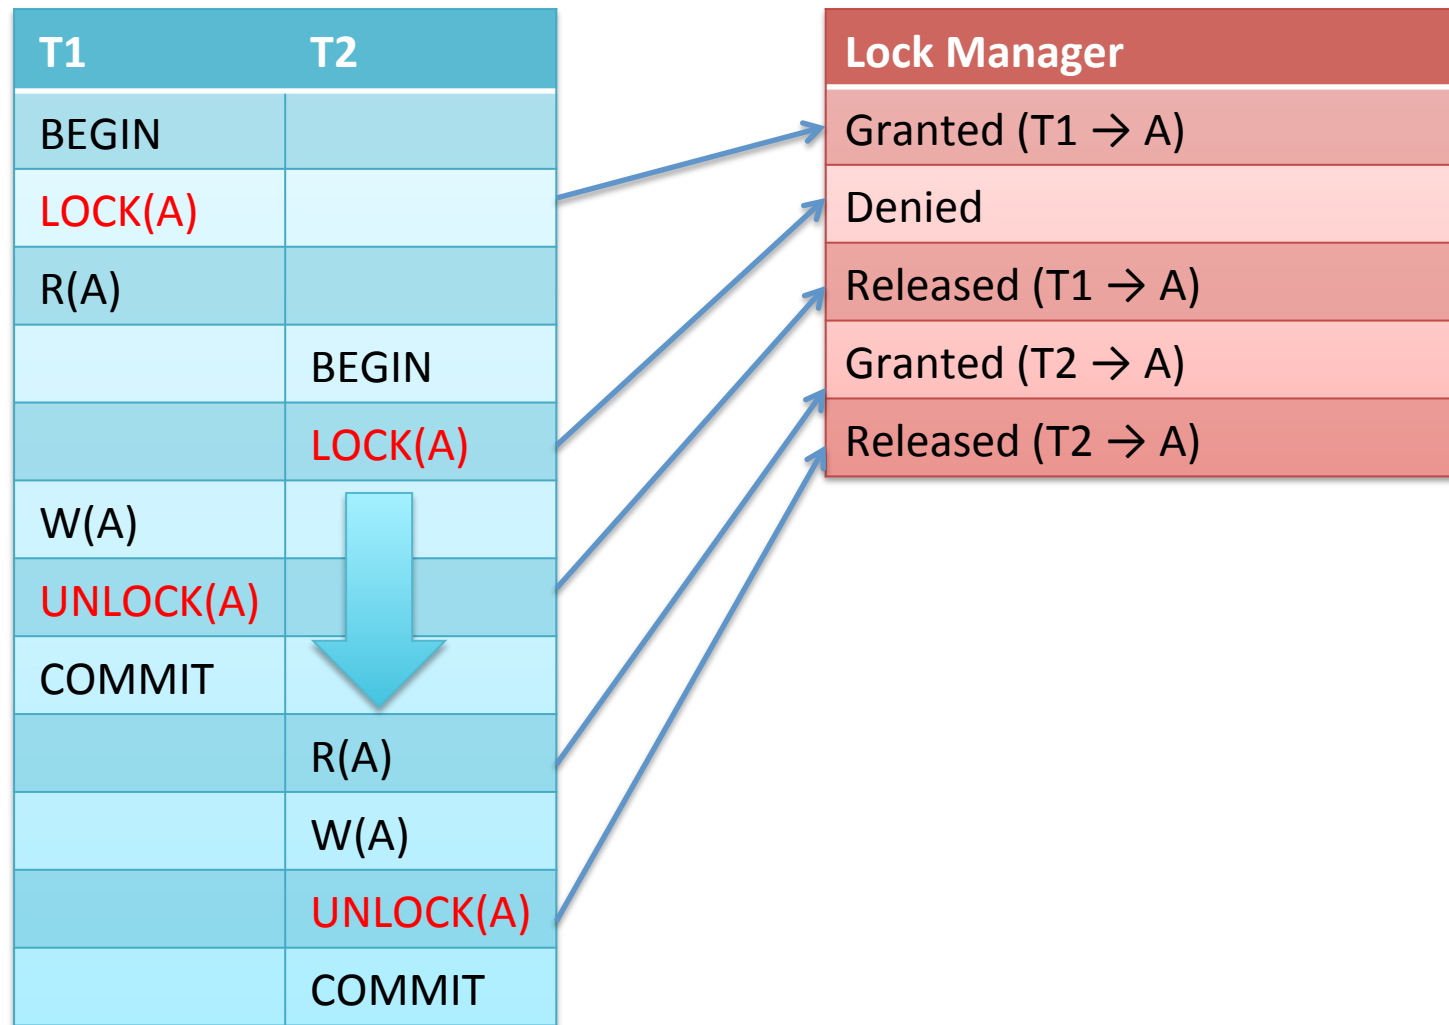
- Protects a resource
 - Row in a table
 - Table in a database
- Operations
 - Lock
 - Unlock



Execution without Locks

T1	T2
BEGIN	
R(A)	
	BEGIN
	R(A)
W(B)	
COMMIT	
	W(B)
	COMMIT

Execution with Locks



Types of Locks

- S-Lock Shared lock (reads)
- X-Lock Exclusive lock (writes)

Compatibility matrix

	Lock requested		
Lock granted		S	X
	S	Yes	No
	X	No	No

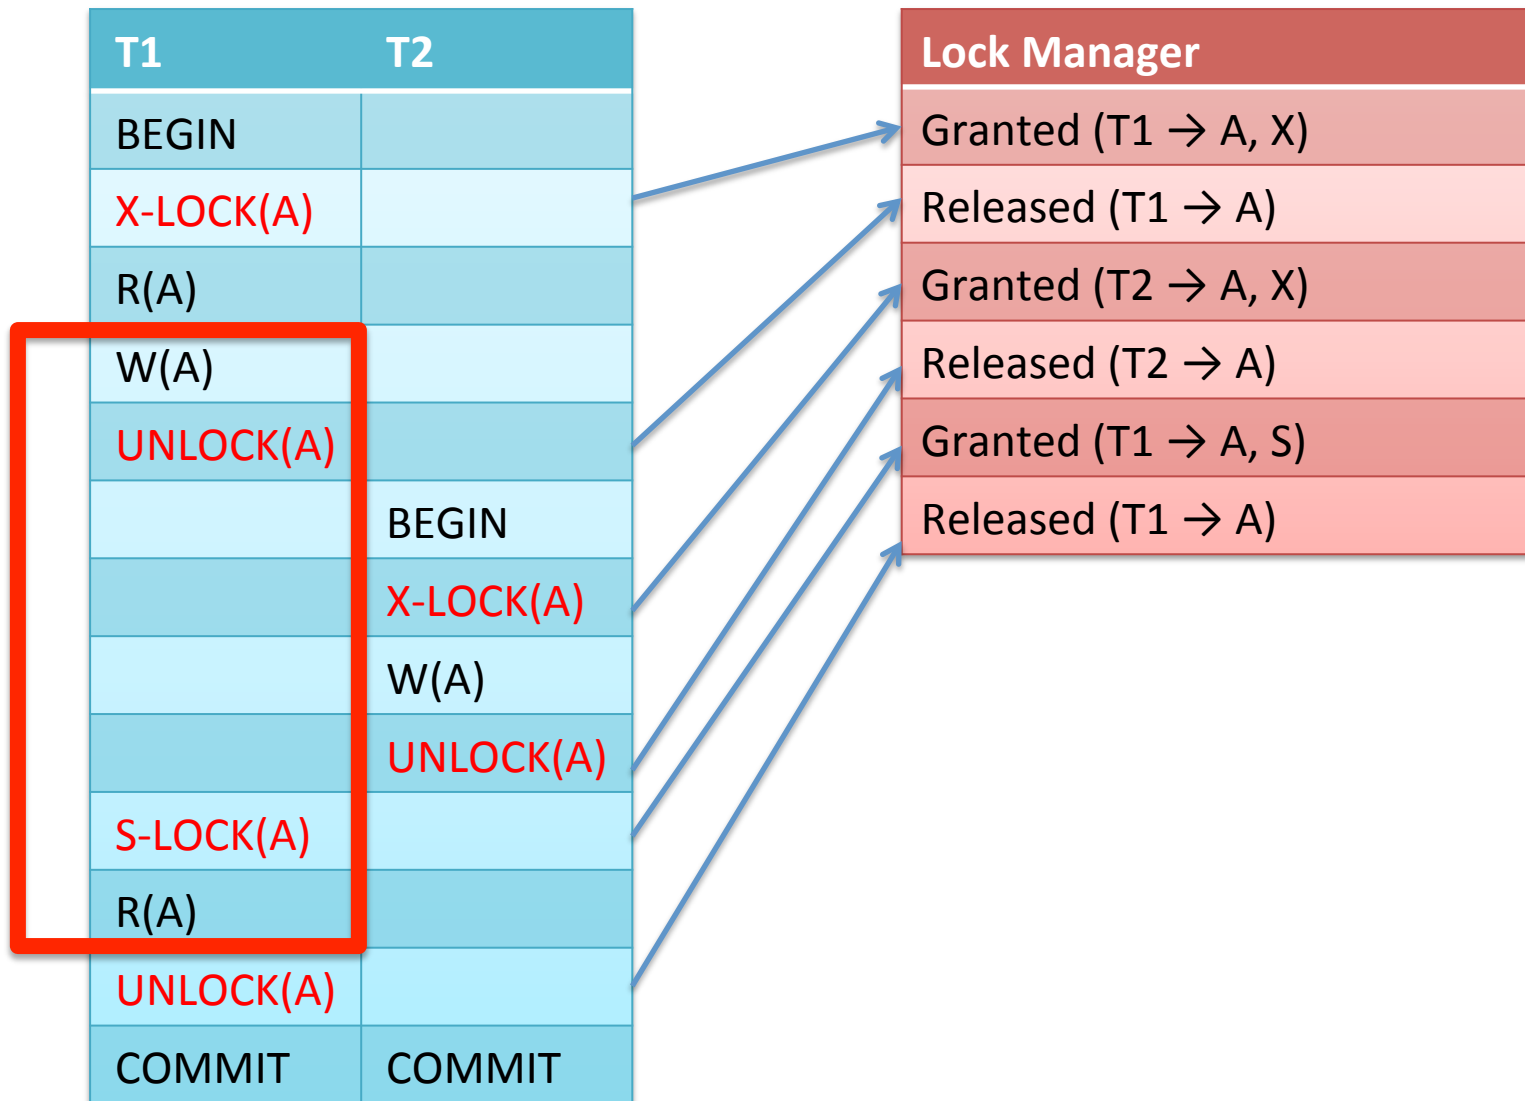
Execution with Locks

- Transactions request locks (or upgrades)
- Lock manager grants or blocks requests
- Transactions release locks
- Lock manager updates lock-table

But

Execution with Locks

Problem
with
Isolation!



Concurrency Control

Locks can help but they need to be requested and released sensibly!

Two Phase Locking (2PL)


Phase 1: Growing

- Transaction requests the locks that it needs from the DBMS's lock manager.
- The lock manager grants/denies lock requests

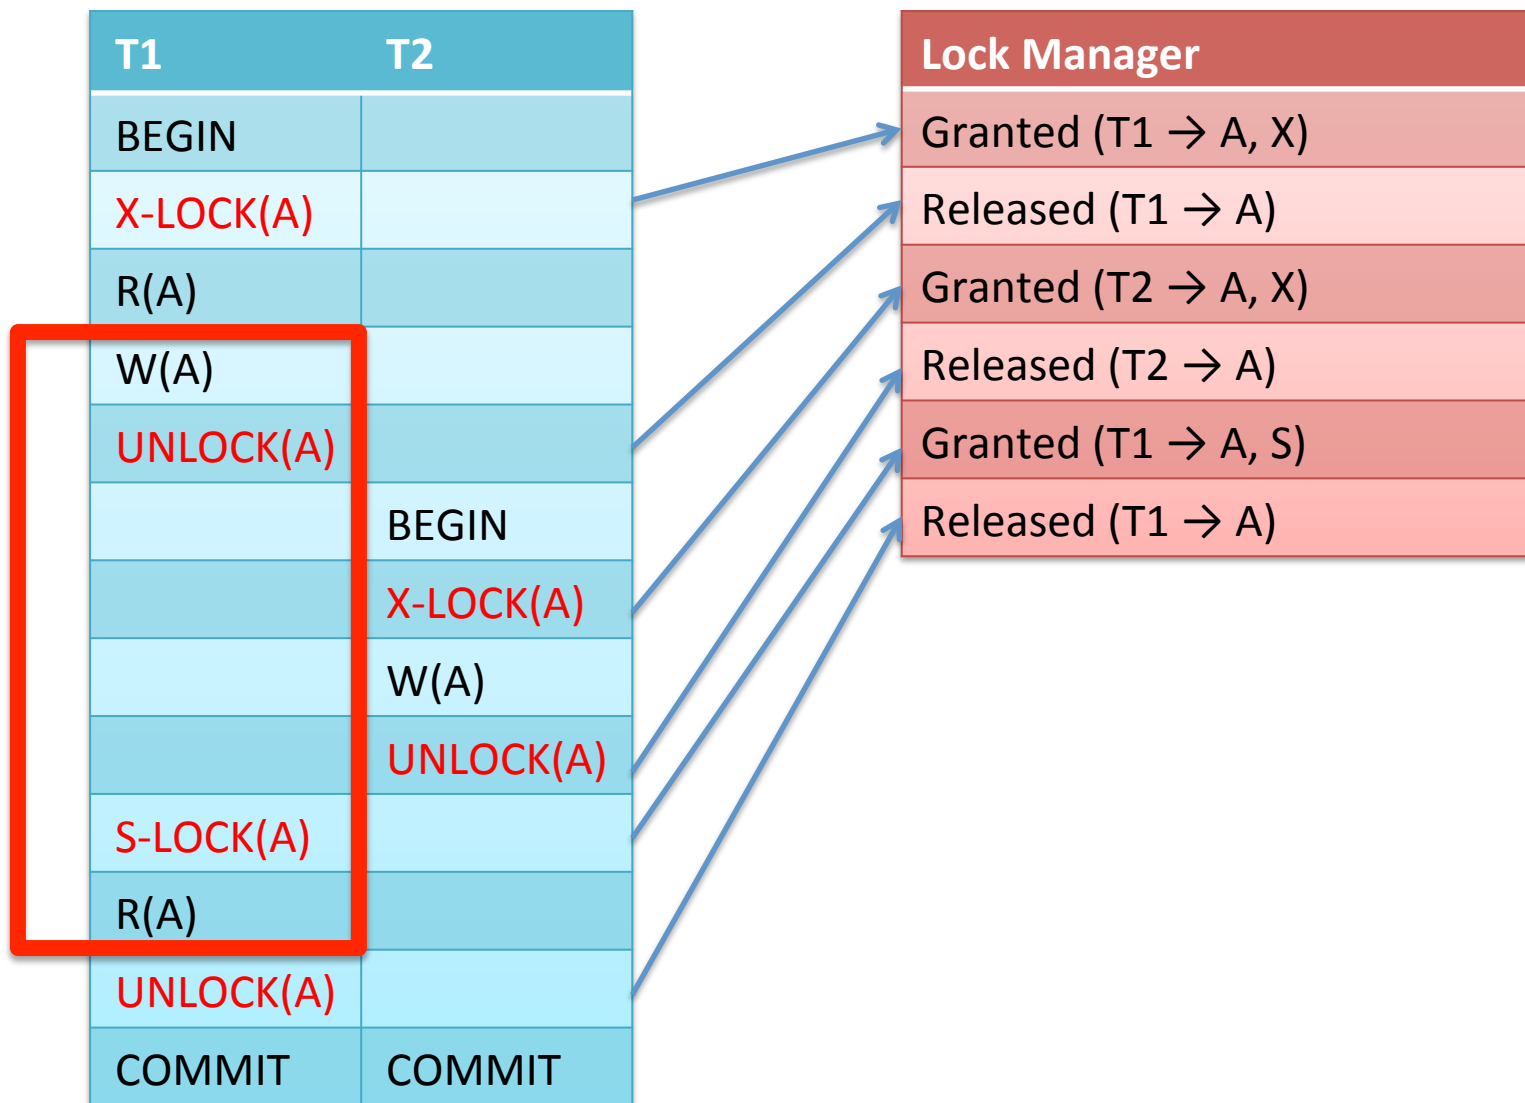
Phase 2: Shrinking

- The transaction is allowed to only release locks that it previously acquired. It cannot acquire new locks.

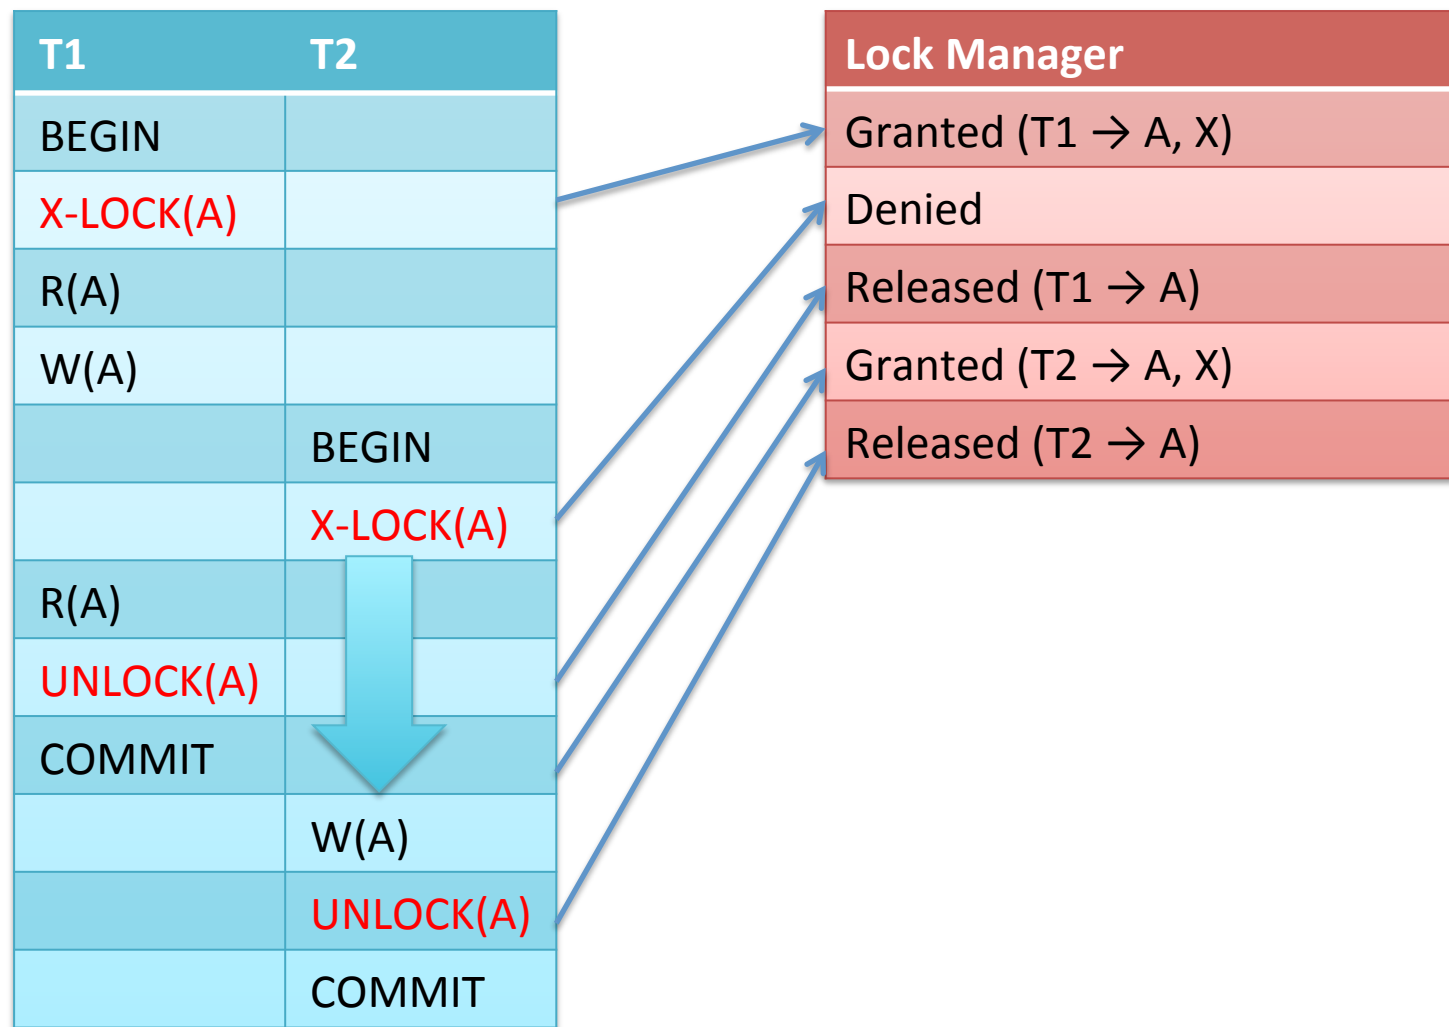
Transaction Lifetime



Violates Two Phase Locking



Violates Two Phase Locking



Two Phase Locking

- There are schedules that are serializable but would not be allowed by 2PL
- Locking limits concurrency
- May lead to deadlocks
- May still have “dirty reads”
- Solution: **Strict 2PL**


Strict Two Phase Locking

A schedule is *strict* if a value written by a transaction is not read or overwritten by other transaction until that transaction finishes.

Advantages:

- Recoverable
- Do not require cascading aborts
- Aborted transactions can be roll-backed by just restoring original values of modified tuples.

Transaction Lifetime



Strict Two Phase Locking

Transactions hold all of their locks until commit.

Good:

- Avoids “dirty reads” etc

Bad:

- Limits concurrency even more
- And still may lead to deadlocks

Locking in Practice

- No need to set locks manually
- But DBMS may require hints to help improve concurrency

ACID

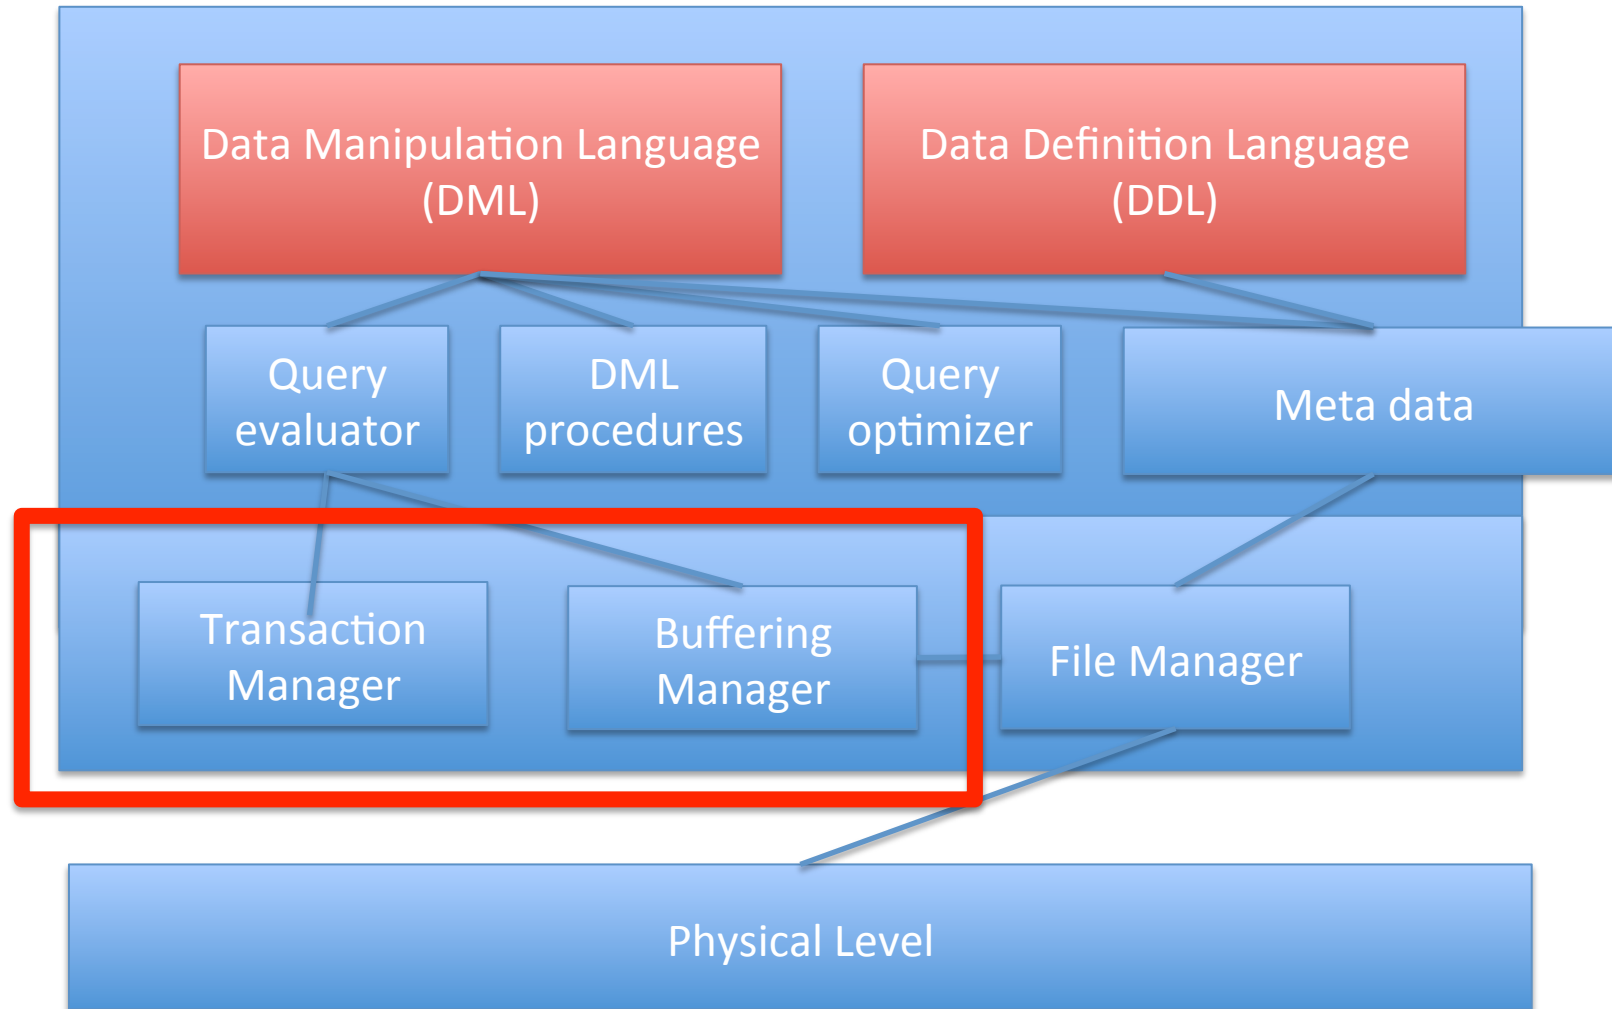
Atomicity: All actions in the defining the transaction happen, or none happen.

Consistency: If each transaction is consistent and the DB starts consistent, then it ends up consistent.

Isolation: Execution of one transaction is isolated from that of other transactions


Durability: If a transaction finishes successfully, its effects persist.

Database System Architecture



T1	T1
BEGIN	BEGIN
A = A	R(A)
B = B	W(A)
COM	R(B)
	W(B)
	COMMIT

Buffer



Transaction Durability

- Records **stored** on disk
- Buffer manager
 - Pages are copied into memory and back to disk at the discretion of the buffer manager
 - One could force flush

This is too slow!

Write-Ahead Log

Log changes before the database is updated

Assume that the log lives on stable storage

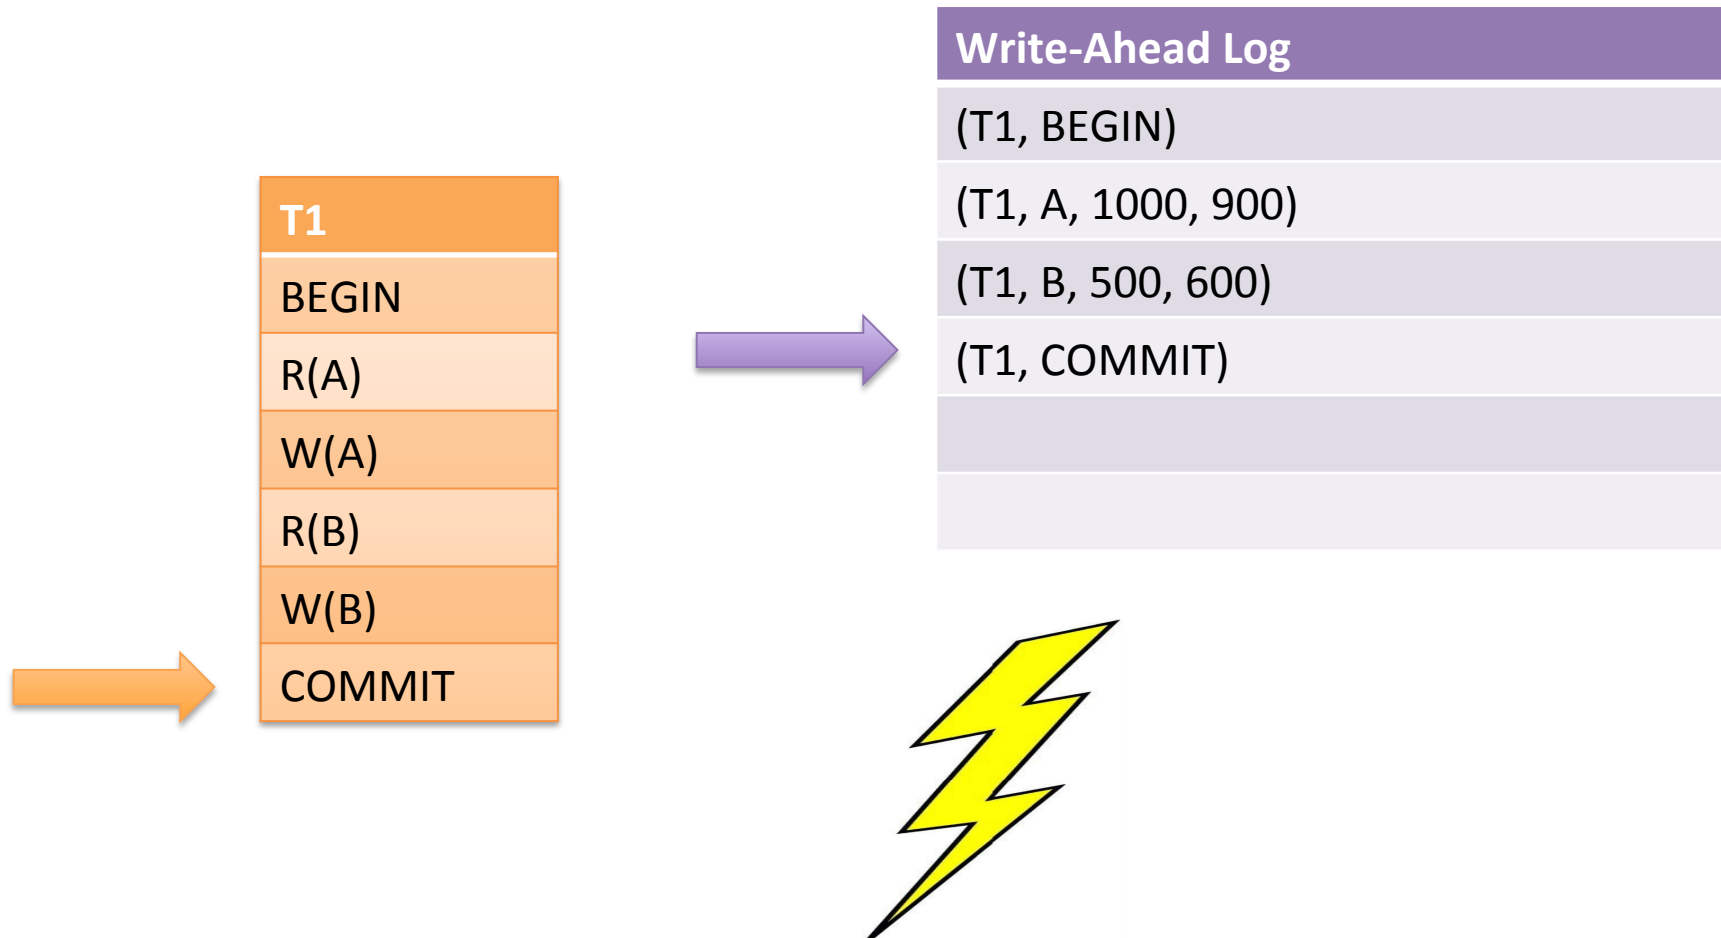
Transaction finishes

- Write a commit to the log
- Make sure that all log messages are flushed before acknowledgment

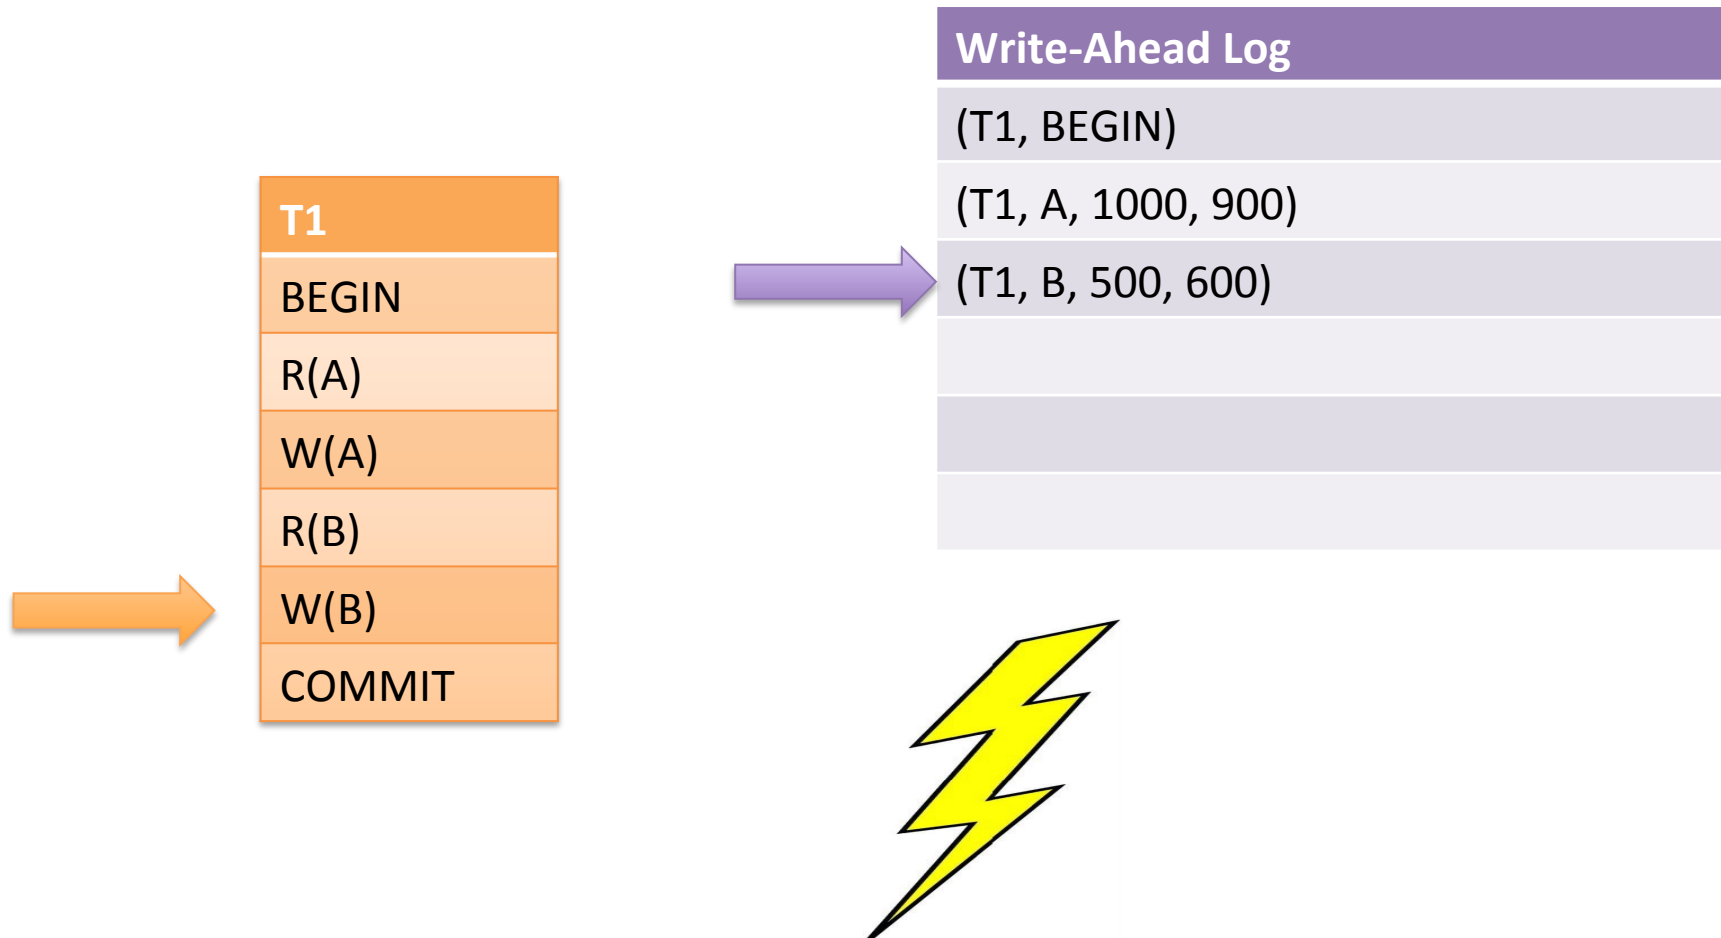
Transaction fails

- Undo uncommitted transactions
- Redo committed transactions

CRASH and REDO



CRASH and UNDO



Summary

- Concurrency control and recovery are among the most important functions provided by a DBMS.
- Concurrency control is automatic
 - System automatically inserts lock/unlock requests and schedules actions of different transactions.
 - Ensures that resulting execution is equivalent to executing the transaction one after the other in some order.

Transactions in Applications

JDBC connection

```
String url = "jdbc:mysql://localhost/";
String dbName = "imdb";
String driver = "com.mysql.jdbc.Driver";
String userName = "root";
String password = "";

try {
    Class.forName(driver);
    Connection conn = DriverManager.getConnection(url+dbName,userName,password);
    System.out.println("Connected to MySQL");

    Database operations

    conn.close();
    System.out.println("Disconnected from MySQL");
} catch (Exception e) {
    e.printStackTrace();
}
```

JDBC dynamic SQL

```
try {  
    Statement st = conn.createStatement();  
    ResultSet rs = st.executeQuery("SELECT gender, count(*) FROM person GROUP BY gender");  
    while (rs.next()) {  
        System.out.println(rs.getString("gender")+": "+rs.getInt(2));  
    }  
  
    st.executeUpdate("DROP TABLE IF EXISTS JDBCtest");  
    st.executeUpdate("CREATE TABLE JDBCtest(id int, string varchar(10))");  
    st.executeUpdate("INSERT INTO JDBCtest VALUES (1, \"Tada!\")");  
}  
catch(SQLException s){  
    System.out.println(s.toString());  
}
```

Use caution when creating SQL based on user input!

JDBC static SQL

```
PreparedStatement insertPerson =  
conn.prepareStatement("INSERT INTO person VALUES (?, ?, ?, ?, ?, ?)"); // Create prepare  
insertPerson.setInt(1, 123456);  
insertPerson.setString(2, "John Doe");  
insertPerson.setString(3, "M");  
insertPerson.setDate(4, new java.sql.Date(1606176000000)); // Set date, given in mil  
insertPerson.setNull(6, java.sql.Types.INTEGER); // Set to NULL  
insertPerson.executeUpdate(); // Execute prepared statement with current parameters
```

Efficiency issues

Connection takes time to establish

- reuse for multiple operations

It takes time to parse dynamic SQL

- prepared statements execute faster

Understand the DBMSd

- ORDER BY may force creation of full result within the DBMS before any output reaches the application.

Cursors

Cursor allows the result to be traversed.

JDBC examples

- `Statement s = con.createStatement
 (ResultSet.TYPE_FORWARD_ONLY,
 ResultSet.CONCUR_READ_ONLY)`
- `Statement s = con.createStatement

 (ResultSet.TYPE_SCROLL_INSENSITIVE,
 ResultSet.CONCUR_UPDATABLE)`

Cursors

TYPE_FORWARD_ONLY (CONCUR_READ_ONLY)

Single (forward) pass through the result set.

TYPE_FORWARD_ONLY (CONCUR_UPDATABLE)

Single (forward) pass through the result set, to update rows.

TYPE_SCROLL_INSENSITIVE

Cursor moves forward and backward. Result set is **not** updated concurrently.

TYPE_SCROLL_SENSITIVE

Cursor moves forward and backward. Result set is updated concurrently.

Four examples

1. Movies by year – imperative way
2. Movies by year – SQL centric way
3. Iterating through a large result set
4. Iterating through a filtered result set

Transactions in JDBC

```
conn.setAutoCommit(false);  
// Disable automatic commit  
conn.commit();  
// Commit all pending updates  
conn.setSavepoint();  
    // Something to roll back to  
conn.rollback();  
// Abort all pending updates
```

Isolation level syntax

Begin transaction with:

```
SET TRANSACTION ISOLATION LEVEL
{ READ UNCOMMITTED |
  READ COMMITTED |
  REPEATABLE READ |
  SERIALIZABLE }
```

SQL isolation levels

READ UNCOMMITTED

No locks are obtained

READ COMMITTED

Read locks are immediately released -
read values may change during the
transaction

REPEATABLE READ

2PL but no lock when adding new tuples

SERIALIZABLE:

2PL with lock when adding new tuples

Language integration

Little languages, database integration

E.g. **Ruby on Rails**

New query sublanguages for mainstream languages such as C#.

E.g. **LINQ**

Automatic translation to SQL.