

Outline

1

- Graphical models (continued from last time)
 - Belief propagation
 - Undirected graphs and factor graphs
- Markov models
- Hidden Markov models

CHAPTER 14:

GRAPHICAL MODELS – PART 2

Last time

3

- Joint distributions over random variables
- general expression:


$$p(X_1, X_2, \dots, X_d) = p(X_1) p(X_2|X_1) p(X_3|X_1, X_2) \cdots p(X_d|X_1, X_2, \dots, X_{d-1})$$

... complex calculations

Last time

4

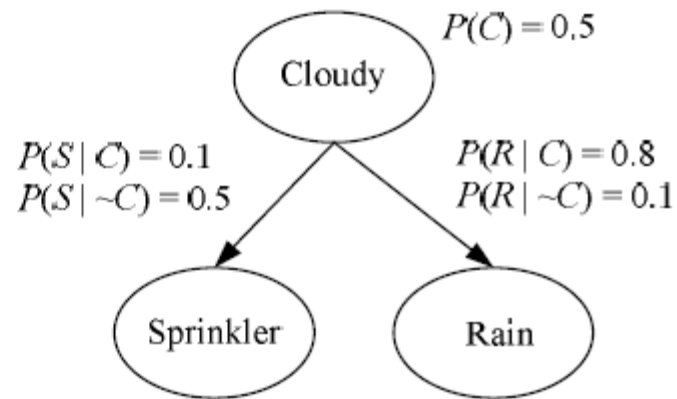

- Graphical models:
 - depict joint distributions
 - conditional independence assumptions
 - simplify distributions



Last time

5


- Probability calculations:
 - Various examples using Bayes' rule



Belief Propagation

6


- We want a *general algorithm* for finding $P(\text{some variable} \mid \text{some evidence})$ in a graphical model
- No designated input and output nodes required



Belief Propagation

7


- We want a *general algorithm* for finding $P(\text{some variable} \mid \text{some evidence})$ in a graphical model
- No designated input and output nodes required
- Make use of the structure of the graph
- Different subgraphs imply different calculations



Belief Propagation

8

- Simplest case: Chain




- We want to find $P(X|E)$
 - X : node of interest
 - E : evidence nodes

Belief Propagation

9


- Simplest case: Chain



- Every (non-leaf or -root) node splits chain into ancestors and descendants
- Every evidence node blocks any evidence further up/down

Belief Propagation


10



- Evidence can be thought of as a message being sent up/down the chain
- Nodes:
 - Receive msg from parent/child
 - Processes the msg
 - Sends it onwards to child/parent

Belief Propagation

11




□ Every node stores:

- evidence via parent: $\pi(X) \equiv P(X|E^+)$
- evidence via child: $\lambda(X) \equiv P(E^-|X)$

$$P(X|E) = \alpha \pi(X) \lambda(X)$$

Belief Propagation

12



- Evidence is propagated recursively


$$\pi(X) = \sum_U P(X|U) \pi(U)$$

$$\lambda(X) = \sum_Y P(Y|X) \lambda(Y)$$

Trees

13


- One parent, multiple children
- Same logic as for chains
 - different λ from each child
 - Different π to each child



Undirected Graphs: Markov Random Fields

14


- So far we looked at directed models
- For symmetric interactions, undirected models can be more convenient.
- E.g. a model of pixels in an image



Undirected Graphs: Markov Random Fields

15



- A and B are conditionally independent if removing C makes them unconnected.



From directed to undirected

16


- Drop directions of edges
- *Marry* the parents of nodes (*moralization*)



Undirected Graphs: Markov Random Fields

17

- Instead of parent/child:
 - Cliques
 - Maximal cliques



Undirected Graphs: Markov Random Fields

18



- Instead of conditional probabilities:
- Potential function $\psi_c(X_c)$
 - Function of variables X in the clique C
- The joint prob. is defined in terms of the max clique potentials

$$p(X) = \frac{1}{Z} \prod_c \psi_c(X_c)$$

Factor Graphs

19

- Define new factor nodes and write the joint in terms of them




$$p(X) = \frac{1}{Z} \prod_s f_s(X_s)$$

CHAPTER 15:

MARKOV MODELS


Billy the Budgie

21



Billy the Budgie


22



Billy the Budgie

23



- Observations over time:
PPSSPPPSSPPPSPPPP.....



Billy the Budgie

24



- Observations over time:
PPSSPPPSSPPPSPPPP.....



Billy the Budgie

25

- Observations over time:
PPSSPPPSSPPPSPPPP.....



		To	
		Swing	Perch
From	Swing	0.5	0.5
	Perch	0.2	0.8

Markov Process

26

- Regularly spaced discrete sequences
 - Time (seconds, minutes, years...)
 - DNA or protein sequences
 - Letters/speech sounds
- Observation at $t+1$ is dependent on previous observation(s)

Markov Process

27

- General case:

$$P(q_{t+1} = S_j | q_t = S_i, q_{t-1} = S_k, \dots)$$

- First order Markov Chain:

$$P(q_{t+1} = S_j | q_t = S_i, q_{t-1} = S_k, \dots) = P(q_{t+1} = S_j | q_t = S_i)$$


- The state at $t+1$ depends only on the state at t

Transition probabilities

28

- Independent of time
- Transition matrix **A**
- Rows sum to 1

$$\mathbf{A} = [a_{ij}]$$



		To	
		Swing	Perch
From	Swing	0.5	0.5
	Perch	0.2	0.8

Initial probabilities

29

- We need a way to model the beginning of the sequence of observations
- Vector of initial probabilities


$$\mathbf{\Pi} = [\pi_i]$$

$$\pi_i \equiv P(q_1 = S_i)$$

$$\sum_{i=1}^N \pi_i = 1$$

Stochastic automaton

30



Note: The circles here are not different variables, they are *states* of the same variable

Observed sequences

31

□ Given:

- Sequence of states: $Q = \{q_1, q_2 \dots q_{T-1}, q_T\}$
- Observed sequence $O = Q$


□ Probability of O is the product of transition probabilities:

$$P(O = Q | \mathbf{A}, \mathbf{\Pi}) = P(q_1) \prod_{t=2}^T P(q_t | q_{t-1}) = \pi_{q_1} a_{q_1 q_2} \cdot \dots \cdot a_{q_{T-1} q_T}$$

Mini-exercise

32

- Calculate the probability of the sequence: PPSSPPP



	Π	To	
		A	
From	0.5	Swing	Perch
	0.5	Swing	Perch
		0.5	0.5
		0.2	0.8

Mini-exercise

33


- Calculate the probability of the sequence:

PPSSPPP

- $\pi_P \cdot a_{PP} \cdot a_{PS} \cdot a_{SS} \cdot a_{SP} \cdot a_{PP} \cdot a_{PP}$

- $0.5 \cdot 0.8 \cdot 0.2 \cdot 0.5 \cdot 0.5 \cdot 0.8 \cdot 0.8$

- $= 0.0128$



π
0.5
0.5


From

A	To	
	Swing	Perch
Swing	0.5	0.5
Perch	0.2	0.8

Mini-exercise II

34

- Given the sequence:
SSSPPSSSPSSSPP
- Estimate the transition probabilities




		To	
		A	
From	A	Swing	Perch
	Swing		
	Perch		

Mini-exercise II

35

- Sequence: SSSPPSSSPSSSP
- $S^* = 9$, $SS = 6/9$, $SP = 3/9$
- $P^* = 4$, $SS = 2/4$, $SP = 2/4$



		To	
		Swing	Perch
From	Counts		
	Swing	6	3
	Perch	2	2

		To	
		Swing	Perch
From	A		
	Swing	2/3	1/3
	Perch	1/2	1/2

Estimating Probabilities

36

□ Initial probabilities

$$\hat{\pi}_i = \frac{\#\{\text{sequences starting with } S_i\}}{\#\{\text{sequences}\}} = \frac{\sum_k 1(q_1^k = S_i)}{K}$$

□ Transition probabilities

$$\hat{a}_{ij} = \frac{\#\{\text{transitions from } S_i \text{ to } S_j\}}{\#\{\text{transitions from } S_i\}} = \frac{\sum_k \sum_{t=1}^{T-1} 1(q_t^k = S_i \text{ and } q_{t+1}^k = S_j)}{\sum_k \sum_{t=1}^{T-1} 1(q_t^k = S_i)}$$

Example: Stock price model

37

- Stock price has three states:
 - Decrease
 - Unchanged
 - Increase




Example: Stock price model

38

□ Stock price has three states:


- Decrease
- Unchanged
- Increase




Markov chains as graphical models

39

1st order



2nd order




time 

Markov chains as graphical models

40


□ What is this then?



Markov chain: US political blog text

41

□ 2.8M words from blog posts



fewer parameters

stronger independence assumptions


richer expressive power

Markov chain: US political blog text

42

□ 0th order (bag of words, unigram)

this trying our putting and funny
and among it herring it obama
but certainly foreign my
c on byron again but from i
i so and i chuck yeah the as but but republicans if
this stay oh so or it mccain bush npr this with what
and they right i while because obama




Markov chain: US political blog text

43

□ 1st order (bigram)

the lack of the senator mccain hadn t keep this story
backwards
while showering praise of the kind of gop weakness
it was mistaken for american economist anywhere in the
white house press hounded the absence of those he s as
a wide variety of this election day after the candidate
b richardson was polled ri in hempstead moderated by
the convention that he had zero wall street journal
argues sounds like you may be the primary
but even close the bill told c e to take the obama on
the public schools and romney
fred flinstone s see how a lick skillet road it s
little sexist remarks




Markov chain: US political blog text

44

□ 2nd order (trigram)

as i can pin them all none of them want to bet that
any of the might be
conservatism unleashed into the privacy rule book and
when told about what paul
fans organized another massive fundraising initiative
yesterday and i don t know what the rams supposedly
want ooh
but she did but still victory dinner
alone among republicans there are probably best not
all of the fundamentalist community
asked for an independent maverick now for
crystallizing in one especially embarrassing



Markov chain: US political blog text

45

□ 4th order (5-gram)

he realizes fully how shallow and insincere
conservative behavior has been he realizes that there
is little way to change the situation

this recent arianna huffington item about mccain
issuing heartfelt denials of things that were actually
true or for that matter about the shi a sunni split and
which side iran was on would get confused about this
any more than someone with any knowledge of us politics
would get confused about whether neo confederates were
likely to be supporting the socialist workers party

at the end of the world and i m not especially
discouraged now that newsweek shows obama leading by
three now

Markov chain: US political blog text

46

□ 99th order (100-gram)

and it would be the work of many hands to catalogue all the ridiculous pronouncements made by this man since his long train of predictions about the middle east has been gaudily disastrously stupefyingly misinformed just the buffoon it seems for the new york times to award with a guest column for if you object to the nyt rewarding failure in quite this way then you re intolerant according to the times editorial page editor andrew rosenthal rosenthal doesn t seem to recognize that his choice of adjectives to describe kristol serious respected are in fact precisely what is at issue for those whom he dismisses as having a fear of opposing views

Hidden Markov Models

47



The Occasionally Dishonest Casino

48

- Game of dice
- Uses a fair die most of the time

Fair

Num	Prob
1	$1/6$
2	$1/6$
3	$1/6$
4	$1/6$
5	$1/6$
6	$1/6$



The Occasionally Dishonest Casino

49

- But will occasionally randomly switch to a loaded die (and randomly back again)

Fair

Num	Prob
1	$1/6$
2	$1/6$
3	$1/6$
4	$1/6$
5	$1/6$
6	$1/6$



Loaded

Num	Prob
1	$1/10$
2	$1/10$
3	$1/10$
4	$1/10$
5	$1/10$
6	$1/2$


The Occasionally Dishonest Casino

50

- The dice are the states, but we only observe the *rolls*

Fair

Num	Prob
1	$1/6$
2	$1/6$
3	$1/6$
4	$1/6$
5	$1/6$
6	$1/6$




Loaded

Num	Prob
1	$1/10$
2	$1/10$
3	$1/10$
4	$1/10$
5	$1/10$
6	$1/2$

The Occasionally Dishonest Casino

51

□ Hidden Markov Model



Hidden Markov Models

52

□ An HMM (λ):

1. N : Number of states in the model

$$S = \{S_1, S_2, \dots, S_N\}$$

2. M : Number of distinct observation symbols in the *alphabet*

$$V = \{v_1, v_2, \dots, v_M\}$$

3. State transition probabilities:

$$\mathbf{A} = [a_{ij}] \text{ where } a_{ij} \equiv P(q_{t+1} = S_j | q_t = S_i)$$

4. Observation probabilities:

$$\mathbf{B} = [b_j(m)] \text{ where } b_j(m) \equiv P(O_t = v_m | q_t = S_j)$$

5. Initial state probabilities:

$$\mathbf{\Pi} = [\pi_i] \text{ where } \pi_i \equiv P(q_1 = S_i)$$

Hidden Markov Models

53



- Two sources of randomness
 - State (fair or loaded)
 - Emission (roll)
- Many possible state sequences for a sequence of observations, e.g:

FFFLLLLFFFFFFL
4352662344636461

Continuous-valued observations

54


- HMM sampling from Gaussian distributions



HMMs as graphical models


55

Simplest case:



HMM unfolded in time

56



Hidden Markov Models

57

- Three main questions:
 - What is the probability of a given sequence of observations?
 - What is the most probable *path* given the observations?
 - What are the parameters of the model given some observations?

Hidden Markov Models

58

- Three main questions:
 - What is the probability of a given sequence of observations? Forward – Backward algorithm
 - What is the most probable *path* given the observations? (Viterbi algorithm)
 - What are the parameters of the model given some observations? (Baum-Welch algorithm)

Probability of a sequence

59

□ Given:

- Some observed sequence O
- A known state sequence Q

□ Calculate $P(O|Q)$

$$P(O|Q, \lambda) = \prod_{t=1}^T P(O_t|q_t, \lambda) = b_{q_1}(O_1) \cdot b_{q_2}(O_2) \cdot \dots \cdot b_{q_T}(O_T)$$

But we don't know Q (the hidden state sequence)!

Probability of a sequence

60

- We want the joint probability

$$\begin{aligned} P(O, Q | \lambda) &= P(q_1) \prod_{t=2}^T P(q_t | q_{t-1}) \prod_{t=1}^T P(O_t | q_t) \\ &= \pi_{q_1} b_{q_1}(O_1) a_{q_1 q_2} b_{q_2}(O_2) \cdots a_{q_{T-1} q_T} b_{q_T}(O_T) \end{aligned}$$

- ... summed over all possible Q

But that's not practical.


Instead use recursive forward-backward procedure

Probability of a sequence

61

- *Forward variable α*
 - P(observations up to and including j)
- *Iterative procedure*

$$\alpha_t(i) \equiv P(O_1 \cdots O_t, q_t = S_i | \lambda)$$



Probability of a sequence

62

□ Initialization:

$$\begin{aligned}\alpha_1(i) &\equiv P(O_1, q_1 = S_i | \lambda) \\ &= P(O_1 | q_1 = S_i, \lambda) P(q_1 = S_i | \lambda) \\ &= \pi_i b_i(O_1)\end{aligned}$$

Probability of a sequence

63

□ Initialization:

$$\begin{aligned}\alpha_1(i) &\equiv P(O_1, q_1 = S_i | \lambda) \\ &= P(O_1 | q_1 = S_i, \lambda) P(q_1 = S_i | \lambda) \\ &= \pi_i b_i(O_1)\end{aligned}$$

□ Recursion:

$$\begin{aligned}\alpha_{t+1}(j) &\equiv P(O_1 \cdots O_{t+1}, q_{t+1} = S_j | \lambda) \\ &= \left[\sum_{i=1}^N \alpha_t(i) a_{ij} \right] b_j(O_{t+1})\end{aligned}$$

Probability of a sequence

64


- $\alpha_T(i)$ is the probability of the whole sequence, ending in state i
- Sum over all states to get $P(O)$

$$\begin{aligned} P(O|\lambda) &= \sum_{i=1}^N P(O, q_T = S_i | \lambda) \\ &= \sum_{i=1}^N \alpha_T(i) \end{aligned}$$

Probability of a sequence

65

- Similar to (π, λ) in Belief Propagation, we can also calculate the *Backward* variable
- Used in state sequence and parameter estimation



Sources & Resources

66

- Recommended books:
 - Bishop: Pattern recognition
 - Durbin, Eddy, Krogh Mitchison: Biological sequence analysis
- Explanation and animations of Markov chains:
 - <http://setosa.io/ev/markov-chains/>