

# **Machine Learning/Advanced Machine Learning**

## **Lecture 6.1: Multilayer Perceptrons**

**Sami S. Brandt**

**Department of Computer Science  
IT University of Copenhagen**

Based on slides originally made by Jes Frellsen

30 September 2019

IT UNIVERSITY OF COPENHAGEN

# Learning Objectives for Week 6

- Reflect the structure of MLP networks
- Apply MLP networks and reflect their role in deep learning
- Use Tensorflow to build simple models and train them
- Explain backpropagation
- Use simple regularisation methods with neural networks
- Explain radial basis function networks
- Explain the principle of convolutional neural networks

# Outline of lecture


Feed-forward neural networks

Training neural networks


# Re-cap: Categories of Machine Learning

**Supervised learning:**  $\mathcal{D} = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_n, t_n)\}$

## Classification




## Regression




**Unsupervised learning:**  $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$

## Clustering




## Density estimation




# Linear models for regression and classification

## Regression



$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=1}^M w_j \phi_j(\mathbf{x}) = \mathbf{w}^\top \boldsymbol{\phi}(\mathbf{x})$$

## Classification



$$y(\mathbf{x}, \mathbf{w}) = f \left( \sum_{j=1}^M w_j \phi_j(\mathbf{x}) \right) = f(\mathbf{w}^\top \boldsymbol{\phi}(\mathbf{x}))$$

where  $f$  is an **activation function** such as the **sigmoid** or **softmax**.


## Advantages:

- Useful computational properties: have **closed form expressions**
- Can model “arbitrary” complicated functions.

## Disadvantages:

- We have to **choose** the basis functions  $\phi_i(\cdot)$ , and these are **not adapted to data**.

# Artificial neuron



**Input:**  $\mathbf{x} = (x_1, \dots, x_D)^\top$

**Weights:**  $\mathbf{w} = (w_1, \dots, w_D)^\top$


**Bias:**  $b$

**Activation:**  $a = \sum_{i=1}^D w_i x_i + b = \mathbf{w}^\top \mathbf{x} + b$


**Output:**  $z = h(a)$

Where  $h(\cdot)$  is a **nonlinear activation function**.

# Combining artificial neurons




# Combining artificial neurons




Where  $W = \begin{bmatrix} w_{11} & w_{12} & \dots & w_{1D} \\ w_{21} & w_{22} & \dots & w_{2D} \end{bmatrix}$  and  $h$  is applied element wise.



# A two-layer neural network



# A two-layer neural network




$\mathbf{x}$

$$\mathbf{z} = h(W^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

$$\mathbf{y} = \sigma(W^{(2)}\mathbf{z} + \mathbf{b}^{(2)})$$

$$\mathbf{y}(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \sigma(W^{(2)}h(W^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) + \mathbf{b}^{(2)})$$

# A deep neural network



$$\mathbf{y}(\mathbf{x}, \mathbf{W}, \mathbf{b}) = f^{(4)}(f^{(3)}(f^{(2)}(f^{(1)}(\mathbf{x}))))$$


where

$$f^{(\ell)}(\mathbf{z}) = h^{(\ell)}(\mathbf{W}^{(\ell)}\mathbf{z} + \mathbf{b}^{(\ell)})$$

# Activations functions

## For hidden layers:

- Sigmoid:  $h(a) = \frac{1}{1+\exp(-a)}$
- Arctangent<sup>1</sup>:  $h(a) = \operatorname{atan}(a)$
- Hyperbolic tangent:  $h(a) = \tanh(a)$
- **Rectified linear**<sup>2</sup>:  $h(a) = \max(0, a)$



## For output layer:

- Regression: identity,  $h(a) = a$
- Binary classification: sigmoid,  $h(a) = \frac{1}{1+\exp(-a)}$
- Multiclass classification: softmax,  $h(a_k) = \frac{\exp(a_k)}{\sum_j \exp(a_j)}$


<sup>1</sup>The inverse of  $\tan$ .

<sup>2</sup>Often called ReLU, but this is a misnomer as it means *rectified linear unit*

# Activations functions

## For hidden layers:

- Sigmoid:  $h(a) = \frac{1}{1+\exp(-a)}$
- Arctangent<sup>1</sup>:  $h(a) = \operatorname{atan}(a)$
- Hyperbolic tangent:  $h(a) = \tanh(a)$
- **Rectified linear**<sup>2</sup>:  $h(a) = \max(0, a)$



## For output layer:

- Regression: identity,  $h(a) = a$
- Binary classification: sigmoid,  $h(a) = \frac{1}{1+\exp(-a)}$
- Multiclass classification: softmax,  $h(a_k) = \frac{\exp(a_k)}{\sum_j \exp(a_j)}$

## Exercise

Why do we need activation functions in the **hidden layers**?

Why do we need activation functions in the **output layer for classification**?

<sup>1</sup>The inverse of  $\tan$ .

<sup>2</sup>Often called ReLU, but this is a misnomer as it means *rectified linear unit*

# Universal approximation theorem

Neural networks are universal approximators (Bishop):

“A two-layer network with linear outputs can uniformly **approximate any continuous function** on a compact input domain (compact subset of  $\mathbb{R}^N$ ) to **arbitrary accuracy** provided the network has **sufficiently large number of hidden units**”

# How do we program neural networks?

We could in principle program them directly in Python / NumPy

However, we need **derivatives** / **gradients** for training (later)

We will use **TensorFlow**, which has automatic differentiations.

**TensorFlow** works in a **declarative** style:

- First you declare/define a **dataflow graph**.
- Then you use a session to run/evaluate operations in the graph.

**Example: Notebook 1**



# Outline of lecture

Feed-forward neural networks

Training neural networks

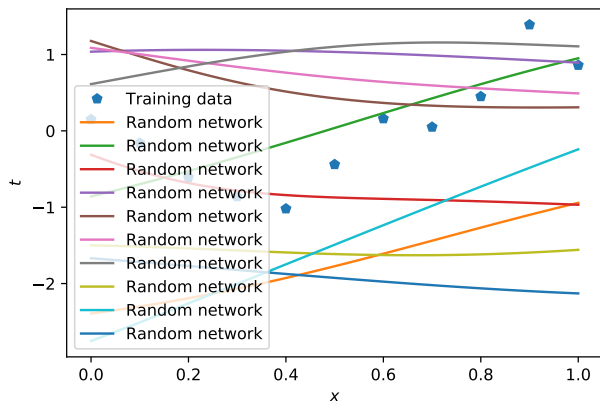


# Training neural networks

We looked at the definition of a neural network, or a **neural network function**  $y(\mathbf{x}, \mathbf{W}, \mathbf{b})$ .

If we are given some data as **input vectors**  $\mathbf{X} = \{\mathbf{x}_n\}_{n=1}^N$  and **target vectors**  $\mathbf{t} = \{\mathbf{t}_n\}_{n=1}^N$ :

- How can we find  $\mathbf{W}$  and  $\mathbf{b}$ ?



Example: Notebook 2

# Regression with neural networks

We find  $\mathbf{W}$  and  $\mathbf{b}$  by minimizing an **error function** that **measures the misfit** between  $\mathbf{y}(\mathbf{x}, \mathbf{W}, \mathbf{b})$  and  $\mathbf{t} = \{\mathbf{t}_n\}_{n=1}^N$ .

The **sum-of-squares** error function is given by



$$E(\mathbf{W}, \mathbf{b}) = \frac{1}{2} \sum_{n=1}^N \|\mathbf{y}(x_n, \mathbf{W}, \mathbf{b}) - \mathbf{t}_n\|^2.$$

**Recall**, this corresponds to **finding the MLE** of  $\mathbf{W}$  and  $\mathbf{b}$  with the likelihood function

$$p(\mathbf{t}|\mathbf{X}, \mathbf{W}, \mathbf{b}, \beta) = \prod_{n=1}^N p(\mathbf{t}_n|\mathbf{x}_n, \mathbf{W}, \mathbf{b}, \beta)$$

where

$$p(\mathbf{t}_n|\mathbf{x}_n, \mathbf{W}, \mathbf{b}, \beta) = \mathcal{N}(\mathbf{t}_n|\mathbf{y}(x_n, \mathbf{W}, \mathbf{b}), \beta^{-1}I)$$



# Binary classification with neural networks

Now consider a classification problem, where  $t_n \in \{0, 1\}$ .

In this cases we use the **sigmoid activation** function

$$y(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \sigma(a^{(L)}) = \frac{1}{1 + \exp(-a^{(L)})}$$

We then use the binomial distribution

$$p(t|\mathbf{x}, \mathbf{W}, \mathbf{b}) = y(\mathbf{x}, \mathbf{W}, \mathbf{b})^t (1 - y(\mathbf{x}, \mathbf{W}, \mathbf{b}))^{1-t} \quad (1)$$



Assuming i.i.d. training data we get the likelihood function

$$p(\mathbf{t}|\mathbf{X}, \mathbf{W}, \mathbf{b}, \beta) = \prod_{n=1}^N p(t_n|\mathbf{x}_n, \mathbf{W}, \mathbf{b}) \quad (2)$$

We then set the error function to the negative log-likelihood

$$E(\mathbf{W}, \mathbf{b}) = - \sum_{n=1}^N t_n \ln y(\mathbf{x}_n, \mathbf{W}, \mathbf{b}) + (1 - t_n) \ln(1 - y(\mathbf{x}_n, \mathbf{W}, \mathbf{b})) \quad (3)$$

Which is also called the **cross-entropy error function**



# Binary classification with neural networks

Now consider a classification problem, where  $t_n \in \{0, 1\}$ .

In this cases we use the **sigmoid activation** function

$$y(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \sigma(a^{(L)}) = \frac{1}{1 + \exp(-a^{(L)})}$$

We then use the binomial distribution

$$p(t|\mathbf{x}, \mathbf{W}, \mathbf{b}) = y(\mathbf{x}, \mathbf{W}, \mathbf{b})^t (1 - y(\mathbf{x}, \mathbf{W}, \mathbf{b}))^{1-t} \quad (1)$$

Assuming i.i.d. training data we get the likelihood function

$$p(\mathbf{t}|\mathbf{X}, \mathbf{W}, \mathbf{b}, \beta) = \prod_{n=1}^N p(t_n|\mathbf{x}_n, \mathbf{W}, \mathbf{b}) \quad (2)$$



We then set the error function to the negative log-likelihood

$$E(\mathbf{W}, \mathbf{b}) = - \sum_{n=1}^N t_n \ln y(\mathbf{x}_n, \mathbf{W}, \mathbf{b}) + (1 - t_n) \ln(1 - y(\mathbf{x}_n, \mathbf{W}, \mathbf{b})) \quad (3)$$

Which is also called the **cross-entropy error function**

## Exercise

Use equation (1) and (2) to derive (3).



# Multi-class classification with neural networks

Now consider a classification problem, where  $t_n \in \{0, 1\}^K$ .

In this cases we use the **softmax activation** function


$$y_k(\mathbf{x}, \mathbf{W}, \mathbf{b}) = \sigma(\mathbf{a}^{(L)}) = \frac{\exp(a_k)}{\sum_{j=1}^K \exp(a_j)}$$

In this case we obtain the error function

$$E(\mathbf{W}, \mathbf{b}) = - \sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln y_k(\mathbf{x}_n, \mathbf{W}, \mathbf{b}). \quad (4)$$

from the log-likelihood function.

This is also called the **cross-entropy error function**



# Parameter optimization and gradient descent

We want to find  $\mathbf{W} = (\mathbf{W}, \mathbf{b})$  that **minimizes**  $E(\mathbf{W})$ , i.e. find  $\mathbf{W}$  such that  $\nabla E(\mathbf{W}) = 0$ .

**Gradient descent** starts with an initial random point  $\mathbf{W}^{(0)}$ , and iteratively refines it by following the steepest descent direction:

$$\mathbf{W}^{(\tau+1)} = \mathbf{W}^{(\tau)} - \eta \nabla E(\mathbf{W}^{(\tau)})$$


where  $\eta$  is called the **learning rate**.

Normally the gradient is calculated on the full dataset (*batch* optimization).


To avoid getting stuck in local minima, we can calculate the gradient on **mini-batches**:

$$\mathbf{W}^{(\tau+1)} = \mathbf{W}^{(\tau)} - \eta \nabla E_s(\mathbf{W}^{(\tau)})$$

where  $E_s$  is the error on a subset of the data.




# Examples: regression



Example: Notebook 3

# Examples: binary classification



Example: Notebook 4



# Next lecture

- Backpropagation
- Regularisation of neural networks
- Radial basis functions
- Brief introduction to CNNs

# References I



Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.