

CHAPTER 10:

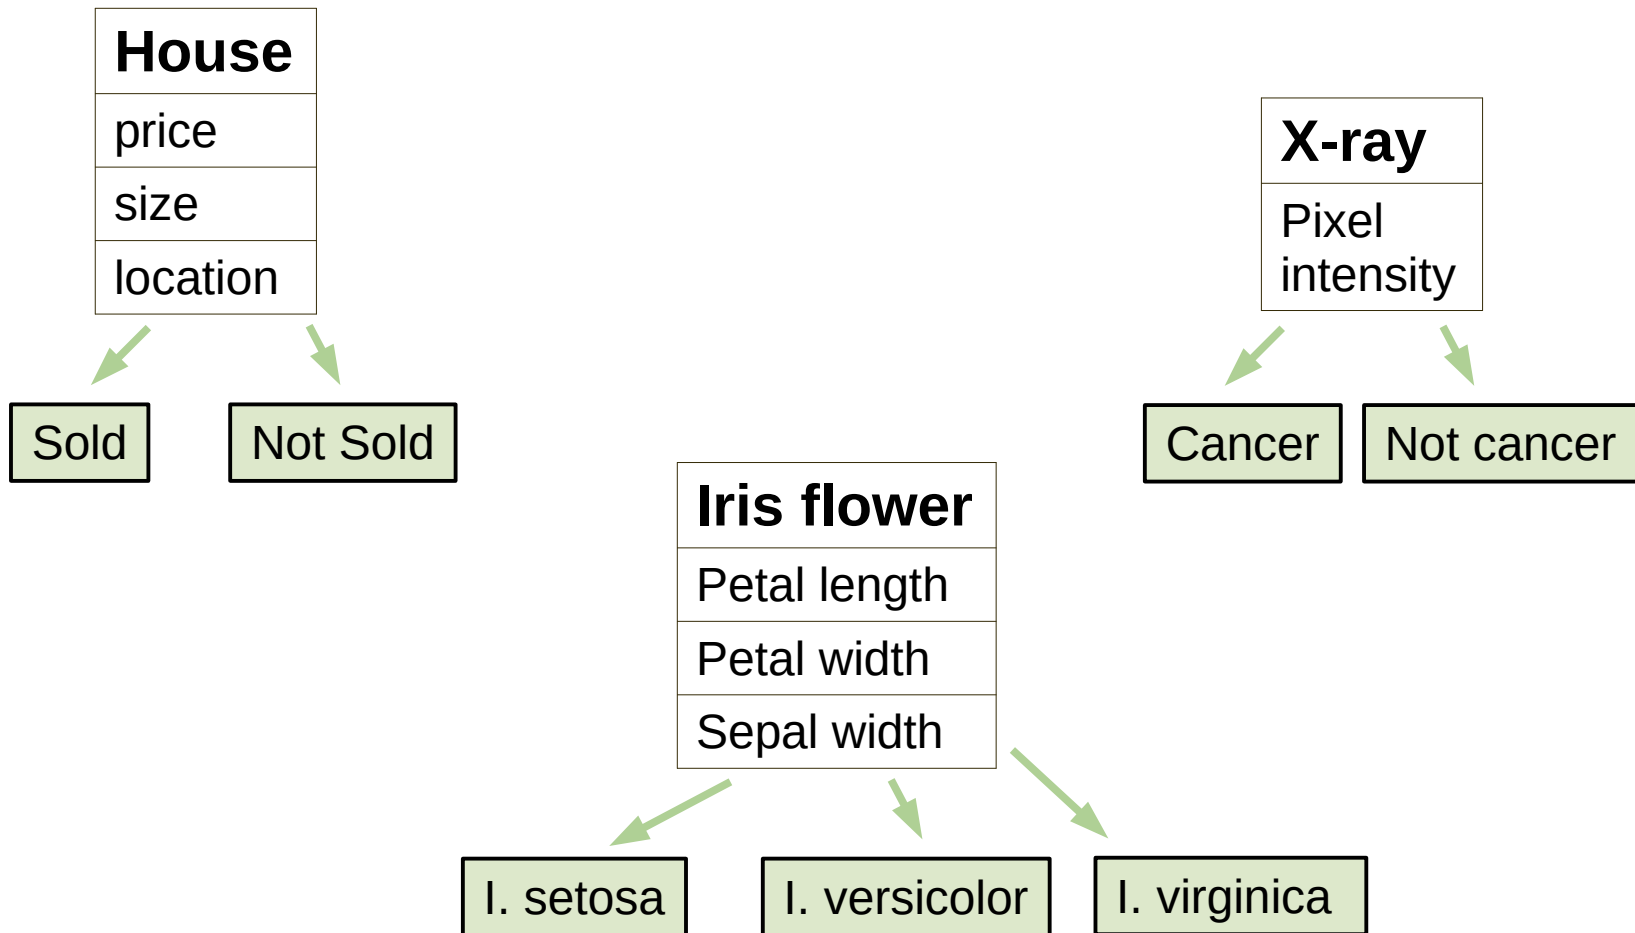
LINEAR DISCRIMINATION

Some slides adapted from:

- *E. Alpaydin, cmpe.boun.edu.tr/~ethem/i2ml3e*
- *M.A Carreira-Perpinan, faculty.ucmerced.edu/mcarreira-perpinan/teaching/CSE176/lecturenotes.pdf*

Classification examples

2



Classification

3

- In previous lectures we looked at classification:
 - Training time: learn a set of discriminant functions $\{g_i(\mathbf{x})\}$
 - Test time: given a new instance \mathbf{x} , choose class k with highest value of $\{g_i(\mathbf{x})\}$

Likelihood- vs. Discriminant-based Classification

4

- There are two approaches to learning the discriminant functions $g_i(\mathbf{x})$:
 - Likelihood-based
 - Discriminant-based


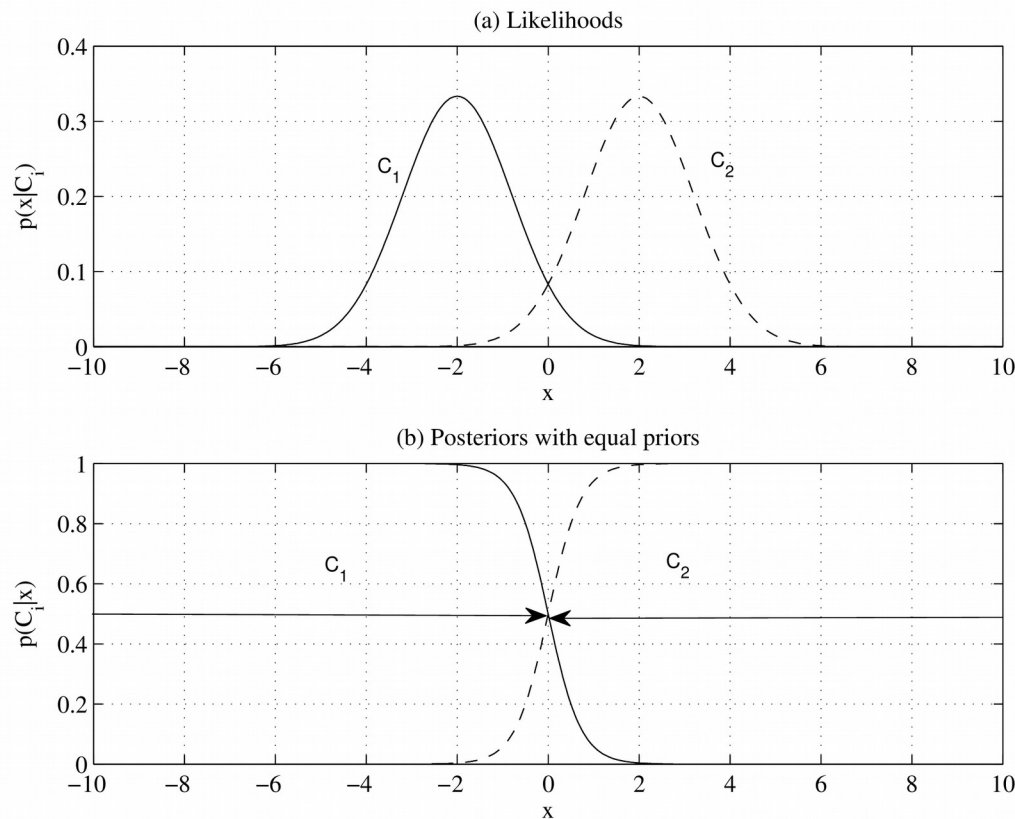
Likelihood-based Classification

5

- Assume a model for $p(\mathbf{x}|C_i)$
- Estimate $p(\mathbf{x}|C_i)$ and $p(C_i)$ from the data
- Use Bayes' rule to calculate $P(C_i|\mathbf{x})$
e.g. $g_i(\mathbf{x}) = \log P(C_i|\mathbf{x})$

Likelihood-based Classification

6





Likelihood-based Classification

7

- We learn both
 - Class boundaries
 - where $p(\mathbf{x}|C_i) = p(\mathbf{x}|C_j)$
 - Class densities
- *Generative* approach
- Previous chapters: parametric and non-parametric methods for finding $p(\mathbf{x}|C_i)$

Likelihood-based Classification

8



Discriminant-based Classification

9

- Learn only class *boundaries*
- Assume a model for $g_i(\mathbf{x}|\Phi_i)$, learn this directly
- No density estimation

Discriminant-based Classification

10

- Estimating the boundaries is enough; no need to accurately estimate the densities inside the boundaries
- A simpler problem than density estimation
- This and following chapters: Discriminant functions

Linear Discriminant

11

- Define a model $g_i(\mathbf{x}|\Phi_i)$ for each class discriminant
- The *linear discriminant* has the form:

$$g_i(\mathbf{x} | \mathbf{w}_i, w_{i0}) = \mathbf{w}_i^T \mathbf{x} + w_{i0} = \sum_{j=1}^d w_{ij} x_j + w_{i0}$$

- Assumption: instances from one class are linearly separable from instances of other classes
- Learning means finding the values of $\{\Phi_i\}$ that optimize separation

Linear Discriminant

12

- Advantages:
 - Faster to train
 - Simple at test time: $O(d)$ space/computation
 - Simple interpretation:
 - Output is weighted sum of attributes
 - Sign of weights: are effects positive/negative
 - Magnitudes of weight: How important

Linear Discriminant

13


- Optimal when $p(\mathbf{x}|C_i)$ are Gaussian with shared covariance matrix
- Useful when classes are (almost) linearly separable
- In many cases, good enough
 - Try linear discrimination before more complex models

Generalizing the Linear Model

14

- When a linear model (linear in \mathbf{x}) is not good enough
 - We can add higher order terms
- Example: Quadratic discriminant:

$$g_i(\mathbf{x} | \mathbf{W}_i, \mathbf{w}_i, w_{i0}) = \mathbf{x}^T \mathbf{W}_i \mathbf{x} + \mathbf{w}_i^T \mathbf{x} + w_{i0}$$



Generalizing the Linear Model

15

- Alternative: preprocess input
- Higher-order (product) terms:

$$z_1 = x_1, z_2 = x_2, z_3 = x_1^2, z_4 = x_2^2, z_5 = x_1 x_2$$

Map from \mathbf{x} to \mathbf{z} using nonlinear basis functions
and use a linear discriminant in \mathbf{z} -space

Generalizing the Linear Model

16

- Linear combination of non-linear functions of \mathbf{x}
- Discriminant:

$$g_i(\mathbf{x}) = \sum_{j=1}^k w_j \phi_{ij}(\mathbf{x})$$

where $\phi_{ij}(\mathbf{x})$ are *basis functions*

Generalizing the Linear Model

17

- Linear combination of non-linear functions of \mathbf{x}
- Discriminant:

$$g_i(\mathbf{x}) = \sum_{j=1}^k w_j \phi_{ij}(\mathbf{x})$$

where $\phi_{ij}(\mathbf{x})$ are *basis functions*


Basis function examples:

(basis functions are used in later chapters; neural nets and SVMs)

- $\sin(x_1)$
- $\exp(-(x_1 - m)^2/c)$
- $\exp(-\|\mathbf{x} - \mathbf{m}\|^2/c)$
- $\log(x_2)$
- $1(x_1 > c)$
- $1(ax_1 + bx_2 > c)$

Two Classes

18




$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

$$\text{choose } \begin{cases} C_1 & \text{if } g(\mathbf{x}) > 0 \\ C_2 & \text{otherwise} \end{cases}$$

Two Classes

19



$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$


\mathbf{w} is orthogonal to the hyperplane

The origin $\mathbf{x} = \mathbf{0}$ is on the

{	positive side	if $w_0 > 0$
	boundary	if $w_0 = 0$
	negative side	if $w_0 < 0$

Two Classes

20



$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

\mathbf{w} is orthogonal to the hyperplane
- determines orientation


Determines location wrt origin

The origin $\mathbf{x} = \mathbf{0}$ is on the

{	positive side	if $w_0 > 0$
	boundary	if $w_0 = 0$
	negative side	if $w_0 < 0$

Geometry

21



Multiple Classes

22

- With $K > 2$ classes
- Two common ways to classify:
 - One-vs-all (linear separability)
 - One-vs-one (pairwise linear separability)

One-vs-all

23

- K discriminants: $g_i(\mathbf{x} | \mathbf{w}_i, w_{i0}) = \mathbf{w}_i^T \mathbf{x} + w_{i0}$
- Training:
 - train each g_i to classify instances of C_i vs all other instances
- Testing:
 - Choose with largest discriminant


Choose C_i if

$$g_i(\mathbf{x}) = \max_{j=1}^K g_j(\mathbf{x})$$

One-vs-all

24

Ideally, classes are completely linearly separable



One-vs-one

- $K(K-1)/2$ discriminants

$$g_{ij}(\mathbf{x} | \mathbf{w}_{ij}, w_{ij0}) = \mathbf{w}_{ij}^T \mathbf{x} + w_{ij0}$$

- One for each pair of classes

- Training:


- train g_{ij} to classify instances of C_i vs instances of C_j
- ignore instances of other classes

- Testing:

- choose C_i if $\forall j \neq i, g_{ij}(\mathbf{x}) > 0$
- Or pick the class with largest summed discriminant

$$g_i(\mathbf{x}) = \sum_{j \neq i} g_{ij}(\mathbf{x})$$

One-vs-one



Both methods divide the input space into K convex decision regions

How do we find \mathbf{w} ?

27

- We know what we want!
 - To find \mathbf{w}
- But how do we do that?
 - In some cases \mathbf{w} can be found analytically
 - But often this is not possible: Iterative optimization instead

Gradient Descent

28



www.mec.ca

Gradient Descent

29

- There are many different iterative optimization methods
- Gradient Descent is a popular one
 - Simple
 - Reasonably effective
 - But can be very slow

Gradient Descent

30

- Error function: $E(\mathbf{w}|X)$ with parameters \mathbf{w} on sample X
- Want to find the optimal \mathbf{w} (with minimal error):

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w} | X)$$

Gradient Descent

31

- Error function: $E(\mathbf{w}|X)$ with parameters \mathbf{w} on sample X
- Want to find the optimal \mathbf{w} (with minimal error):
 $\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w} | X)$
- Begin: Set random \mathbf{w}
- Iterate:
 - Calculate update term: $\Delta w_i = -\eta \frac{\partial E}{\partial w_i}, \forall i$
 - Set new (updated) w :
$$w_i = w_i + \Delta w_i$$


Gradient Descent

32

- *Gradient* (vector of partial derivatives)

$$\nabla_w E = \left[\frac{\partial E}{\partial w_1}, \frac{\partial E}{\partial w_2}, \dots, \frac{\partial E}{\partial w_d} \right]^T$$

- Learning factor (η):


$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}, \forall i$$


- Stop iterating when:

$$\nabla E(\mathbf{w}) \approx \mathbf{0}$$

Gradient-Descent

33



Logistic Regression

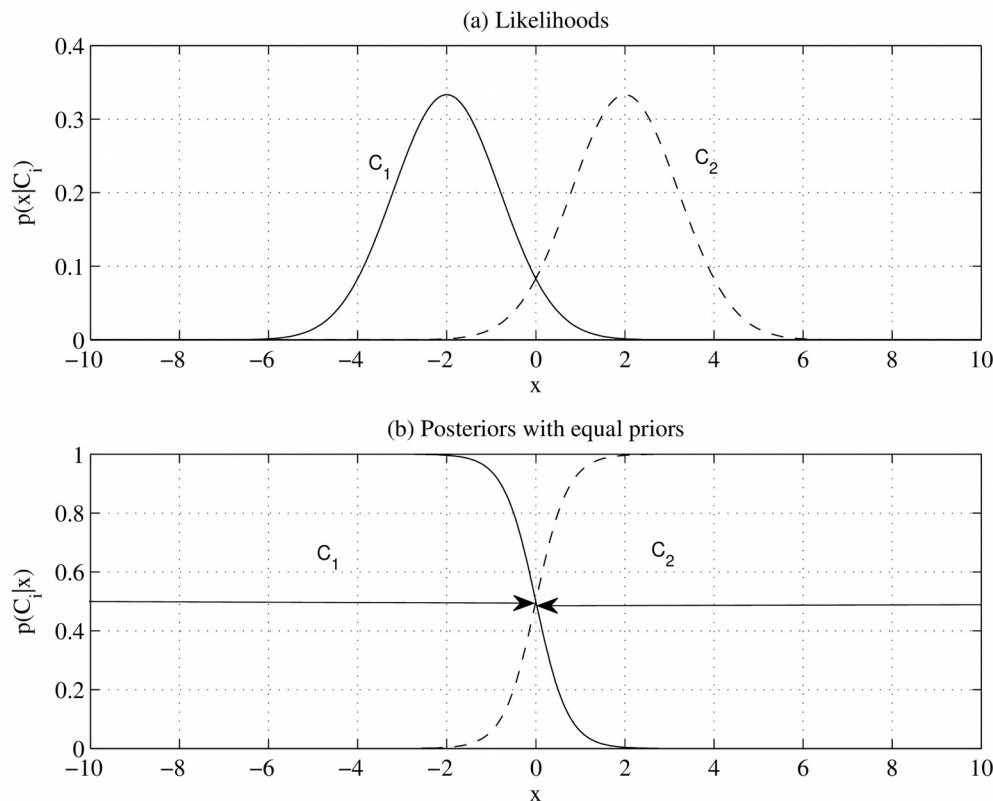
34



Image via flickr, CC

Previous chapters: Sigmoid posterior


35



Instead of first estimating densities, then calculating posteriors:
We want to model the posterior directly using the logistic function

Sigmoid (Logistic) Function

36



Transform a linear discriminant to a posterior probability:

Calculate $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$ and choose C_1 if $g(\mathbf{x}) > 0$, or

Calculate $y = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0)$ and choose C_1 if $y > 0.5$

Logistic Discrimination

37

Why does the logistic approach work?

We know that:

- For Gaussian class densities with equal covariance:

$$g_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + w_{i0}$$

- Parameters can be calculated analytically:

$$\begin{aligned} \mathbf{w}_i &= \Sigma^{-1} \boldsymbol{\mu}_i \\ w_{i0} &= -\frac{1}{2} \boldsymbol{\mu}_i^T \Sigma^{-1} \boldsymbol{\mu}_i + \log P(C_i) \end{aligned}$$

Logistic Discrimination

38

Log odds is linear (again for Gaussian with shared cov):

$$\begin{aligned}\text{logit}(P(C_1|\mathbf{x})) &= \log \frac{P(C_1|\mathbf{x})}{1 - P(C_1|\mathbf{x})} = \log \frac{P(C_1|\mathbf{x})}{P(C_2|\mathbf{x})} \\ &= \log \frac{p(\mathbf{x}|C_1)}{p(\mathbf{x}|C_2)} + \log \frac{P(C_1)}{P(C_2)} \\ &= \log \frac{(2\pi)^{-d/2} |\Sigma|^{-1/2} \exp[-(1/2)(\mathbf{x} - \boldsymbol{\mu}_1)^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_1)]}{(2\pi)^{-d/2} |\Sigma|^{-1/2} \exp[-(1/2)(\mathbf{x} - \boldsymbol{\mu}_2)^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_2)]} + \log \frac{P(C_1)}{P(C_2)} \\ &= \mathbf{w}^T \mathbf{x} + w_0\end{aligned}$$

where

$$\begin{aligned}\mathbf{w} &= \Sigma^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) \\ w_0 &= -\frac{1}{2}(\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2)^T \Sigma^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) + \log \frac{P(C_1)}{P(C_2)}\end{aligned}$$

Logistic Discrimination

39

- So the Log-Odds ($\text{logit}(P(C_1|\mathbf{x}))$) is linear:

$$\log \frac{P(C_1|\mathbf{x})}{1 - P(C_1|\mathbf{x})} = \mathbf{w}^T \mathbf{x} + w_0$$

- The *inverse* of the log-odds is the logistic function
- Which means that:

$$P(C_1|\mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

Logistic Discrimination

40

- So the posterior can be described by a sigmoid:

$$P(C_1|\mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

- We showed this in the case of two Gaussian class densities with shared cov
- It also holds in some other cases

K>2 classes, softmax

41

- When there are more than two classes:




Softmax

$$y_i = \hat{P}(C_i|\mathbf{x}) = \frac{\exp[\mathbf{w}_i^T \mathbf{x} + w_{i0}]}{\sum_{j=1}^K \exp[\mathbf{w}_j^T \mathbf{x} + w_{j0}]}, \quad i = 1, \dots, K$$

- If the 'linear output' for one class (C_k) is much larger than for the others:
 - $P(C_k|\mathbf{x})$ will be close to one
 - Close to 0 for other classes
 - Together they all sum to 1

Examples


42



Training a logistic classifier

43

How do we find the correct parameters?



Univariate example

Training a logistic classifier

44

K= 2: We have a training set with input data and labels
e.g [house price, size,] \rightarrow {sold, not sold}

We want to learn the parameters for

$$P(C_1|\mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

Training a logistic classifier

45

K= 2: We have a training set with input data and labels
e.g [house price, size,] \rightarrow {sold, not sold}

We want to learn the parameters for

$$P(C_1|\mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

There are two ways of doing that:

- Model the outcome as a Bernoulli distribution. Find the Maximum Likelihood parameter estimates.
- Use least-square regression, considering the labels {0,1} real values.

Training a logistic classifier

46

$K=2$: We have a training set with input data and labels
e.g [house price, size,] \rightarrow {sold, not sold}

We want to learn the parameters for

$$P(C_1|\mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

There are two ways of doing that:

- Model the outcome as a Bernoulli distribution. Find the Maximum Likelihood parameter estimates.
- Use least-square regression, considering the labels $\{0,1\}$ real values.

In both cases: No closed-form solution. Need Gradient Descent

The ML & Bernoulli approach

47

- Labels are a Bernoulli distribution

- $r^t | \mathbf{x}^t \sim \text{Bernoulli}(y^t)$

- Sample likelihood is

- $$l(\mathbf{w}, w_0 | \mathcal{X}) = \prod_t (y^t)^{r^t} (1 - y^t)^{(1-r^t)}$$

- Turn this into an error function to minimize

- $$E(\mathbf{w}, w_0 | \mathcal{X}) = - \sum_t r^t \log y^t + (1 - r^t) \log(1 - y^t)$$

- (this is *Cross Entropy*)

The ML & Bernoulli approach

48

- To do Gradient Descent we need the derivative $\frac{dy}{da} = y(1 - y)$

- This gives update equations for the descent:


$$\begin{aligned}\Delta w_j &= -\eta \frac{\partial E}{\partial w_j} = \eta \sum_t \left(\frac{r^t}{y^t} - \frac{1 - r^t}{1 - y^t} \right) y^t(1 - y^t) x_j^t \\ &= \eta \sum_t (r^t - y^t) x_j^t, j = 1, \dots, d\end{aligned}$$

$$\Delta w_0 = -\eta \frac{\partial E}{\partial w_0} = \eta \sum_t (r^t - y^t)$$

The ML & Bernoulli approach

49

- Initialize with weights near 0
- Stop when all samples are classified correctly
 - or even before (regularization)



$K > 2$ Classes

50

- Similar, but labels are modelled as multinomial instead of Bernoulli
- We have the softmax instead of a single logistic function

$$y_i = \hat{P}(C_i | \mathbf{x}) = \frac{\exp[\mathbf{w}_i^T \mathbf{x} + w_{i0}]}{\sum_{j=1}^K \exp[\mathbf{w}_j^T \mathbf{x} + w_{j0}]}, \quad i = 1, \dots, K$$

Training by least-squares


51

- Another way to train the logistic classifier: Least squares regression
- “Discrimination by regression”
- Can be useful when classes are not mutually exclusive and exhaustive


Training by least-squares

52

- Two-way classification:



- Use the class codes $\{0,1\}$ as numeric target values for regression
- Fit to the logistic function by minimizing sum-of-squares



Training by least-squares

53

- Function to fit:

$$y^t = \text{sigmoid}(\mathbf{w}^T \mathbf{x}^t + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x}^t + w_0)]}$$

- Fit by minimizing sum of square Error:

$$E(\mathbf{w}, w_0 | \mathcal{X}) = \frac{1}{2} \sum_t (r^t - y^t)^2$$

Training by least-squares

54

- Update equations for gradient descent

$$\Delta \mathbf{w} = \eta \sum_t (r^t - y^t) y^t (1 - y^t) \mathbf{x}^t$$

$$\Delta w_0 = \eta \sum_t (r^t - y^t) y^t (1 - y^t)$$

Training by least-squares

55

- Can also be used for $K > 2$ classes
- A separate logistic function is fit for each class
- The resulting class probabilities $P(C|x)$ don't necessarily sum to one!!
 - So x may be predicted to belong to more than one class, or none at all

Sources & resources

56

- Linear and GLM-intro:
byclb.com/TR/Tutorials/neural_networks/ch9_1.htm
- Math animations (esp. linear algebra & calculus)
youtube.com/channel/UCYO_jab_esuFRV4b17AjtAw
- Various stats topics explained
youtube.com/user/joshstarmer
- Machine learning lectures, Cornell
[youtube.com/playlist?
list=PLI8OIHZGYOQ7bkVbuRthEsaLr7bONzbXS](https://youtube.com/playlist?list=PLI8OIHZGYOQ7bkVbuRthEsaLr7bONzbXS)