

Lab – Deep Learning
Data Mining, Spring 2018

Carol (berm@itu.dk) / Daniel (dafr@itu.dk) / Mathias (jams@itu.dk)

Today's Lab: Deep Learning Framework

Deep Learning Framework

- In today's lab you will try your hand at the deep learning framework, DL4J
 - <https://deeplearning4j.org/documentation>



Two part exercise

- Purpose: get experience with deep learning
- Part 1: Experiment with an existing CNN topology and get familiar with Deeplearning4j.
 - A very high accuracy should be achieved
 - Should be quite straight forward
- Part 2: See if you can create a well performing CNN topology yourself – hard


Part 1: DL4J - Code

- All the code is already provided, with some default parameters and variables
- How to import and run the code can be found at <https://deeplearning4j.org/mnist-for-beginners>
 - Also contains documentation and information regarding the provided code
- All dependencies taken care of (using Maven)

```
public class MLPmnistSingleLayerExample {  
  
    private static Logger log = LoggerFactory.getLogger(MLPMnistSingleLayerExample.class);  
  
    public static void main(String[] args) throws Exception {  
        //number of rows and columns in the input pictures  
        final int numRows = 28;  
        final int numColumns = 28;  
        int outputNum = 10; // number of output classes  
        int batchSize = 128; // batch size for each epoch  
        int rngSeed = 123; // random number seed for reproducibility  
        int numEpochs = 15; // number of epochs to perform  
  
        //Get the DataSetIterators:  
        DataSetIterator mnistTrain = new MnistDataSetIterator(batchSize, true, rngSeed);  
        DataSetIterator mnistTest = new MnistDataSetIterator(batchSize, false, rngSeed);  
  
        log.info("Build model....");  
        MultiLayerConfiguration conf = new NeuralNetConfiguration.Builder()  
            .seed(rngSeed) //include a random seed for reproducibility  
            // use stochastic gradient descent as an optimization algorithm  
            .optimizationAlgo(OptimizationAlgorithm.STOCHASTIC_GRADIENT_DESCENT)  
            .iterations(1)  
            .learningRate(0.006) //specify the learning rate  
            .updater(Updater.NESTEROVS).momentum(0.9) //specify the rate of change of the learning rate.  
            .regularization(true).l2(1e-4)  
            .list()  
    }  
}
```

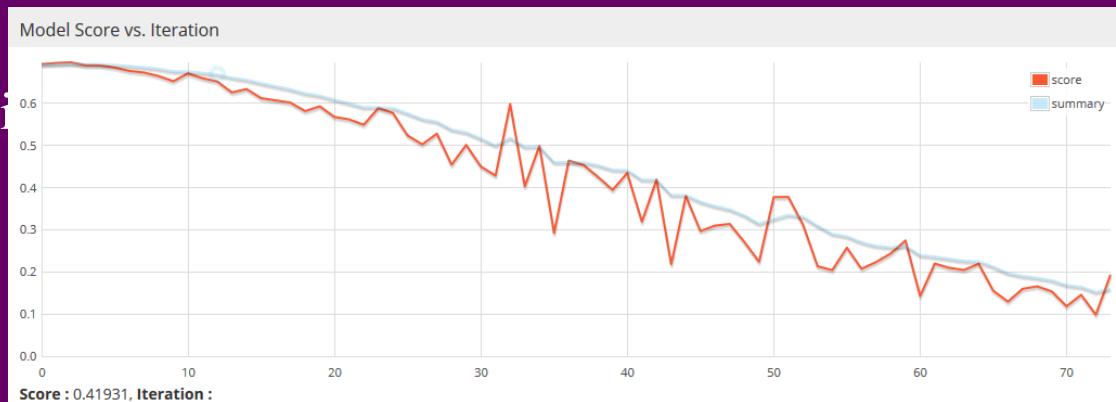
Part 1: Data

- MNIST dataset
 - Giant dataset containing images of handwritten integers (grayscale)
 - $|\text{TrainingSet}| = 60,000$
 - $|\text{TestSet}| = 10,000$
 - Input size: $28 \times 28 \times 1$
 - Output size: 10



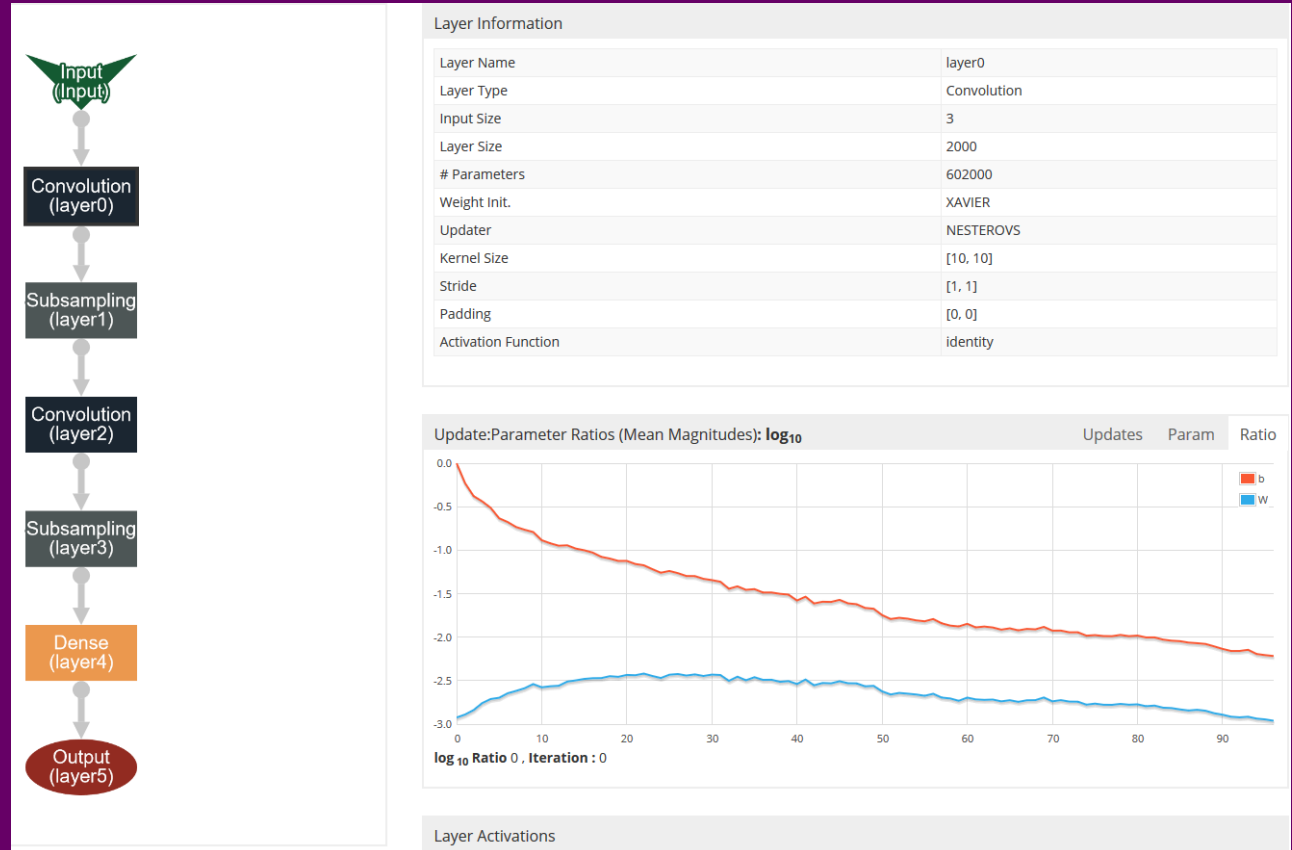
Part 1: Plan of Attack

- Read the tutorial found on previous slide to understand the code
- Download the .zip
- Import the project in your IDE
- Run the code
- Fiddle around with parameters and values
- Get lowest score (in lowest iteration)
- Visualize data




Visualize the learning

- UI server can be accessed at: <http://localhost:9000>
- Visualize learning and model score
- Somewhat visualize the CNN being trained and the layers



Part 2: Design your own CNN

- Cifar10 dataset
- 60,000 images
- 10 distinct object categories
- Colors - 3 channels (RGB)
– a whole new dimension



Part 2: Dataset and Preprocessing

- Already preprocessed.
 - Preprocessing has a big influence on the result!
- Dataset split into test (10,000 images) and training (50,000 images)
- [3x32x32]
- All categories have around 5,000 examples – important!
- Normalize image data [0..1] – done when reading the images
- The dataset is provided



Part 2: Plan of attack

- Fill in the missing parts in Cifar10Example.java
- Design the CNN – get started here:
<https://deeplearning4j.org/convolutionalnets.html#dl4j-code-example>
- Experiment with feature maps.
- Compare your accuracy with others
<https://www.kaggle.com/c/cifar-10/leaderboard>




Performance

- The code provided does not utilize the GPU
- Can take a while to train – work on group projects meanwhile



Deep Learning Framework

- Optionally you can use either Keras or Tensorflow, which uses python
 - <https://keras.io/>
 - <https://www.tensorflow.org/>



Optional: Keras / Tensorflow

- Keras is simple, Tensorflow more complicated
 - Keras uses Tensorflow, meaning installation of Tensorflow is required to use Keras
 - Check <https://www.tensorflow.org/install/> on how to install Tensorflow
- No code provided, but for Keras, check <https://github.com/keras-team/keras/tree/master/examples> for examples

Thanks for listening!