

# Intelligent Systems Programming

Lecture 9: Local Search (Meta-Heuristics)



# Local Search: Overview & Motivation

- Makes small changes to a **complete solution**
- Has low memory consumption
- Effective at solving large optimization problems
- Easy to implement real-world constraints
- Widely used in practice
- However:
  - Incomplete method, i.e. no guarantees of optimality

# Outline

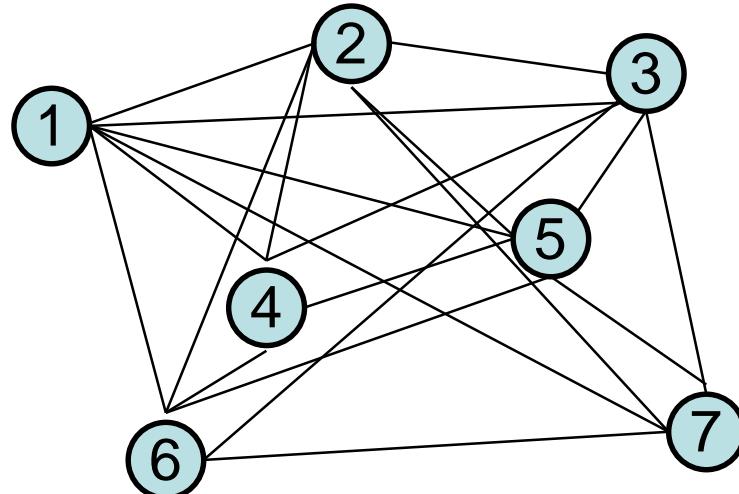
1. Local Search Basics
2. Meta Heuristics
  1. Hill Climbing
  2. Simulated Annealing
  3. Tabu Search
  4. Genetic Algorithms
  5. Constraint Based Local Search

# Local Search Basics

1. Begin with a complete assignment to variables.
  - (A possibly infeasible solution to the problem)
2. Search by moving to other complete assignments.
  - (Explore the “neighborhood”)
3. Repeat the previous step until the assignment is “Good enough” wrt. feasible and/or optimal
  - (Termination condition)

# Traveling Salesman Problem (TSP)

- Given: A fully connected undirected graph.
  - Edge costs: distance between nodes

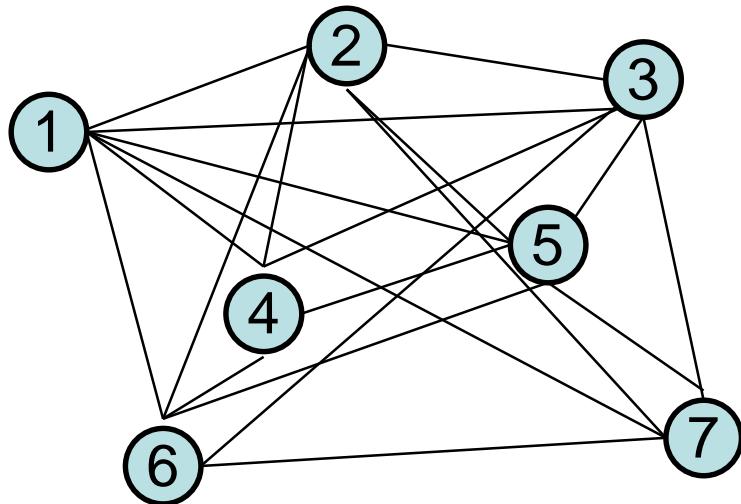


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- Task: Find the minimum cost path that visits all nodes and returns to the start.

# TSP: Initial Solution

- Local searches need an initial solution.
- We can store a TSP solution as a permutation.

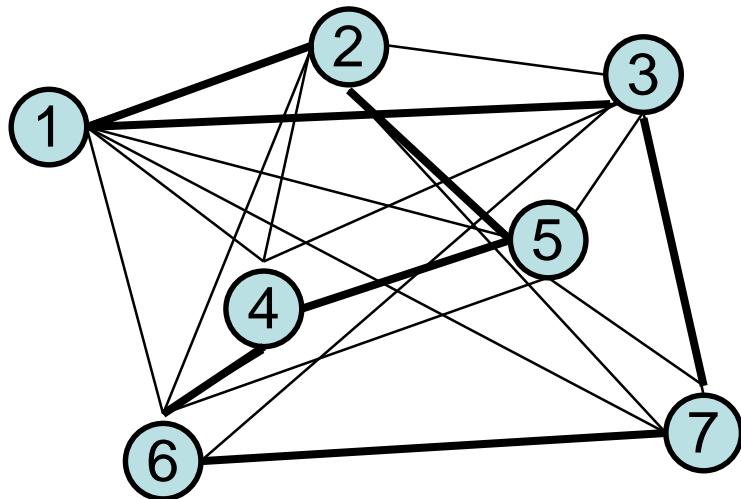


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- How would you construct an initial solution?

# TSP: Initial Solution

- Nearest neighbor heuristic

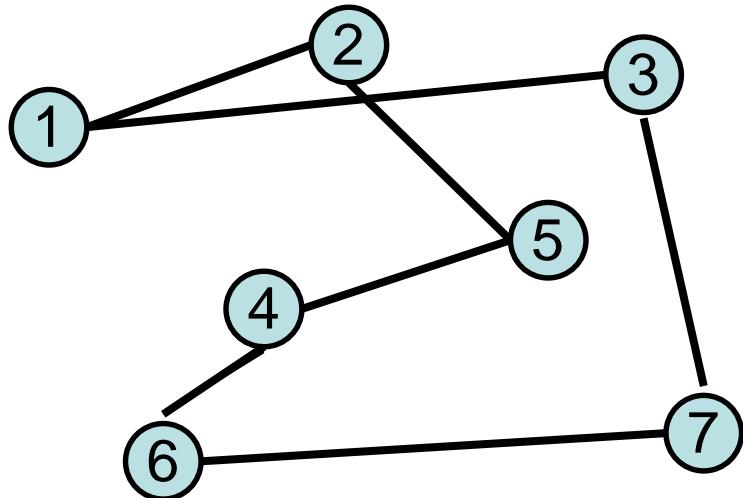


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- Initial solution: 1 2 5 4 6 7 3
- Cost: 27.4

# TSP: Neighborhoods

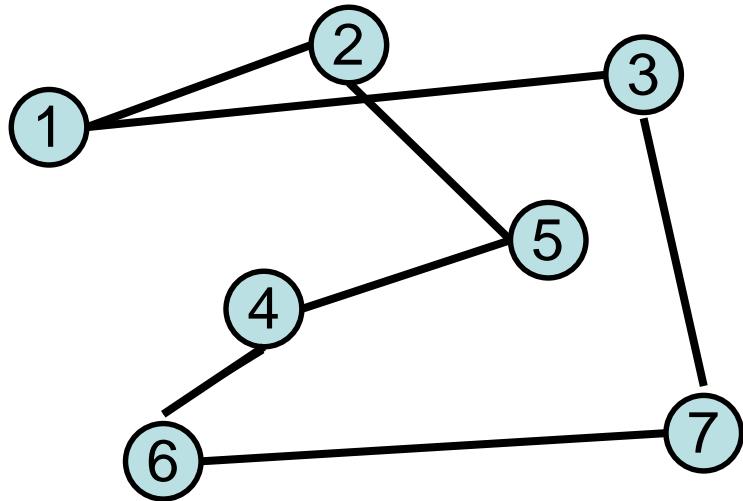
- How can we improve our initial solution?



	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

# TSP: 2-Opt Neighborhood

- 2-Opt removes 2 edges that cross each other and replaces them with non-crossing edges.

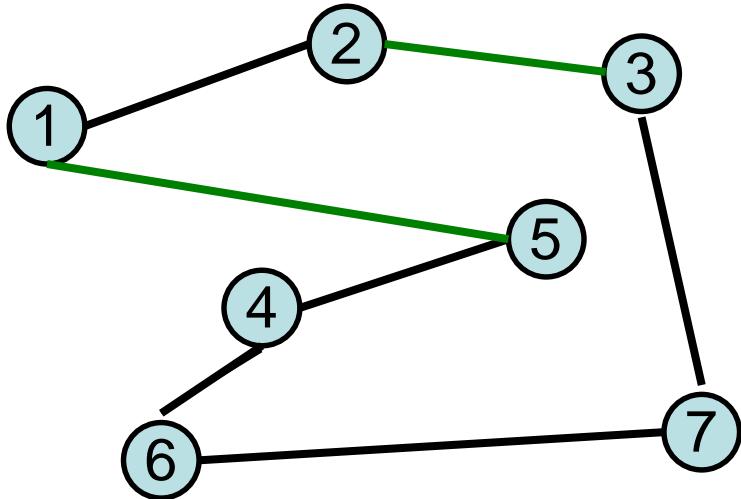


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- Which edges should we pick to remove?

# TSP: 2-Opt Neighborhood

- 2-Opt removes 2 edges that cross each other and replaces them with non-crossing edges.

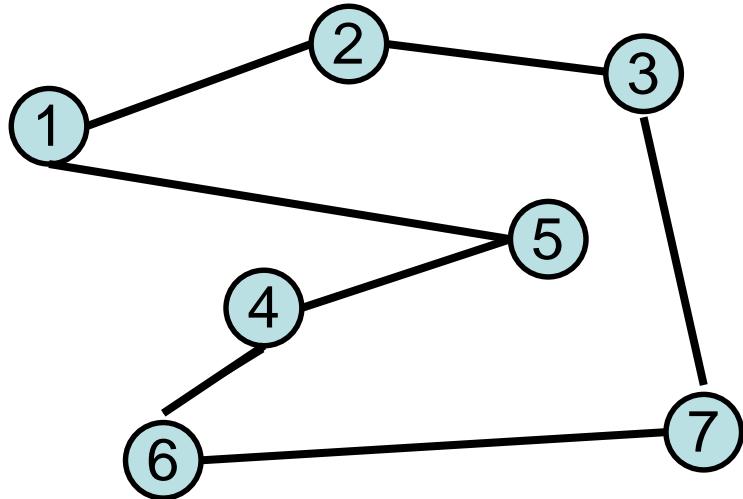


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- New solution: 1 2 3 7 6 4 5
- New cost: 26.7 (previous cost: 27.1)

# TSP: $k$ -Opt Neighborhood

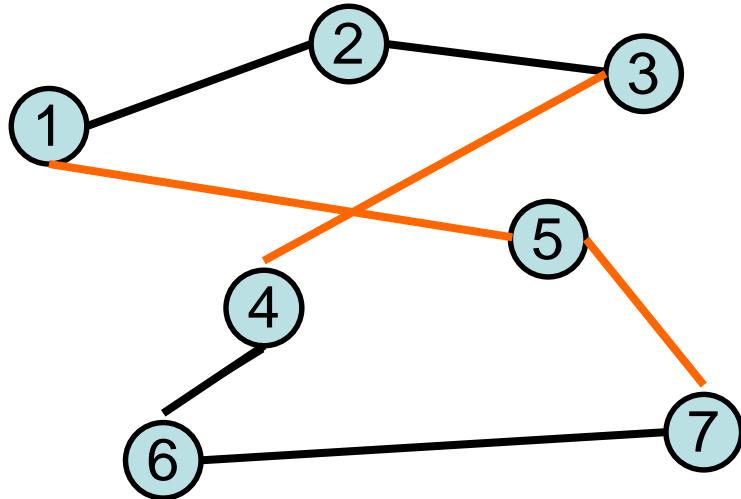
- Remove  $k$  edges and repair the path.
- Lets try  $k=3$



	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

# TSP: $k$ -Opt Neighborhood

- Remove  $k$  edges and repair the path.
- Lets try  $k=3$

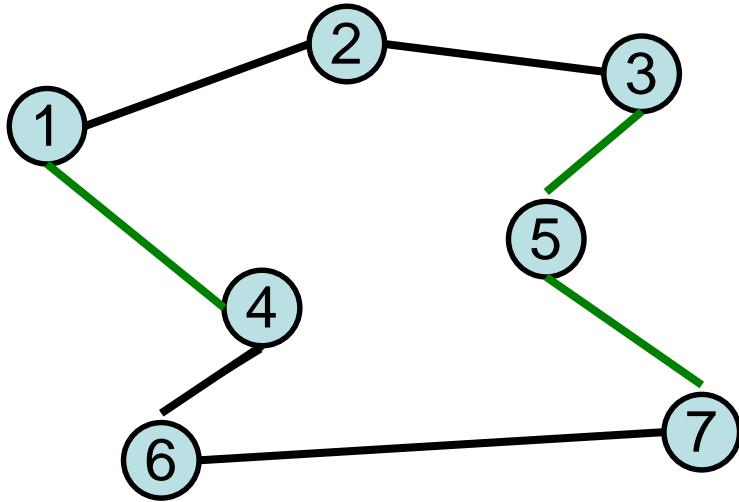


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- Possible solution: 1 2 3 4 6 7 5
- Cost: 27.1

# TSP: $k$ -Opt Neighborhood

- Remove  $k$  edges and repair the path.
- Lets try  $k=3$

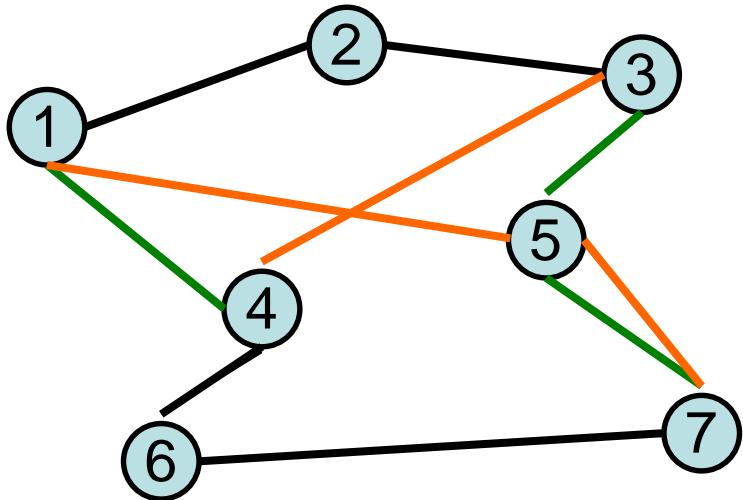


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- Possible solution: 1 2 3 5 7 6 4
- Cost: 23.8

# TSP: Neighbor Selection

- Which neighbor should we choose?

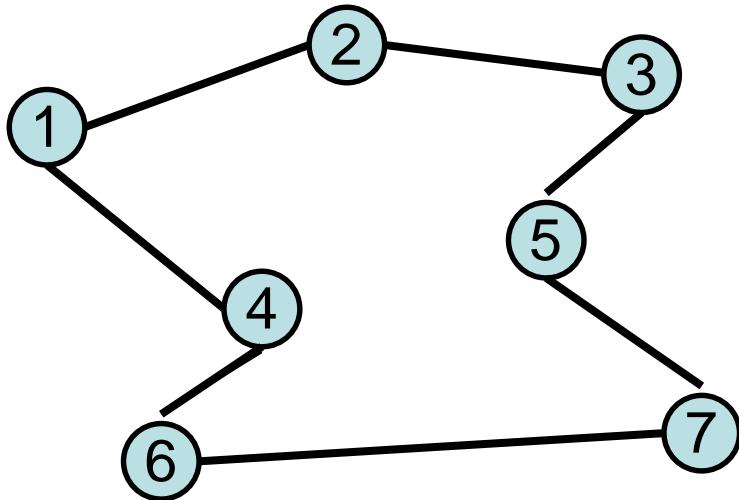


	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- 23.8 vs. 27.1

# TSP: Termination

- When should we stop performing improvements?



	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

# TSP: Termination

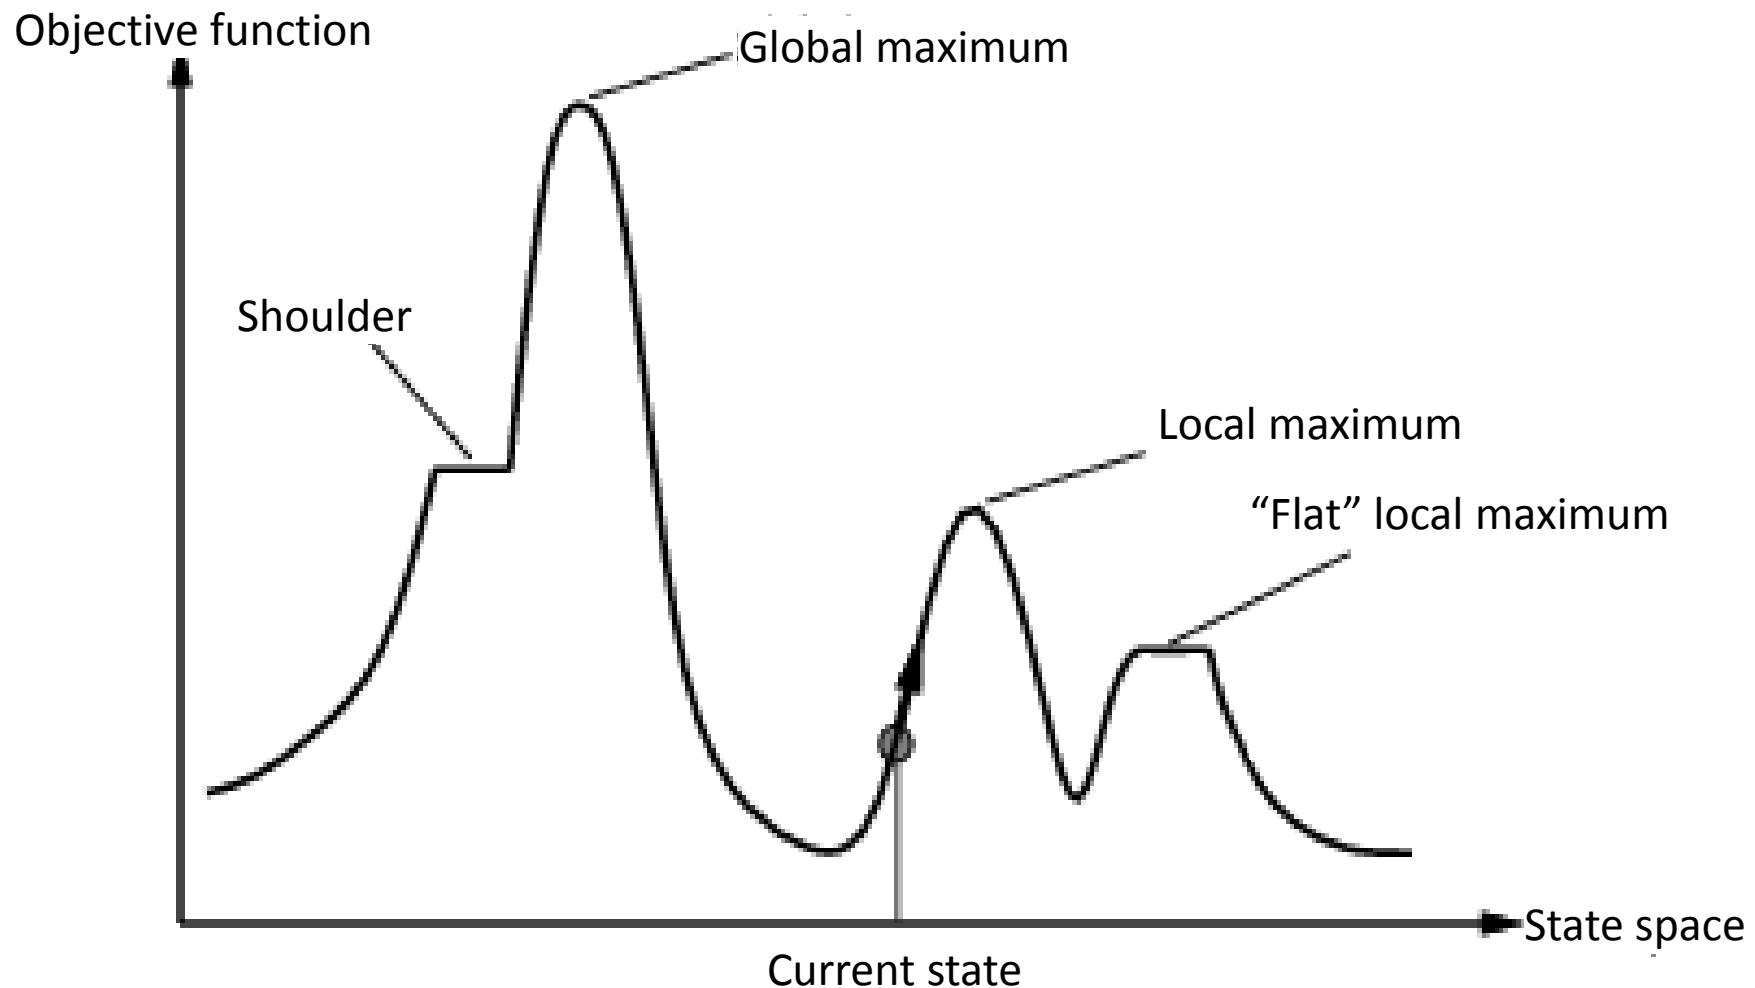
- When should we stop performing improvements?
- Budget based:
  - Number of iterations
  - CPU time
- Solution quality
  - Within  $X\%$  of a lower bound
  - Business requirements satisfied
- Convergence criteria:
  - No improvement in last  $Y$  iterations.
  - Average improvement below threshold  $\varepsilon$

# Hill Climbing

# Hill Climbing Algorithm

```
function HILL-CLIMBING( problem ) return a state that is a local maximum
  current  $\leftarrow$  MAKE-NODE( problem.INITIAL-STATE )
  loop do
    neighbor  $\leftarrow$  a highest-value successor of current
    if neighbor.VALUE  $\leq$  current.VALUE then return current.STATE
    current  $\leftarrow$  neighbor
```

# State-Space Landscape



# Hill Climbing: Pro & Con

- Advantages
  - Fast convergence to a local maximum
  - Often results in good (but not optimal) solutions
- Disadvantages
  - Gets stuck in local maxima
  - Gets stuck on shoulders and plateaus

# Exploitation vs. Exploration

- Exploitation (Intensification)
  - Greedy; always select most improving neighbor
- Exploration (Diversification)
  - Also select less improving and non-improving neighbors

Hill Climbing

Random Walk



Exploitation

Exploration

# Escaping Local Maxima

- Variations of Hill-Climbing
  - **Sideways move:** allow non-improving moves to traverse plateaus.
  - **Stochastic Hill-Climbing:** random choice of uphill moves.
  - **First-Choice Hill-Climbing:** random generation of neighbors.
  - **Random-restart Hill-Climbing:** restart the search from a different initial state

# Simulated Annealing

# Simulated Annealing



- Inspired by annealing in metals
  - Used to harden metals by gradually cooling them, allowing atoms to find a low-energy crystalline state
- Idea:
  - Escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

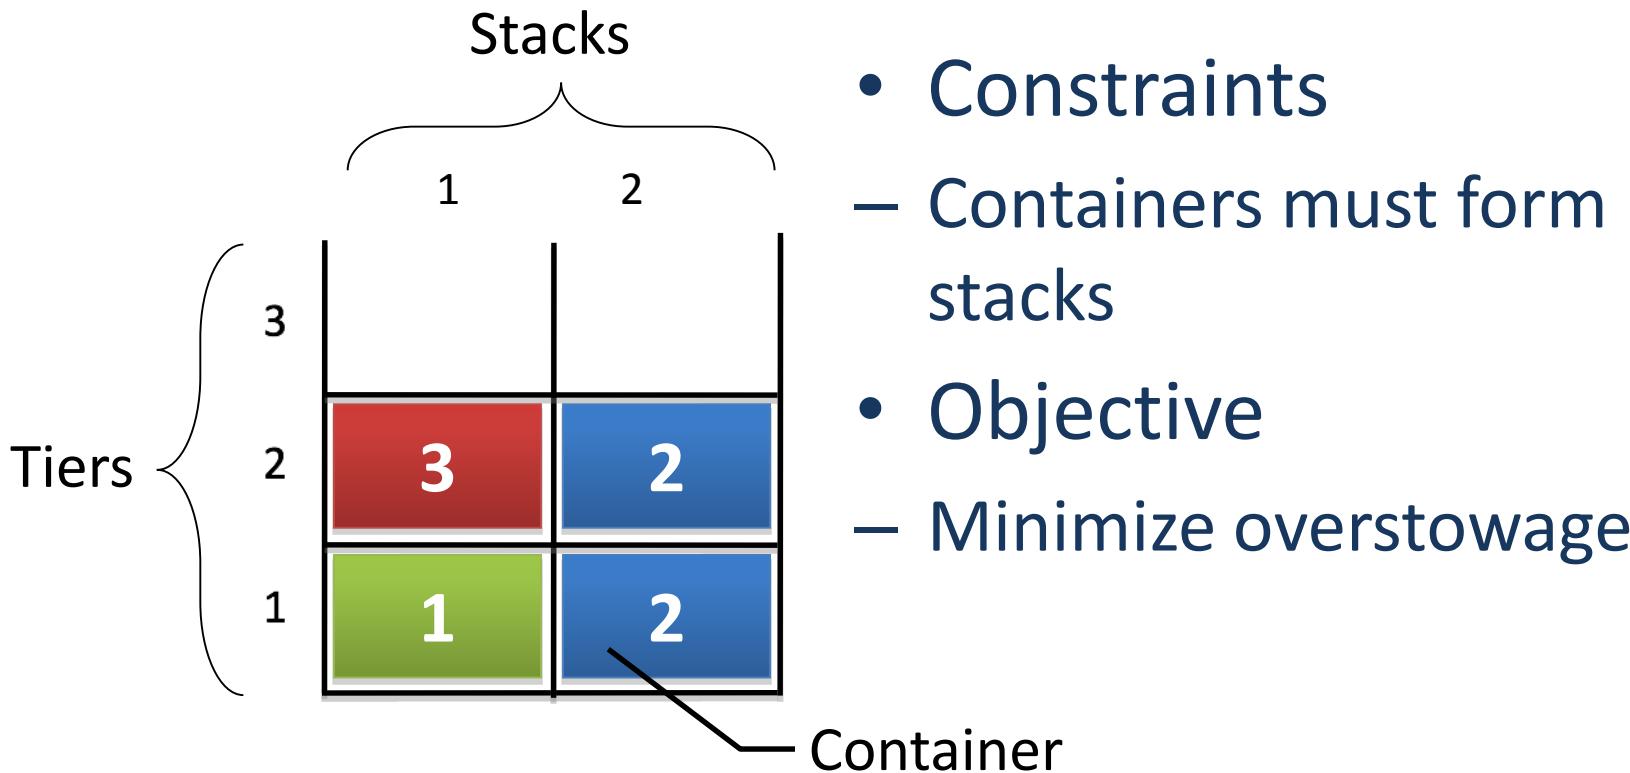
# Simulated Annealing Implementation

```
function SIMULATED-ANNEALING( problem, schedule ) returns a solution state
  input: problem, a problem
         schedule, a mapping from time to “temperature”

  current  $\leftarrow$  MAKE-NODE( problem.INITIAL-STATE )
  for t = 1 to  $\infty$  do
    T  $\leftarrow$  schedule(t)
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected neighbor of current
     $\Delta E \leftarrow$  next.VALUE – current.VALUE
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

# Ex. Container Stowage Problem

- Given a feasible container configuration, find the one which minimizes overstowage.

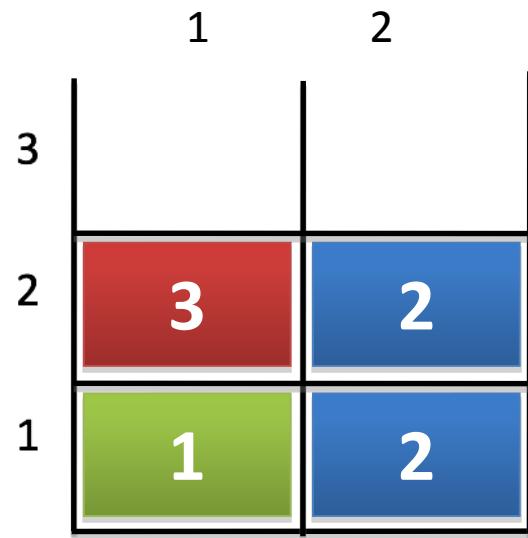


# Ex. Container Stowage Problem

- State: A container configuration
- Neighborhood: Container swaps
- Objective function (*Value*): Number of overstowed containers.
- Termination criteria: *Value* = 0
- $\Delta E := \text{current.VALUE} - \text{next.VALUE}$

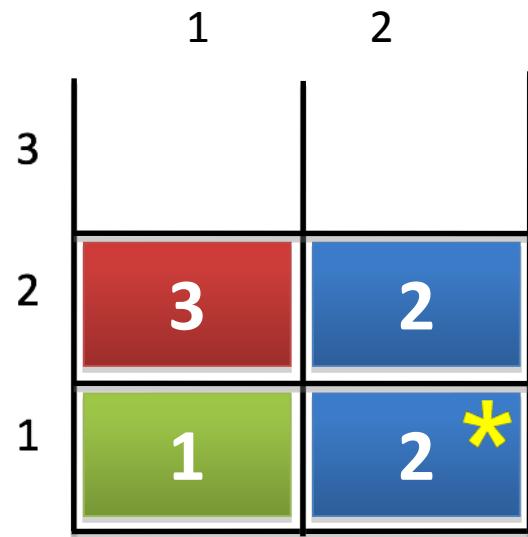
Obs: Minimization problem!

# Ex. Container Stowage Problem



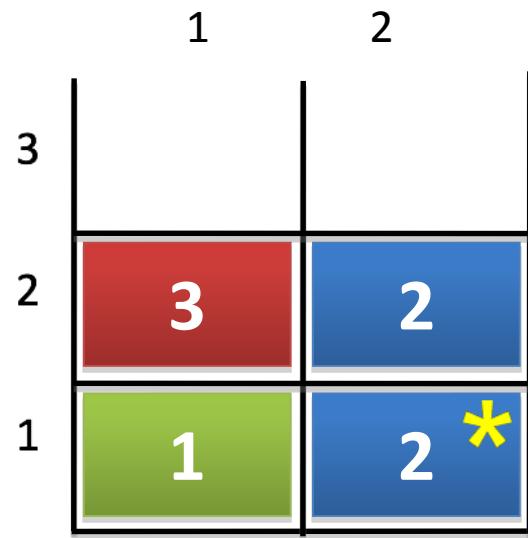
Objective: 1  
Temperature: 10

# Ex. Container Stowage Problem

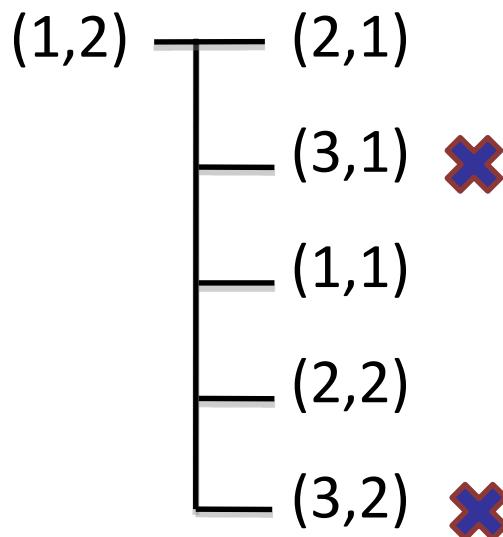


Objective: 1  
Temperature: 10

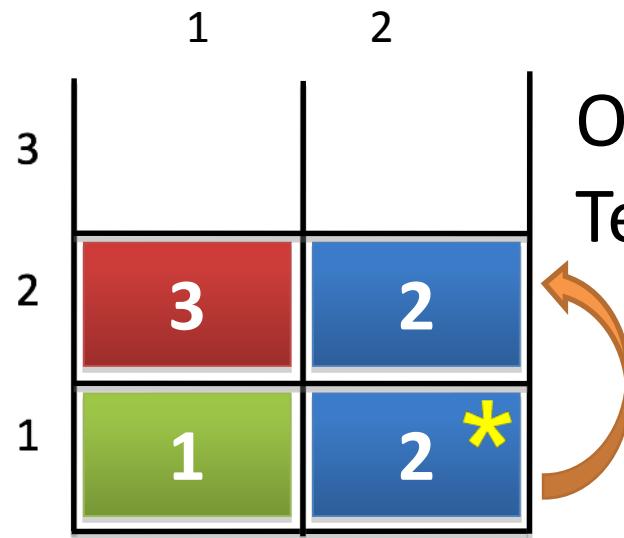
# Ex. Container Stowage Problem



Objective: 1  
Temperature: 10



# Ex. Container Stowage Problem



Objective: 1  
Temperature: 10

(1,2) — (2,1)

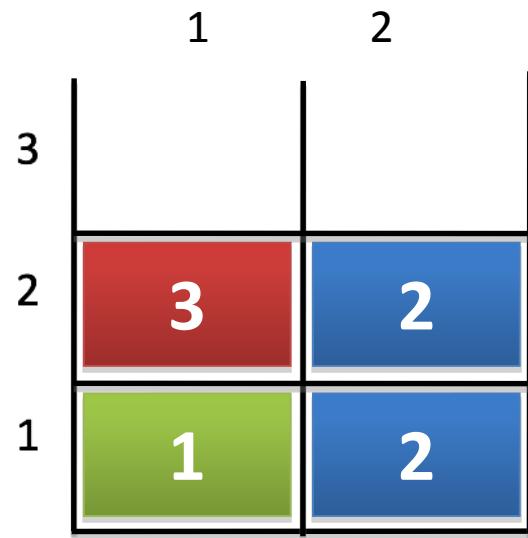
(3,1) ✗

(1,1)

(2,2)  $\Delta E = 0 : e^{\frac{0}{10}} = 1$  ✓

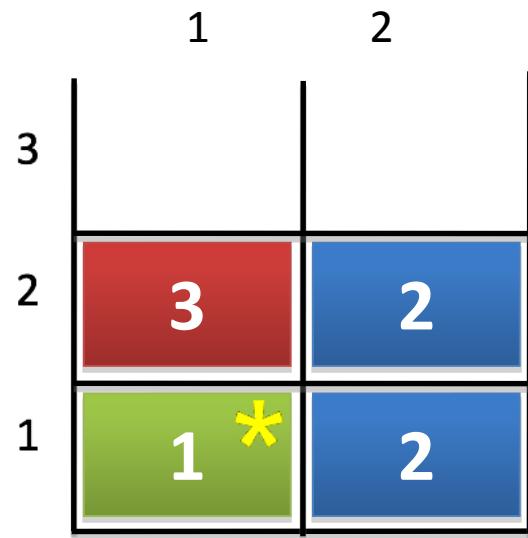
(3,2) ✗

# Ex. Container Stowage Problem



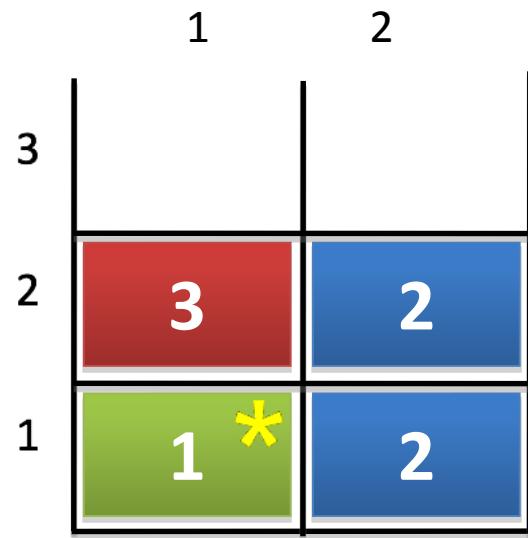
Objective: 1  
Temperature: 2

# Ex. Container Stowage Problem

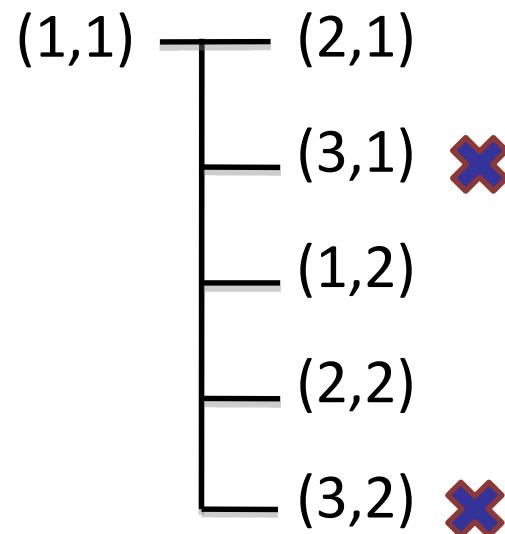


Objective: 1  
Temperature: 2

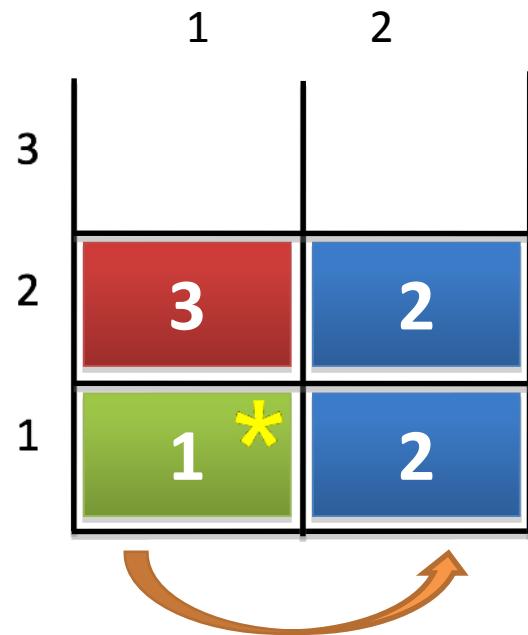
# Ex. Container Stowage Problem



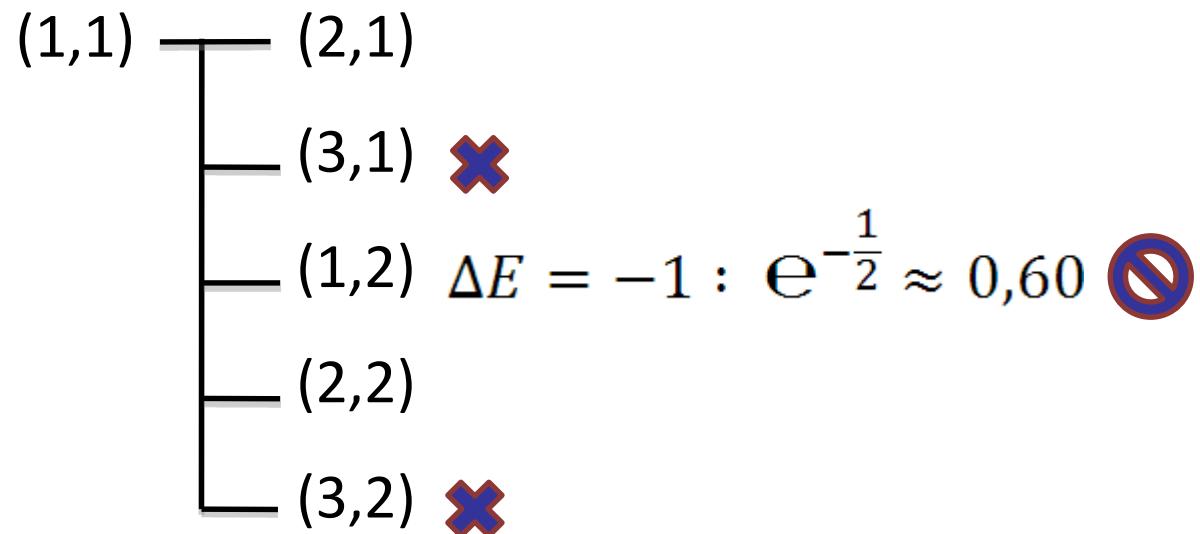
Objective: 1  
Temperature: 2



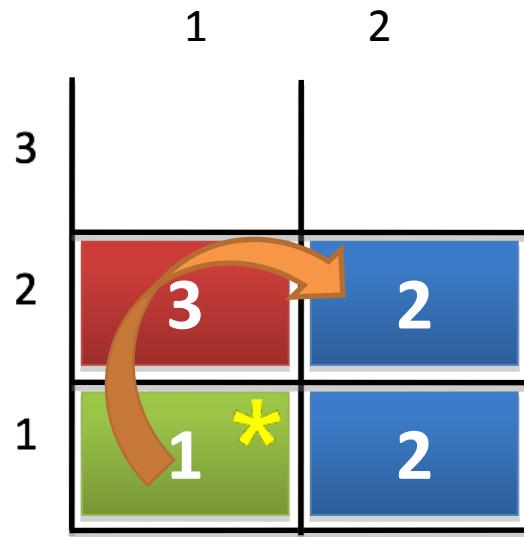
# Ex. Container Stowage Problem



Objective: 1  
Temperature: 2



# Ex. Container Stowage Problem



Objective: 1  
Temperature: 2

(1,1) — (2,1)

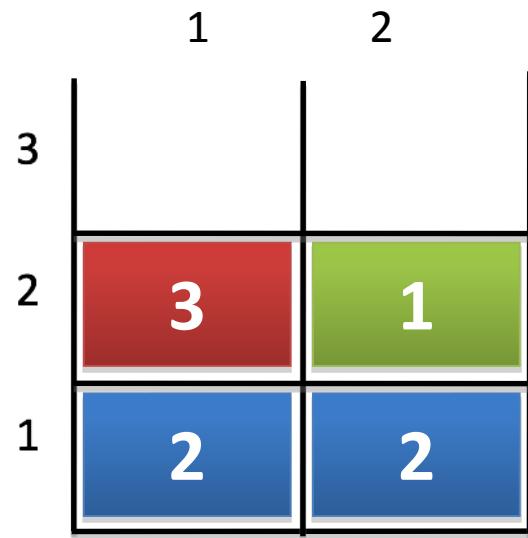
(3,1) ✘

(1,2)  $\Delta E = -1 : e^{-\frac{1}{2}} \approx 0,60$  ✘

(2,2)  $\Delta E = 0 : e^{\frac{0}{2}} = 1$  ✓

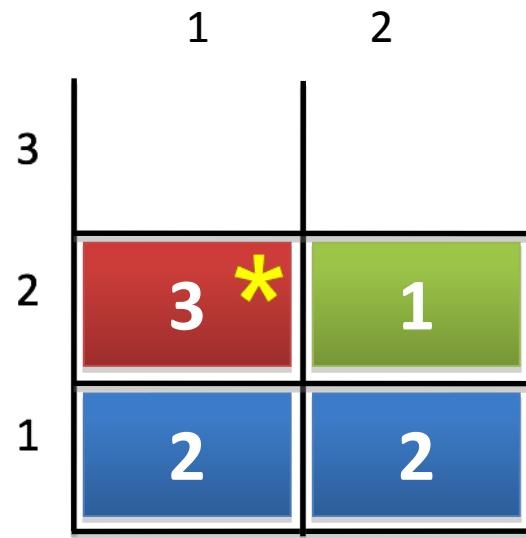
(3,2) ✘

# Ex. Container Stowage Problem



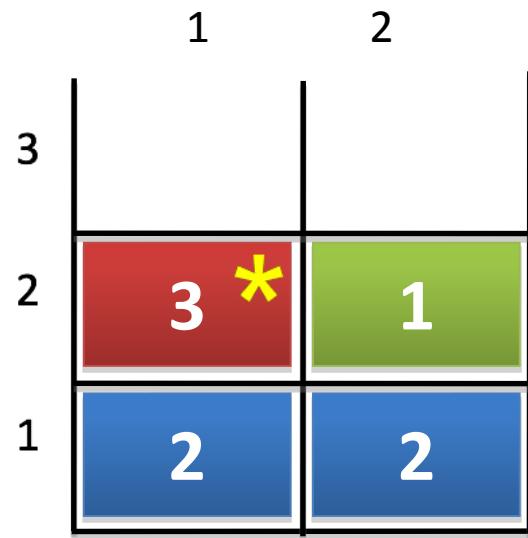
Objective: 1  
Temperature: 0,5

# Ex. Container Stowage Problem

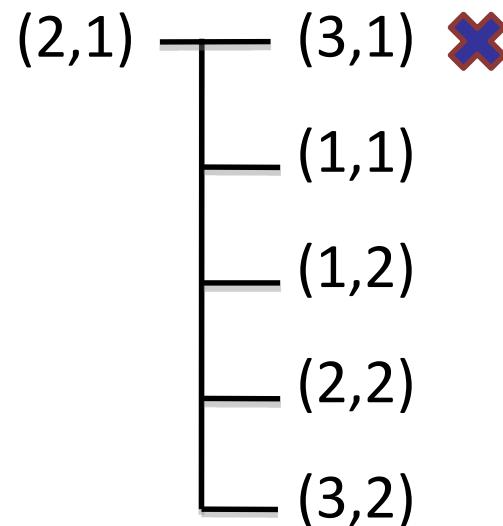


Objective: 1  
Temperature: 0,5

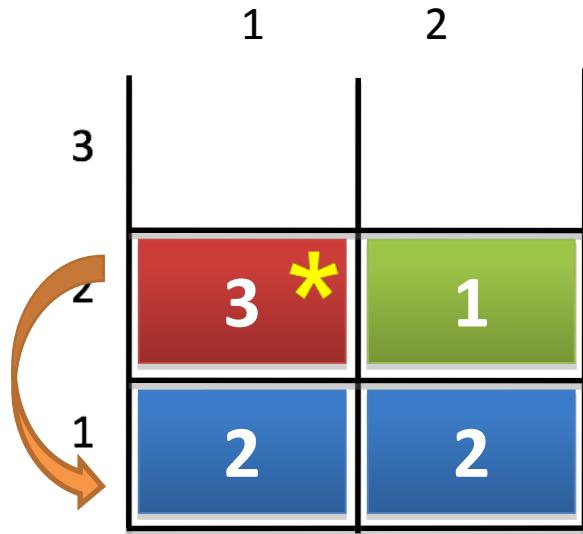
# Ex. Container Stowage Problem



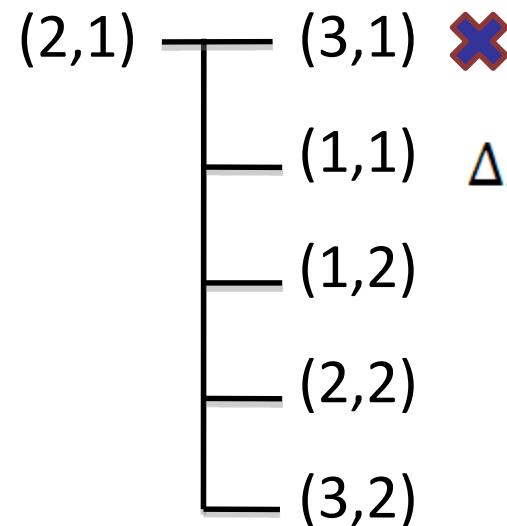
Objective: 1  
Temperature: 0,5



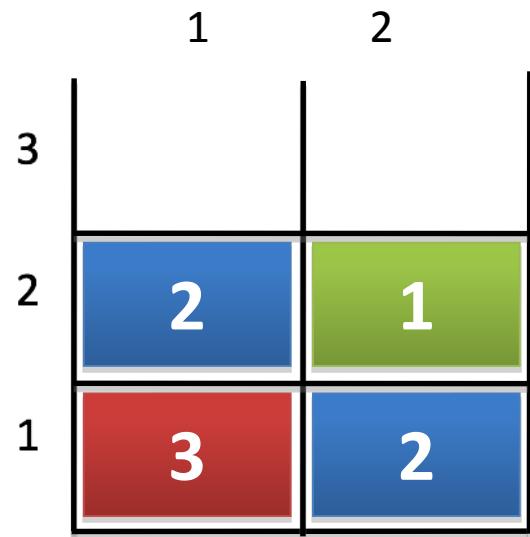
# Ex. Container Stowage Problem



Objective: 1  
Temperature: 0,5



# Ex. Container Stowage Problem



Objective: 0  
Temperature: 0,2



Lower bound ! Terminate!

# Tabu Search (TS)

# Tabu Search

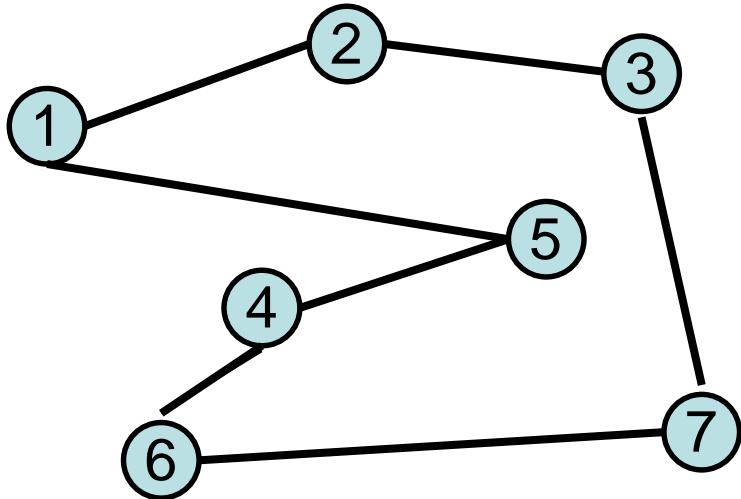
- A **tabu** (also spelled **taboo**) is a strong social **prohibition** (or **ban**) against words, objects, actions, or discussions that are considered undesirable or offensive by a group, culture, society, or community.
  - “Taboo” Wikipedia

# Tabu Search

- Idea:
  - Accept the best neighbor at each iteration
  - Avoid previously seen solutions by keeping a memory (tabu list) of previous states

# Tabu Search: TSP Example

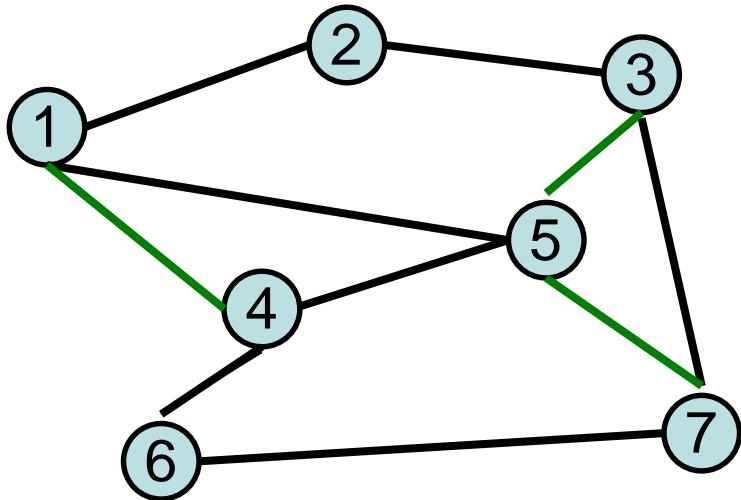
- What could we store in our tabu list?



	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

# Tabu Search: TSP Example

- We could store an entire solution
- Or, just store the changes we made



	1	2	3	4	5	6	7
1	0	2.2	5	2.8	4.1	5	8.5
2	2.2	0	3	3	2.8	6	8
3	5	3	0	4.2	2.2	7.2	7
4	2.8	3	4.2	0	2.2	3.1	5.7
5	4.1	2.8	2.2	2.2	0	5	5.4
6	5	6	7.2	3.1	5	0	5.1
7	8.5	8	7	5.7	5.4	5.1	0

- Tabu list:
  1.  $-(1,5), -(4,5), -(3,7)$
  2.  $+(1,4), +(5,7), +(3,5)$

# Tabu Search Implementation

```
function TABU-SEARCH( problem ) returns a solution state
  inputs: problem, a problem

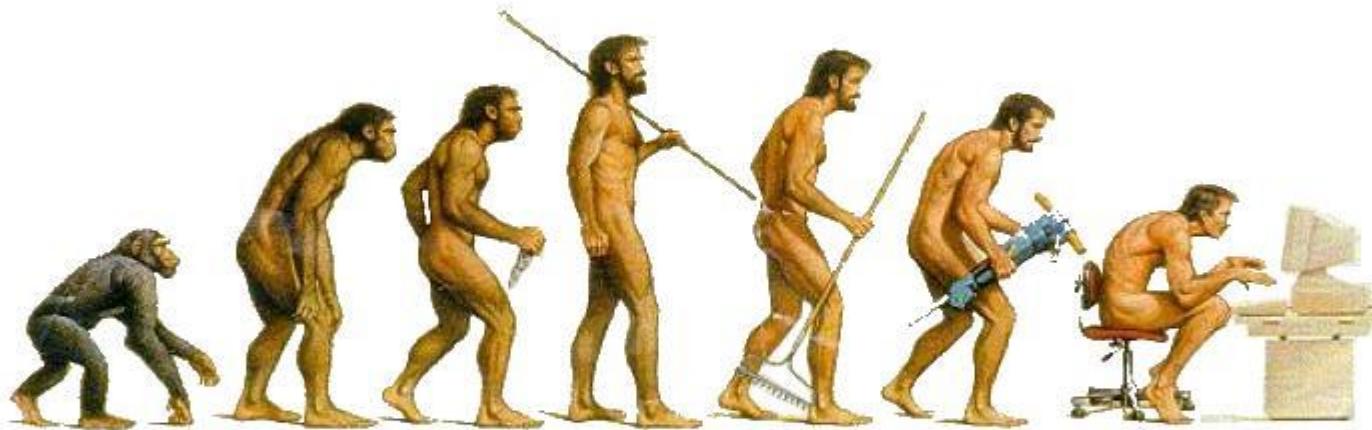
  current  $\leftarrow$  MAKE-NODE( problem.INITIAL-STATE)
  best  $\leftarrow$  current
  T  $\leftarrow$  Empty tabu list
  while ( termination criterion not satisfied ) do
    current  $\leftarrow$  a highest-valued successor of current legal wrt. T
    if current.VALUE > best.VALUE then
      best  $\leftarrow$  current
      ADD( T , ACTION-TO( current ) )
  return best
```

# Tabu Search Variations

- Tabu list
  - Variable length
  - Aspiration criteria (override tabu, e.g. if improving move)
- Probabilistic tabu search
  - Only consider a random sample of the neighborhood

# Genetic Algorithms

## Evolution



(or is it?)

# Genetic Algorithms

- **Individual:** A variable assignment (“Genome”)
  - Often represented by a bit string
- **Population:**  $n$  individuals
  - Initialize to randomly generated states
- **Fitness function:** Evaluates the “fitness” of an individual
- **Selection:** Identify the most fit members of a population
- **Crossover:** Form new individuals out of multiple individuals
- **Mutation:** Randomly change a value in an individual’s genome

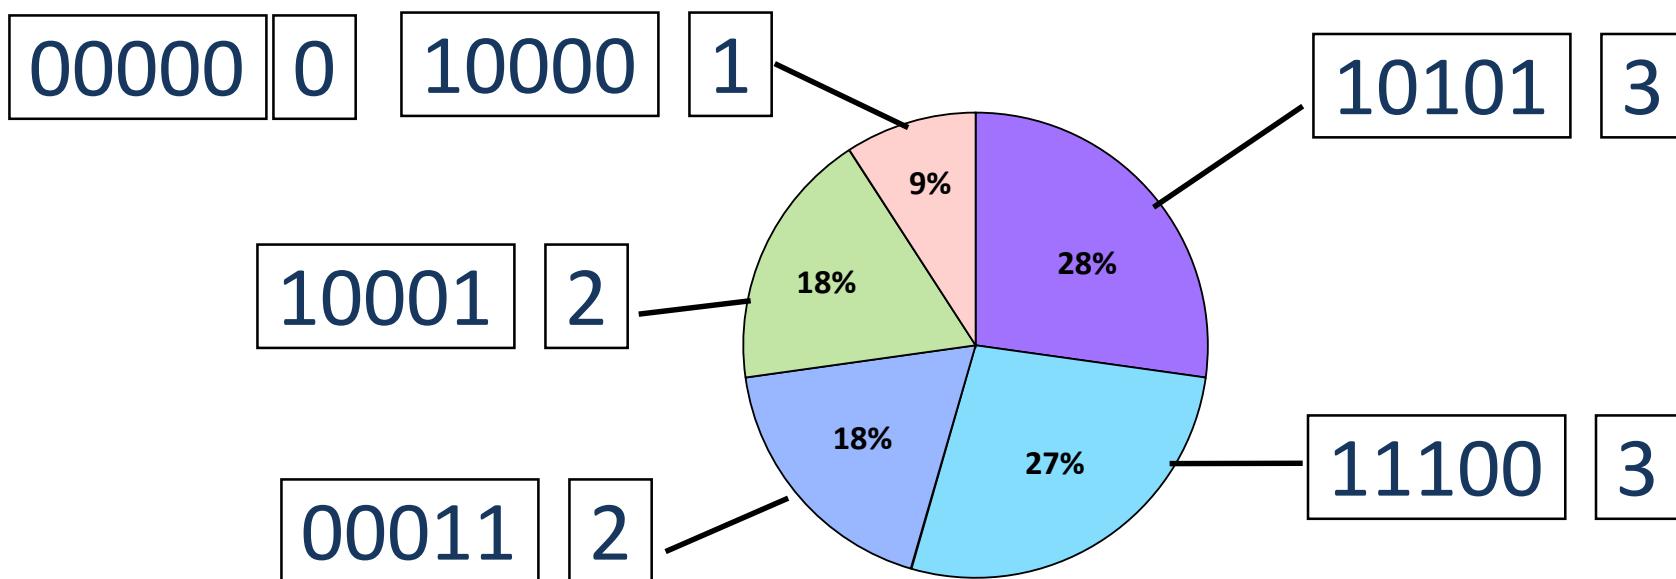
# Example - 1 max

- **Genome:** bit string of length 5
- **Fitness function:**  $f(x) = \# \text{ of } 1\text{s in the bit string}$ 
  - e.g.  $f(00110) = 2, f(111111) = 5$
- Goal: Maximize  $f(x)$
- Step 1: Initialize population

10000	10101	10001	00011	11100	00000
-------	-------	-------	-------	-------	-------

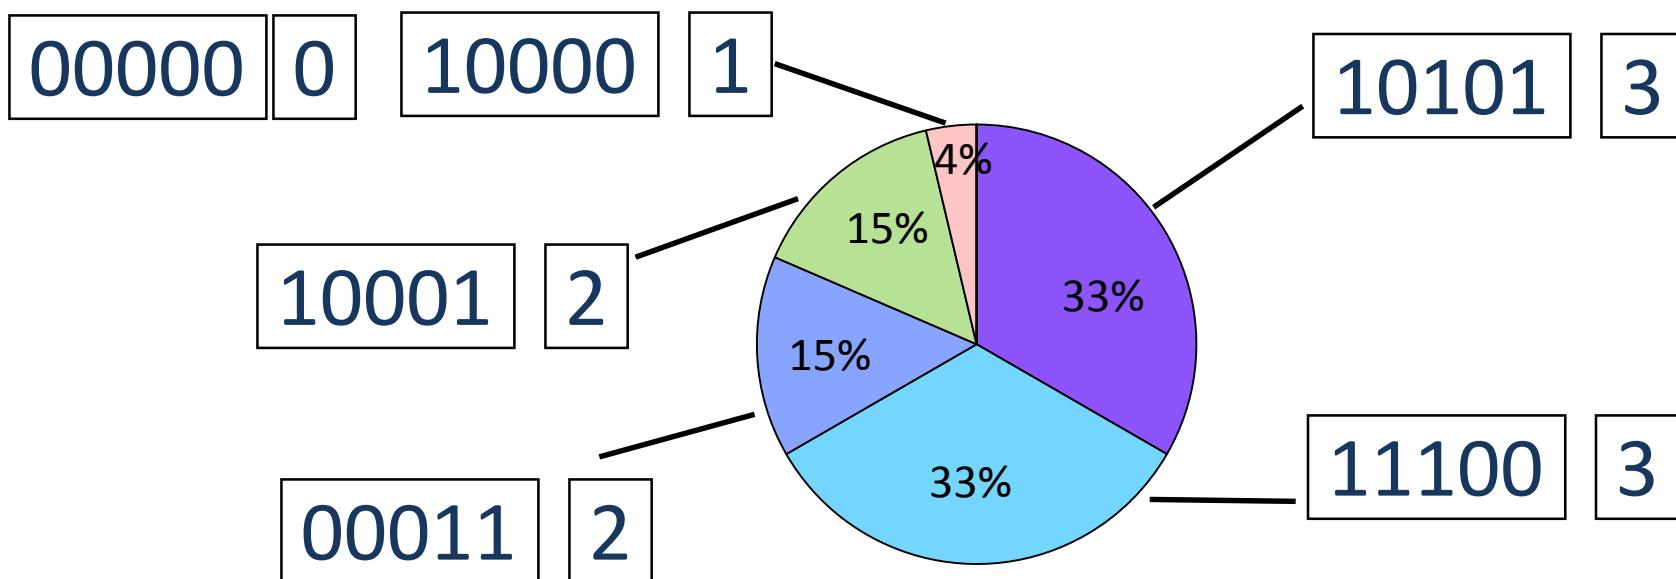
# Example - 1 max

- Step 2: Evaluate the population's fitness
- Step 3: Selection. (Roulette wheel selection)
  - Select genomes with probability proportional to their fitness



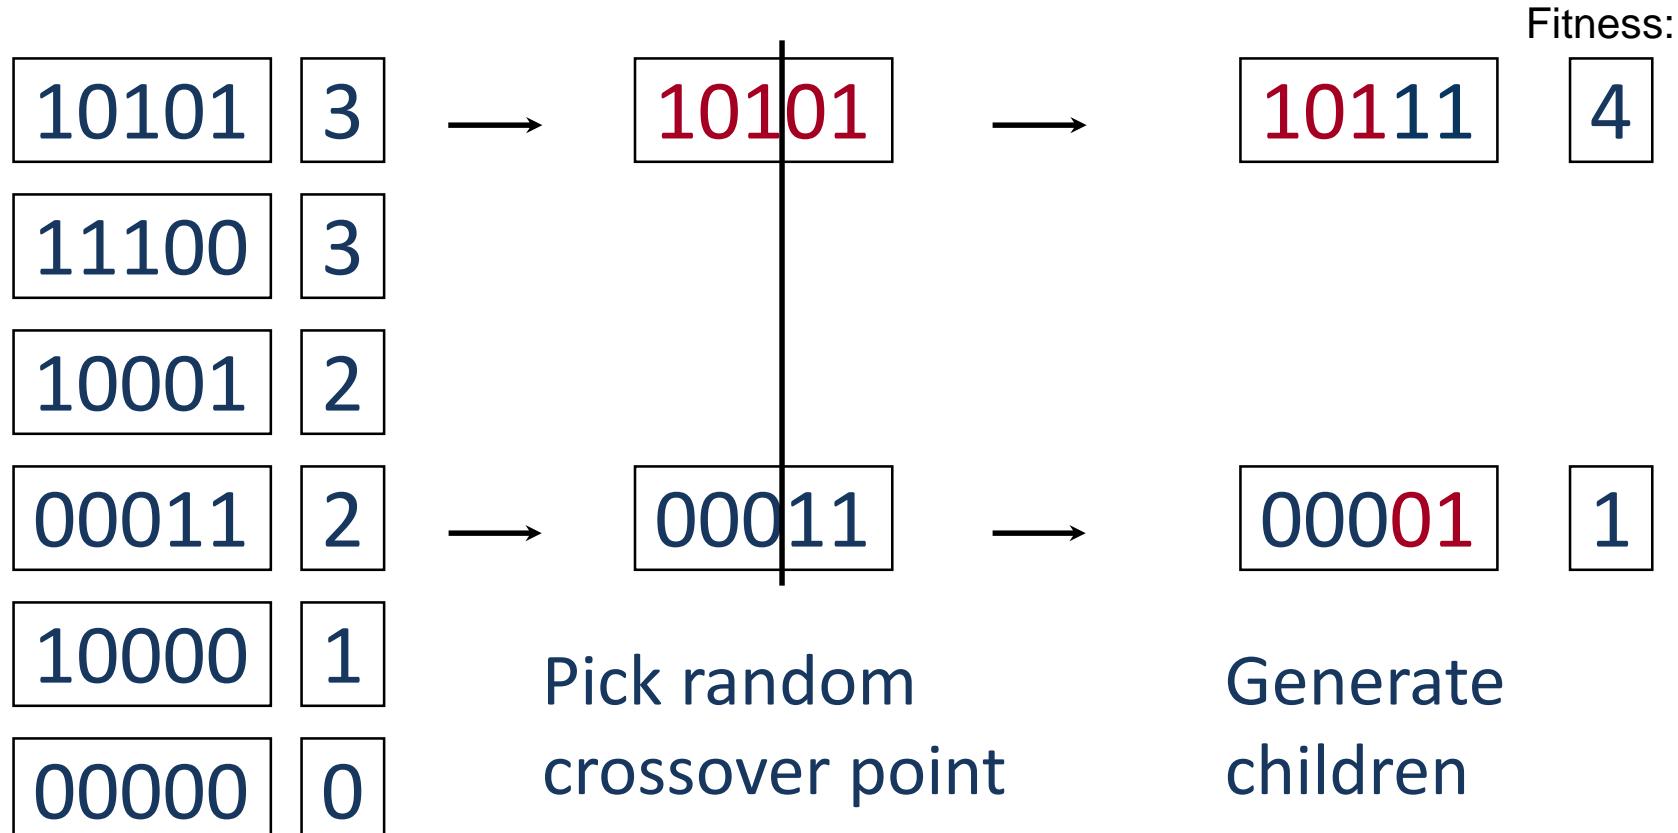
# Example - 1 max

- Fitness scaling can improve performance.
  - Square (or cube) the fitness of each individual before performing selection



# Example - 1 max

- Select genomes and perform crossover.



# Example - 1 max

- Continue selection until new population is formed
  - Individuals are often allowed to be part of multiple selections.
- Step 4: Mutation
  - With some probability, usually  $< 0.1$  make a small change in the genome

00001



10001

Flip random bit

# Example - 1 max

- Population at the end of the generation:

10111	4
11100	3
10001	2
10001	2
10000	1
00000	0

- Unless a termination criteria has been reached, continue with the next generation.

# Genetic Algorithms in Practice

- Necessary that “Genome” forms **meaningful** components of the problem
- Population size difficult to determine
  - Between 25 and 100, depending on the problem
- Terminate criteria vary
  - Little change in average fitness of the population over last  $n$  generations, where  $n \approx 5$
- Choice of crossover operator extremely important
  - Single point vs. multiple point, etc.

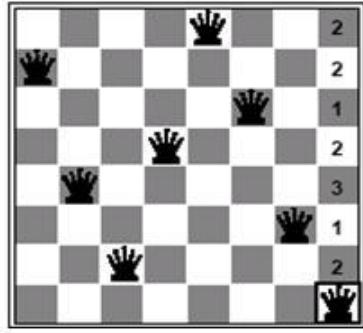
Is multi-restart LS better than population-based LS?

# Constraint Based Local Search

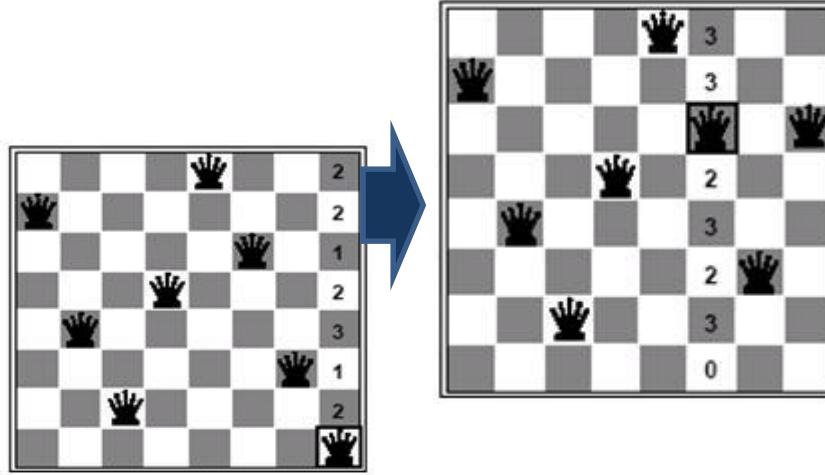
# Constraint Based Local Search

- Constraint satisfaction problems
  - allow states with unsatisfied constraints
  - operators **reassign** variable values
- Variable selection: randomly select any conflicted variable
- Value selection: *min-conflicts heuristic*
  - Select new value that results in a minimum number of conflicts with the other variables

# Min-conflict algorithm



# Min-conflict algorithm



# Min-conflict algorithm

